



# Y O L O

MALEKE CHAKER  
BİLEL EL HADDAD



# Introduction

In the vast realm of computer vision, Object Detection takes center stage, allowing computers to understand images in a way that goes beyond simple classification. It's like giving machines the ability to not only recognize objects but also pinpoint their exact locations within an image.

# Introduction

From self-driving cars and security cameras to healthcare tools and your self-checkout lane at the grocery store, Object Detection is everywhere. It helps cars see pedestrians, keeps an eye on who's coming and going, assists in medical diagnoses, and even speeds up your trip through the store. Let's take a journey into the domain of Object Detection and discover the state-of-art object detection model YOLO and its variants.

# Object Detection :

**Classification**



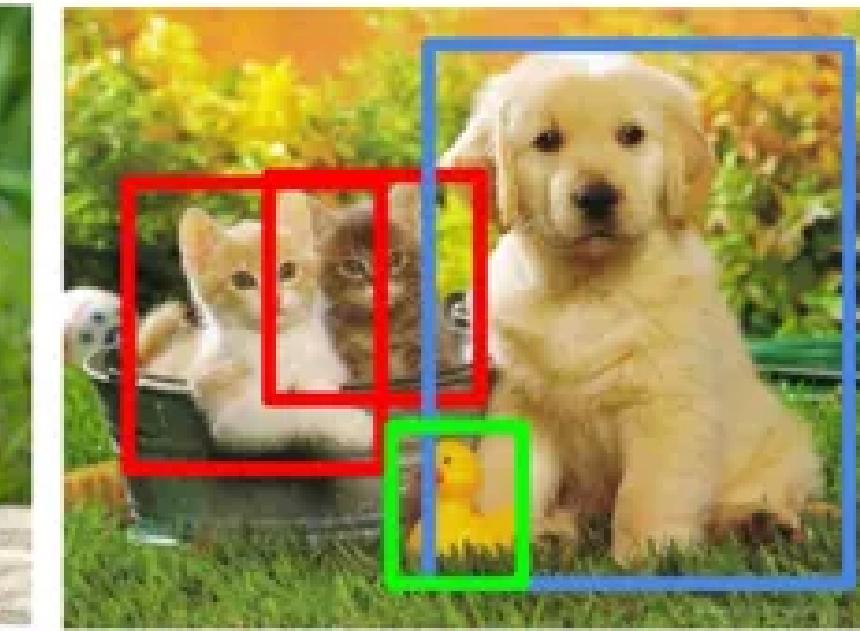
CAT

**Classification + Localization**



CAT

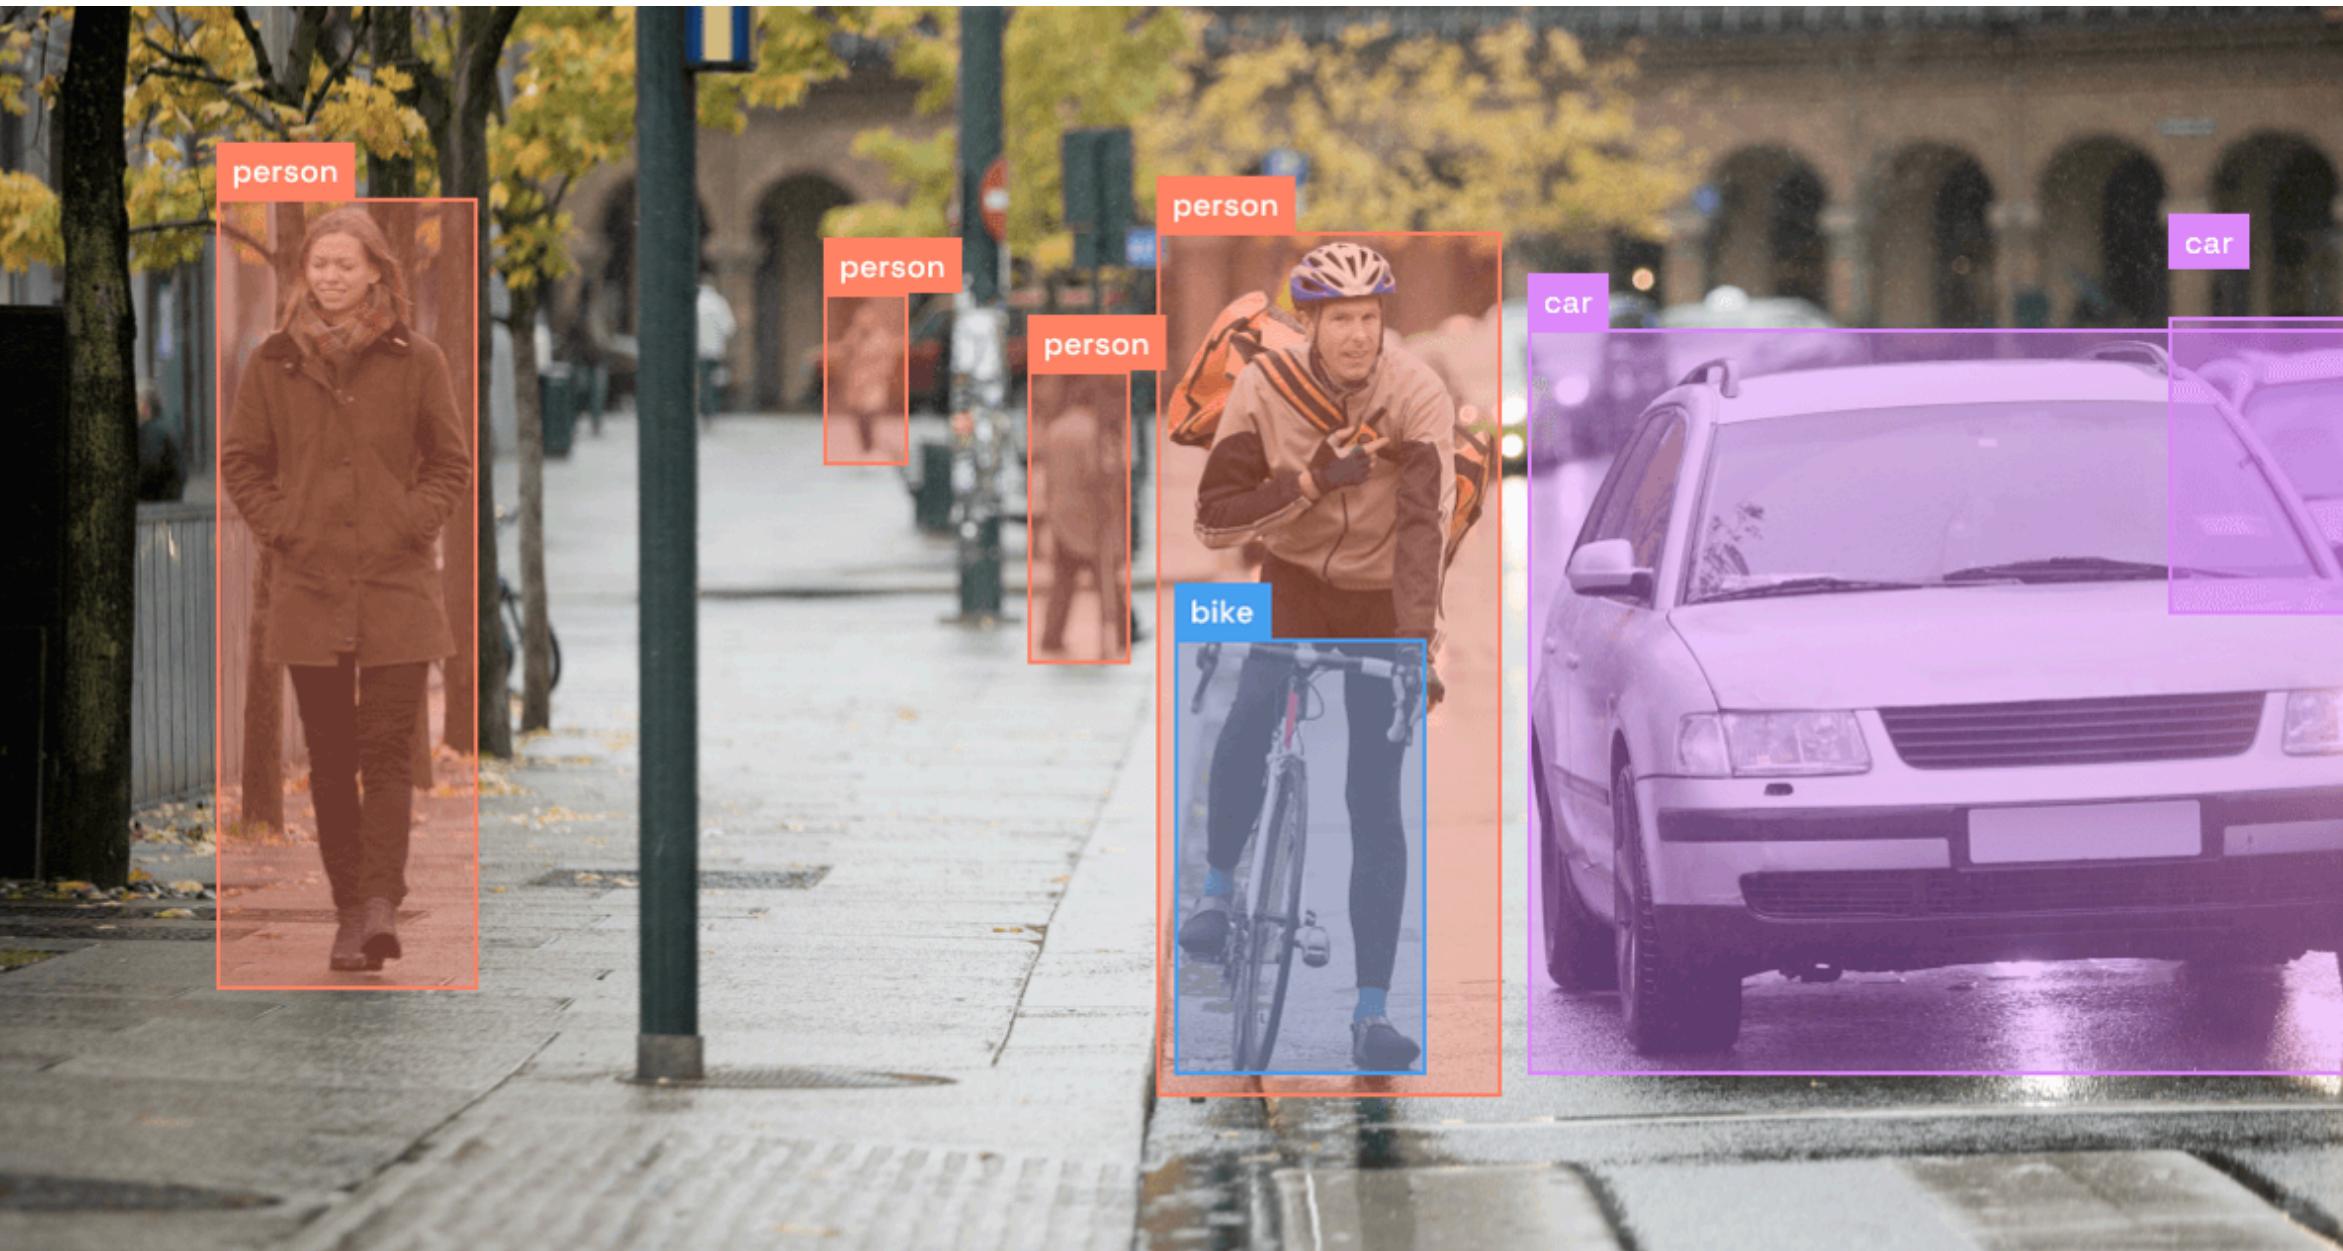
**Object Detection**



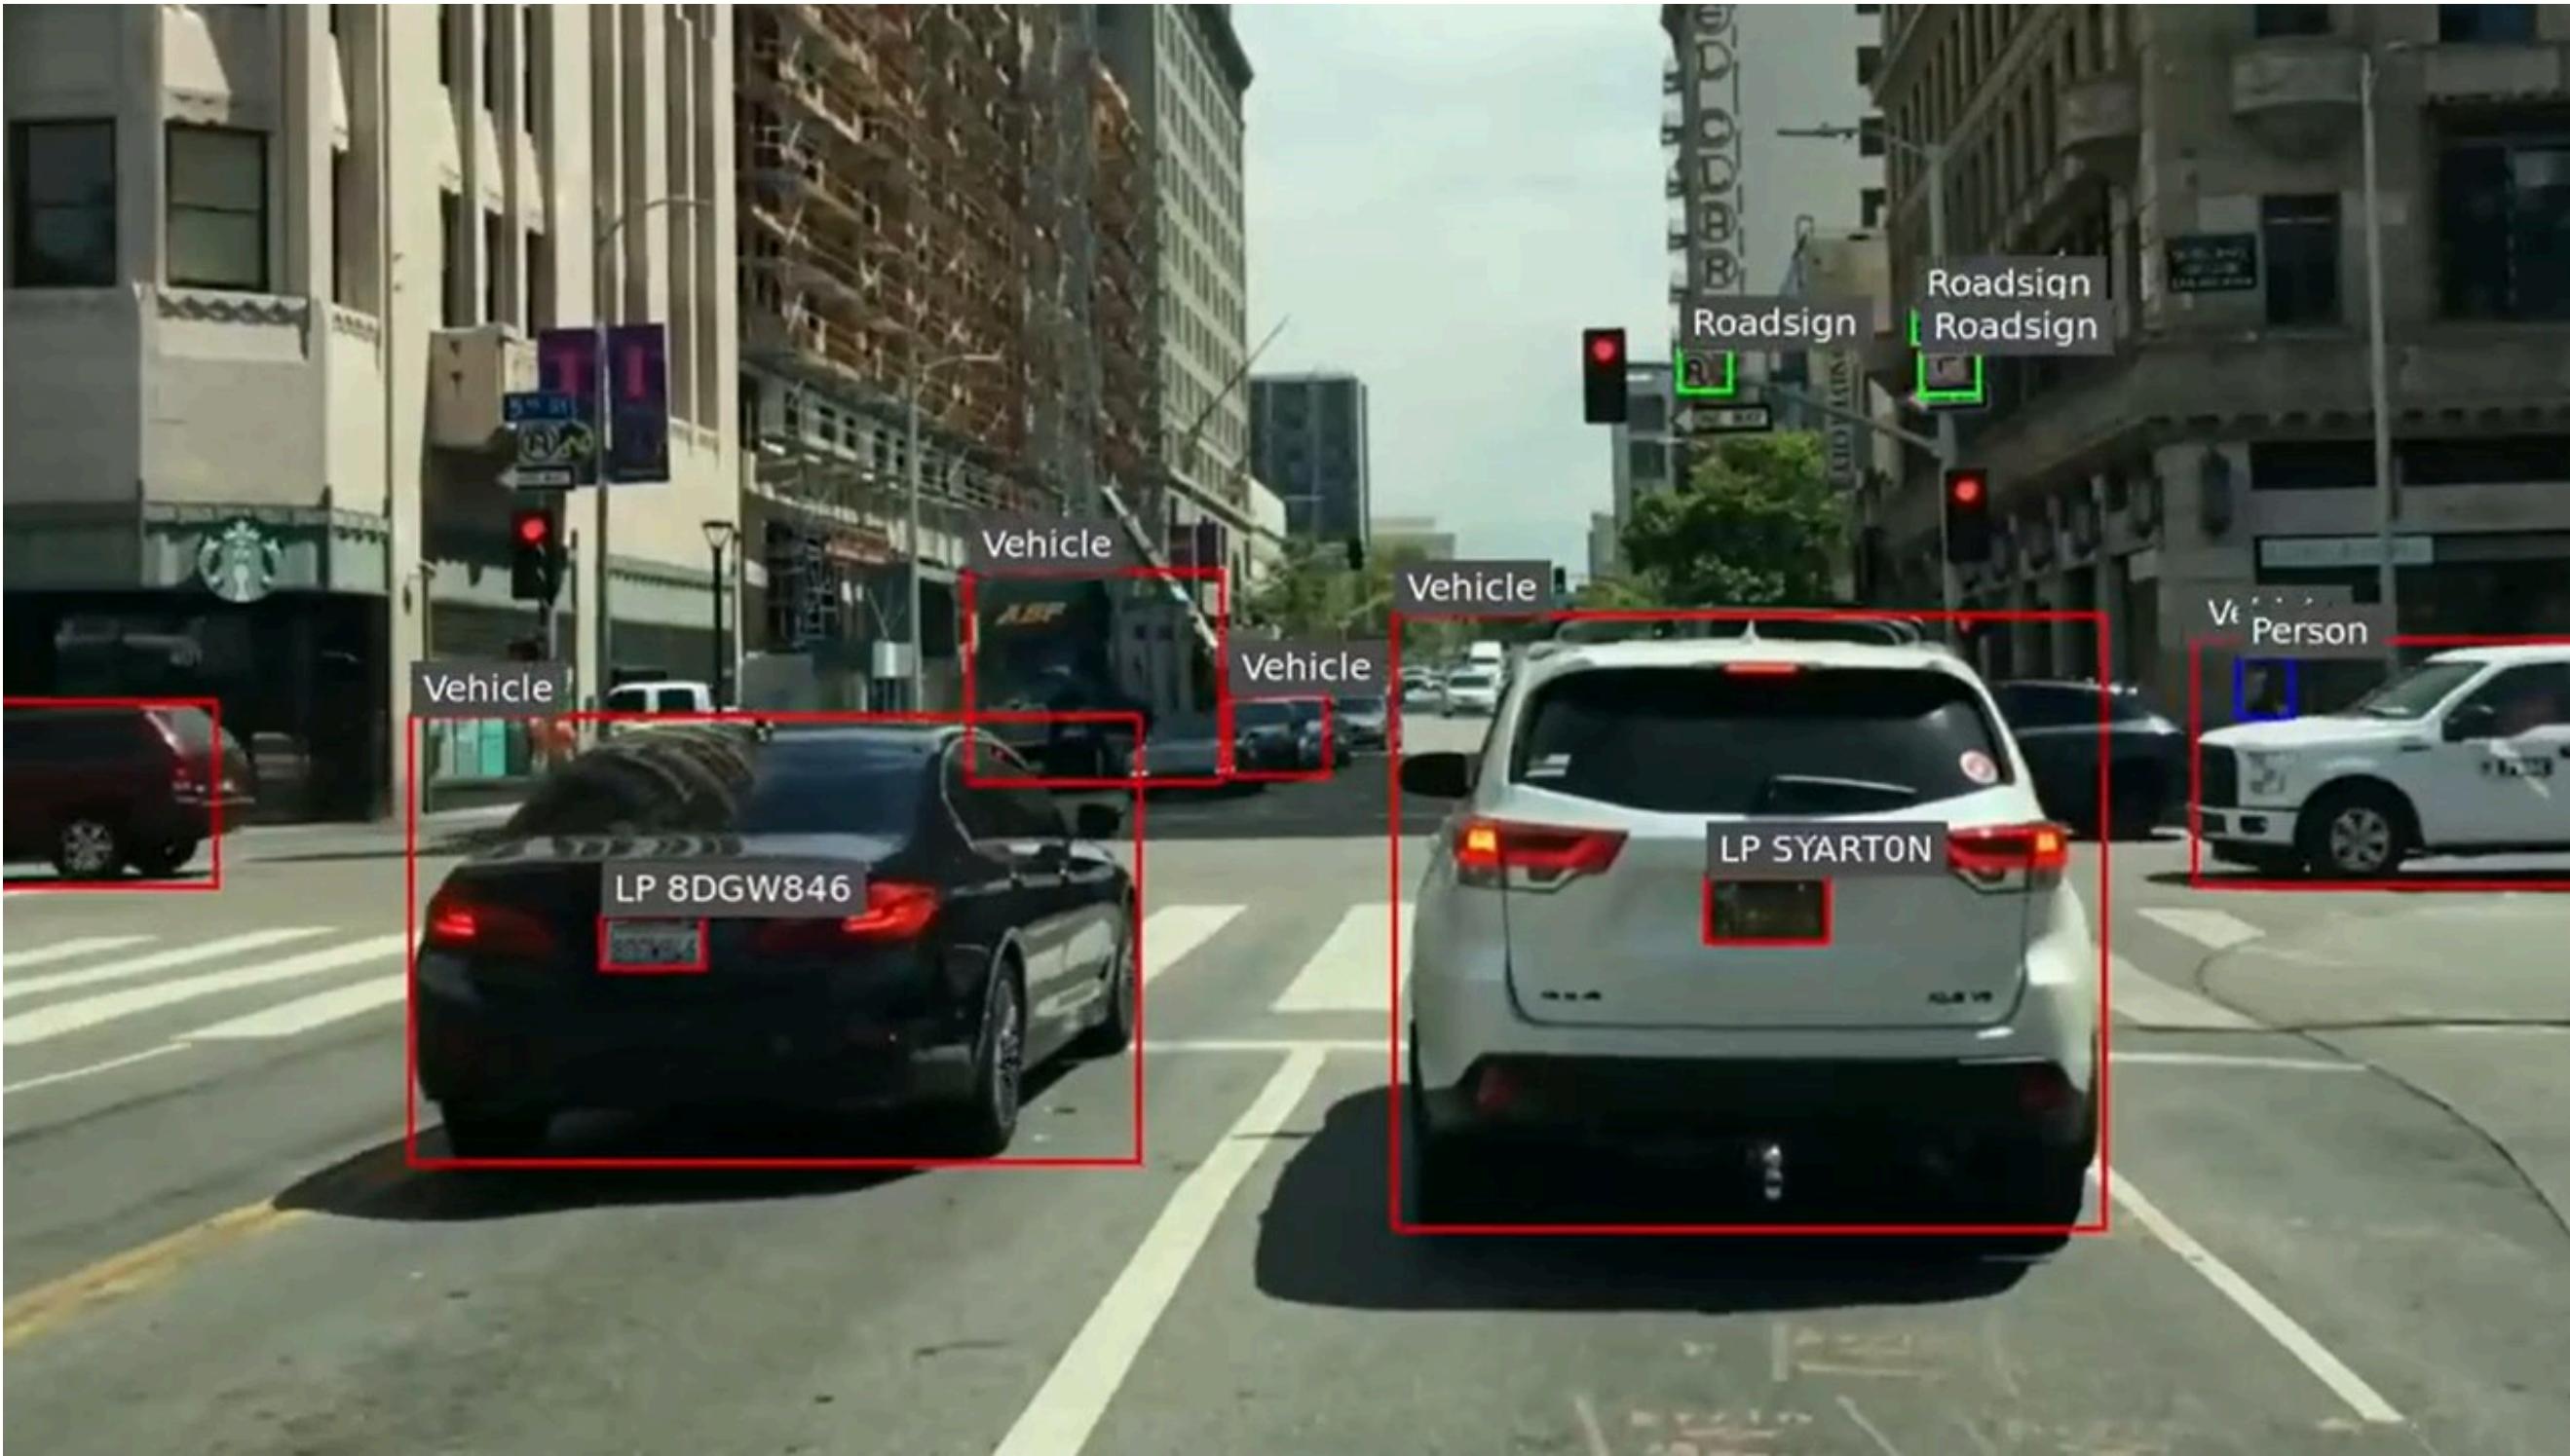
CAT, DOG, DUCK

Object detection is an important task in computer vision. In layman's terms, Object detection is defined as Object Localization + Object Classification.

# Object Detection :



Object Localization is the method of locating an object in the image using a bounding box and Object Classification is the method that tells what is in that bounding box.



There are numerous real-life applications for object detection. For example, in surveillance and monitoring, it is used in detecting trespassers, vehicle license plates... In medical imaging, it's used for detecting certain cells, cancer detection, tumor detection, etc. and the list goes on.

# Object Detection using Classical Computer Vision

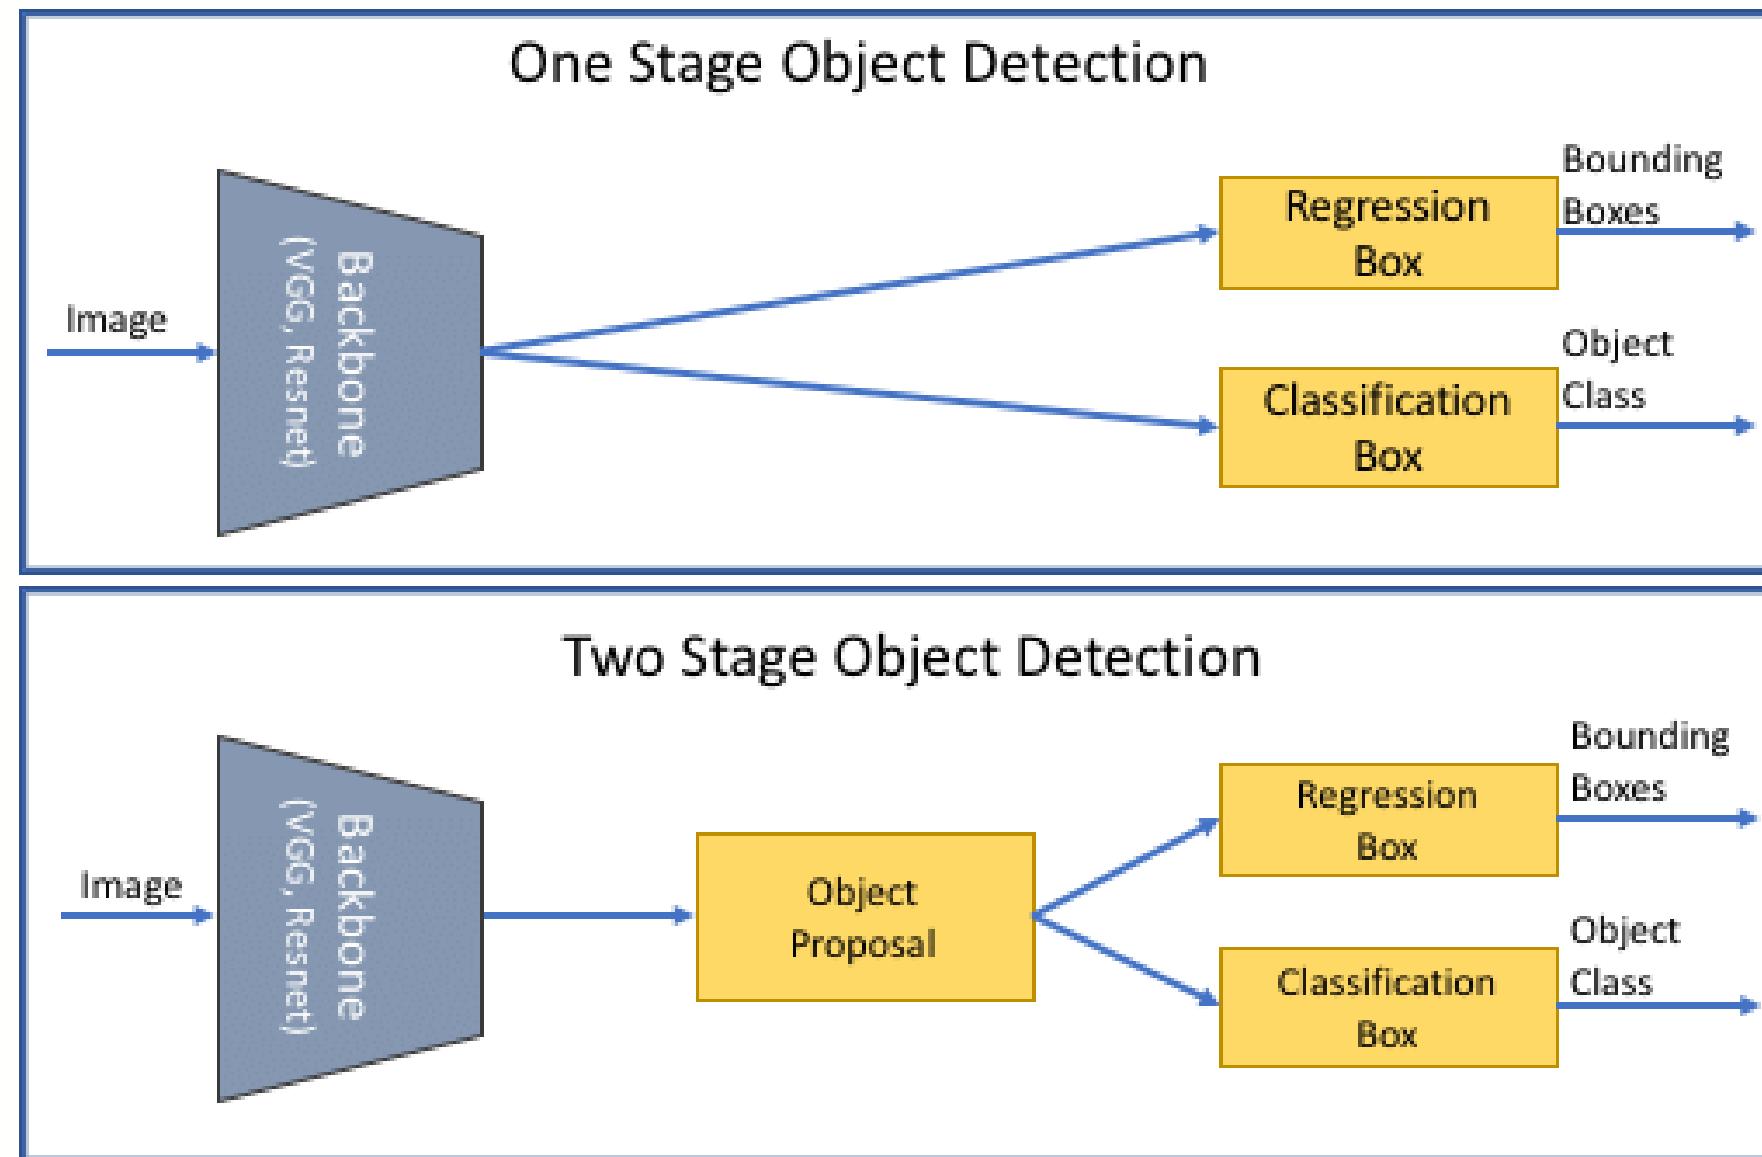
Earlier, detection algorithms used to be based on classical computer vision techniques such as template matching, Haar cascade, Feature detection and matching using SIFT or SURF, HOG Detector, Deformable Part-based Model (DPM), etc. However, most of these methods are very specific to the use case and don't have generalization capabilities, which redirected the research toward deep learning-based methods.

# Object Detection using Deep Learning

The Overfeat paper was a pioneer in deep learning-based object detection as it was a single network model employed to perform object classification and localization.

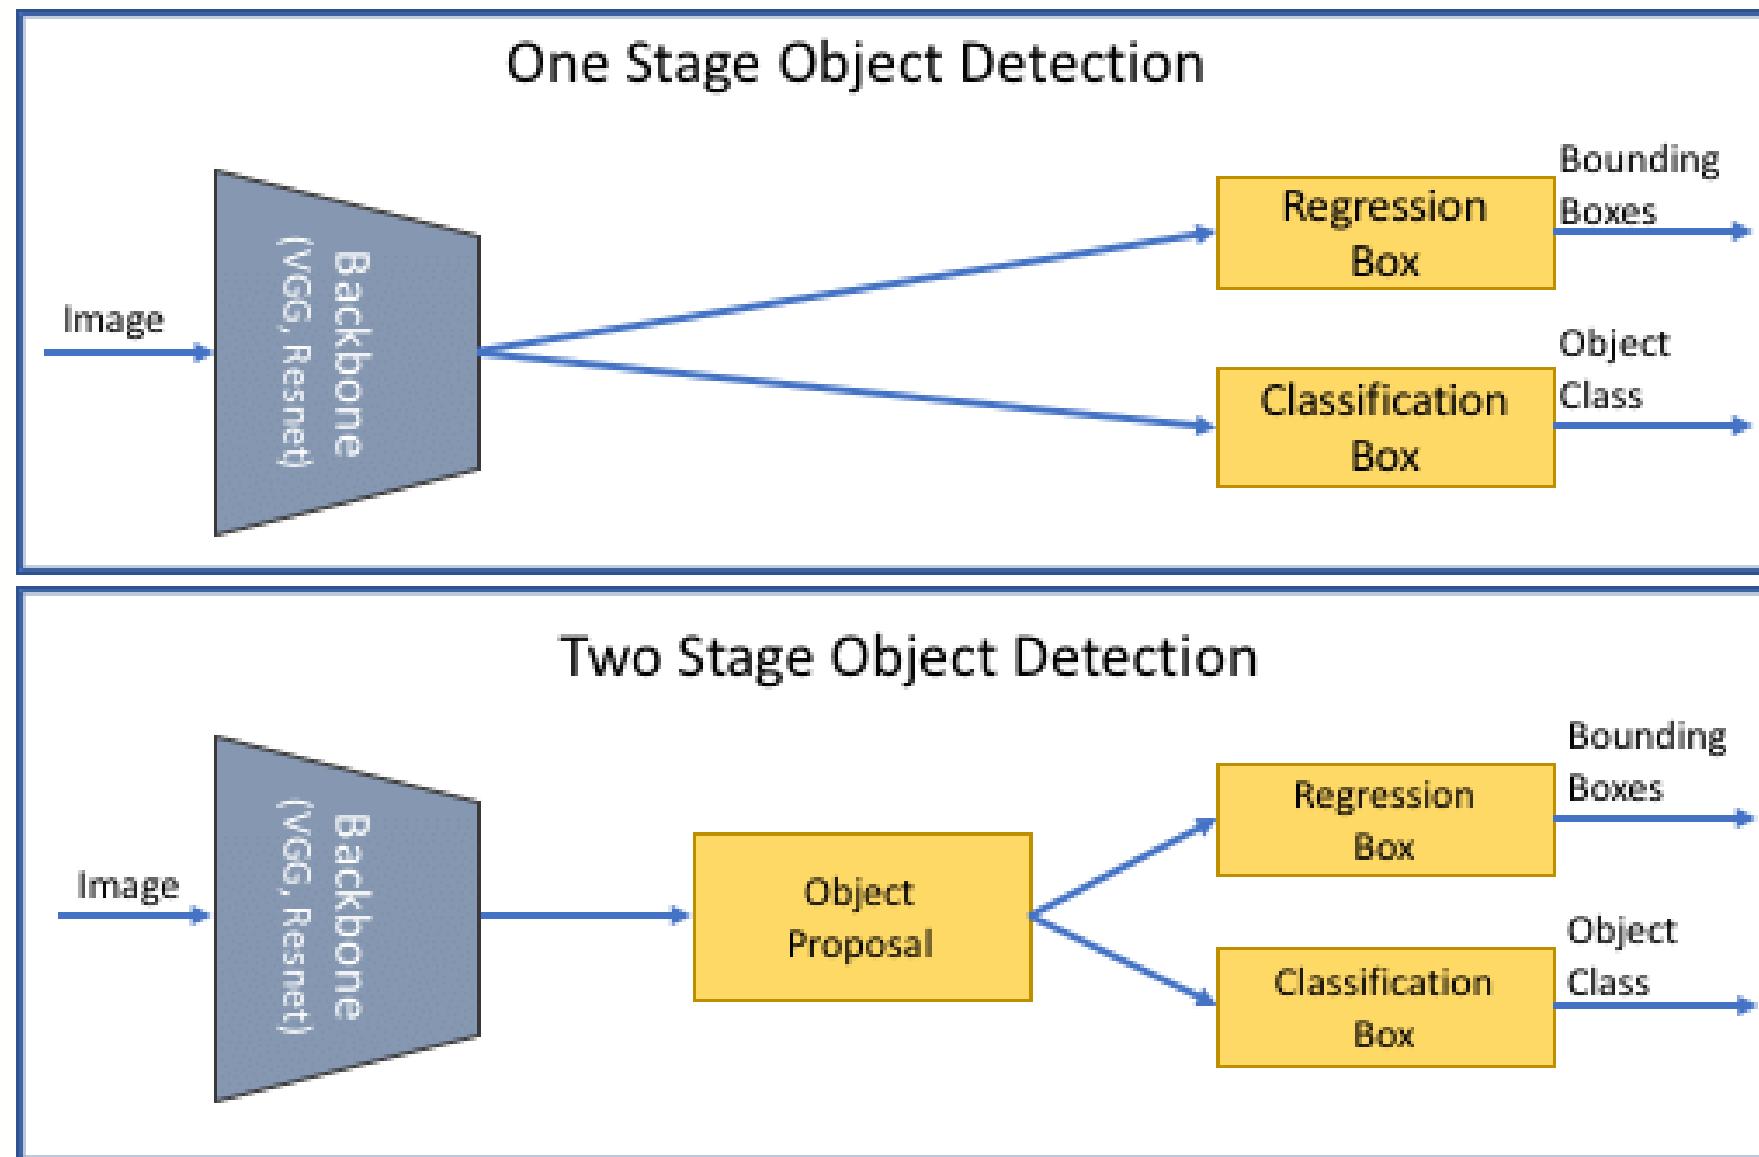
The Overfeat architecture closely resembles the **AlexNet** architecture. It does image classification on different scales in a sliding window fashion and carries out bounding box regression on the same CNN layer. Later other models like RCNN, FastRCNN, SPPNet, YOLO etc. emerged in this field.

# One-Stage Vs. Two-Stage Detectors



Initial models were two-stage detectors, which means there was a step for region proposals and finding out the region of interest. It was also a step for predicting the bounding box coordinates and class using the proposed region.

# One-Stage Vs. Two-Stage Detectors



Later the YOLO authors proposed a single-stage strategy that overlooked the Region Proposal step and ran detection directly over the entire image. This helped speed up the training and inference, preserve global context and keep the accuracy on par.

## Object Detection Techniques

### Traditional object Detector

VJ detector

HOG

DPM

### Deep learning based object detectors

#### Two stage detector

RCNN

SPP

Fast RCNN

Faster RCNN

#### One stage detector

YOLO

SSD

RetinaNet

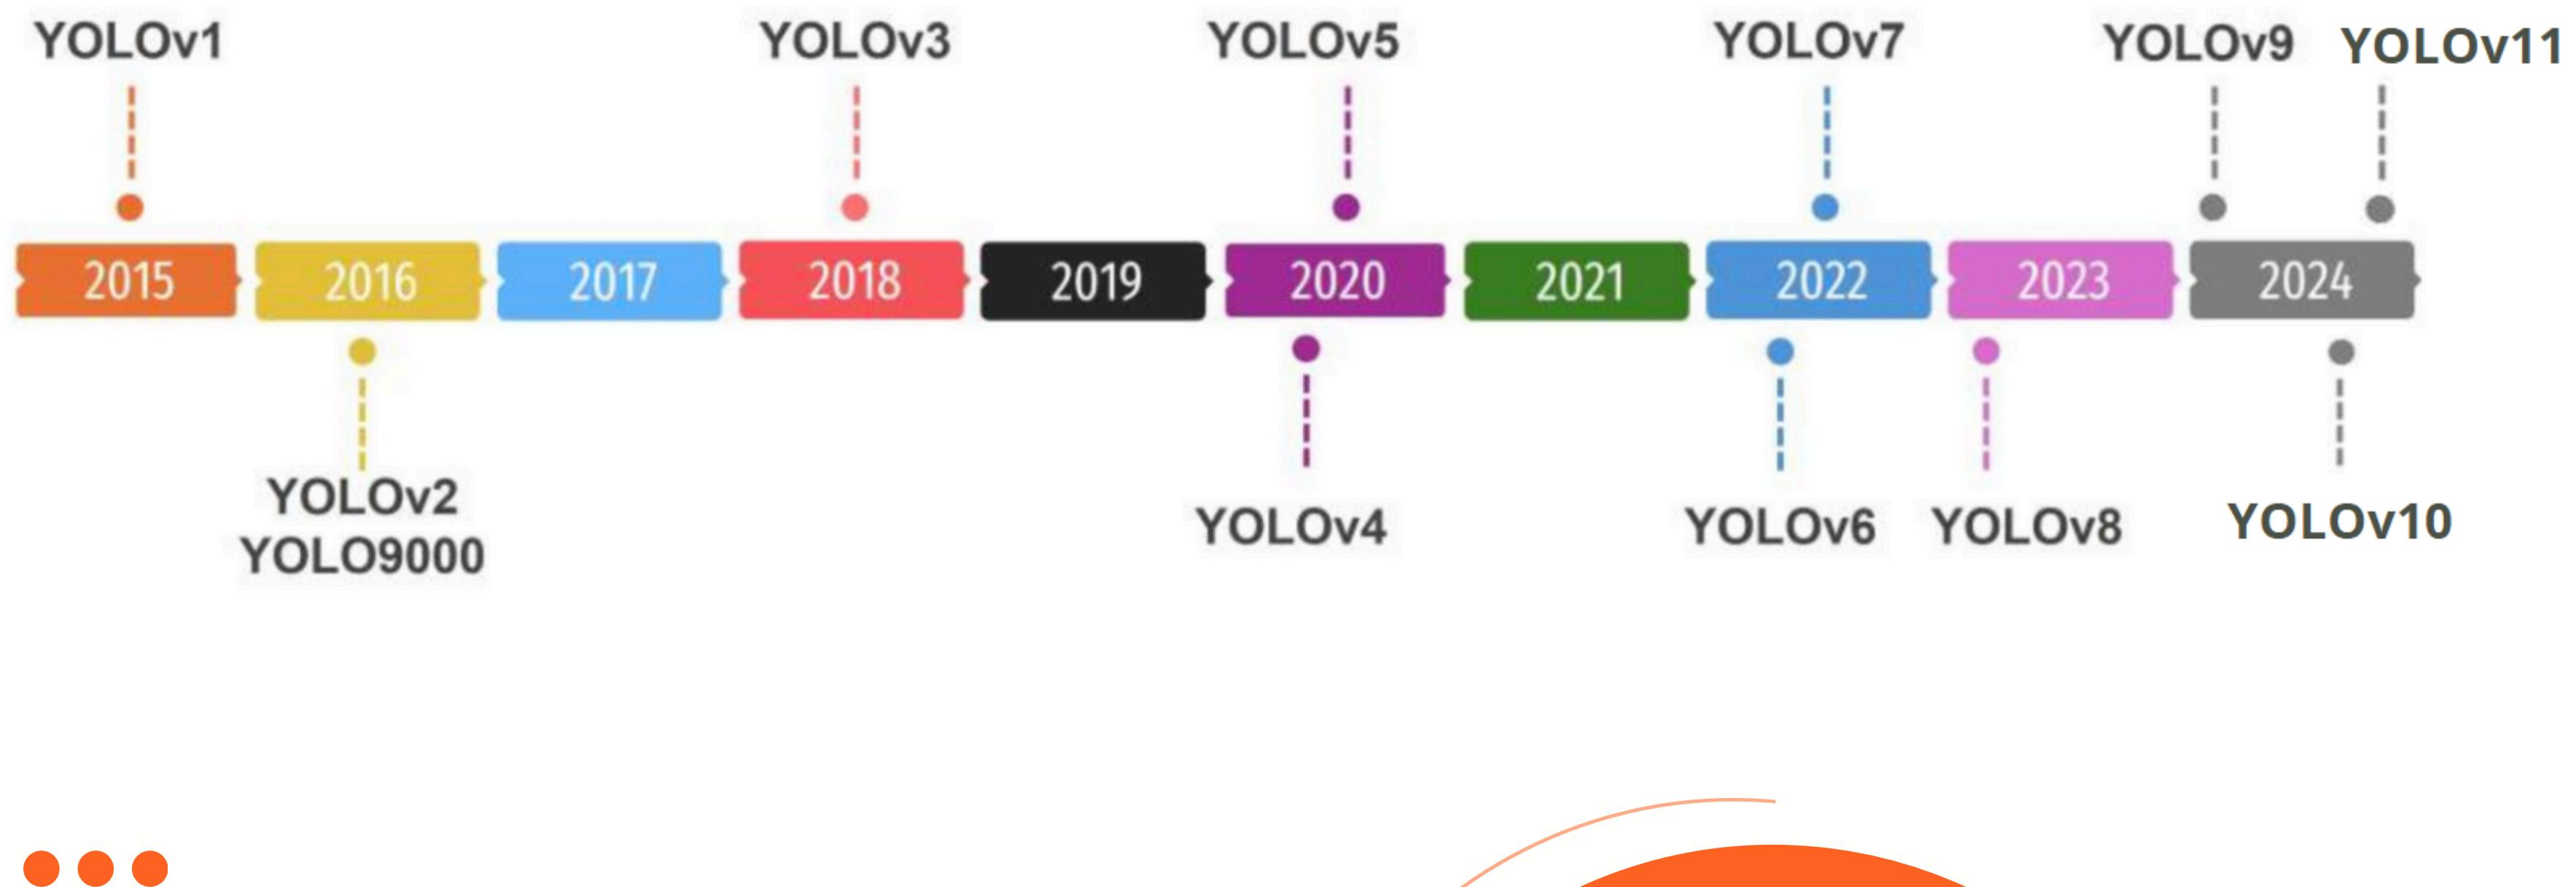
LADet

# One-Stage Vs Two Stage Detectors

Aspect	One-Stage Detector	Two Stage Detector
Process	Directly predicts bounding boxes & classes	Region proposals + classification
Examples	YOLO, SSD, RetinaNet	Faster R-CNN, Mask R-CNN
Speed	Faster, real-time capable	Slower, not ideal for real-time
Accuracy	Slightly lower, especially on small objects	Higher accuracy, especially for small or dense objects
Use Case	Real-time detection, applications needing speed	Applications requiring accuracy over speed

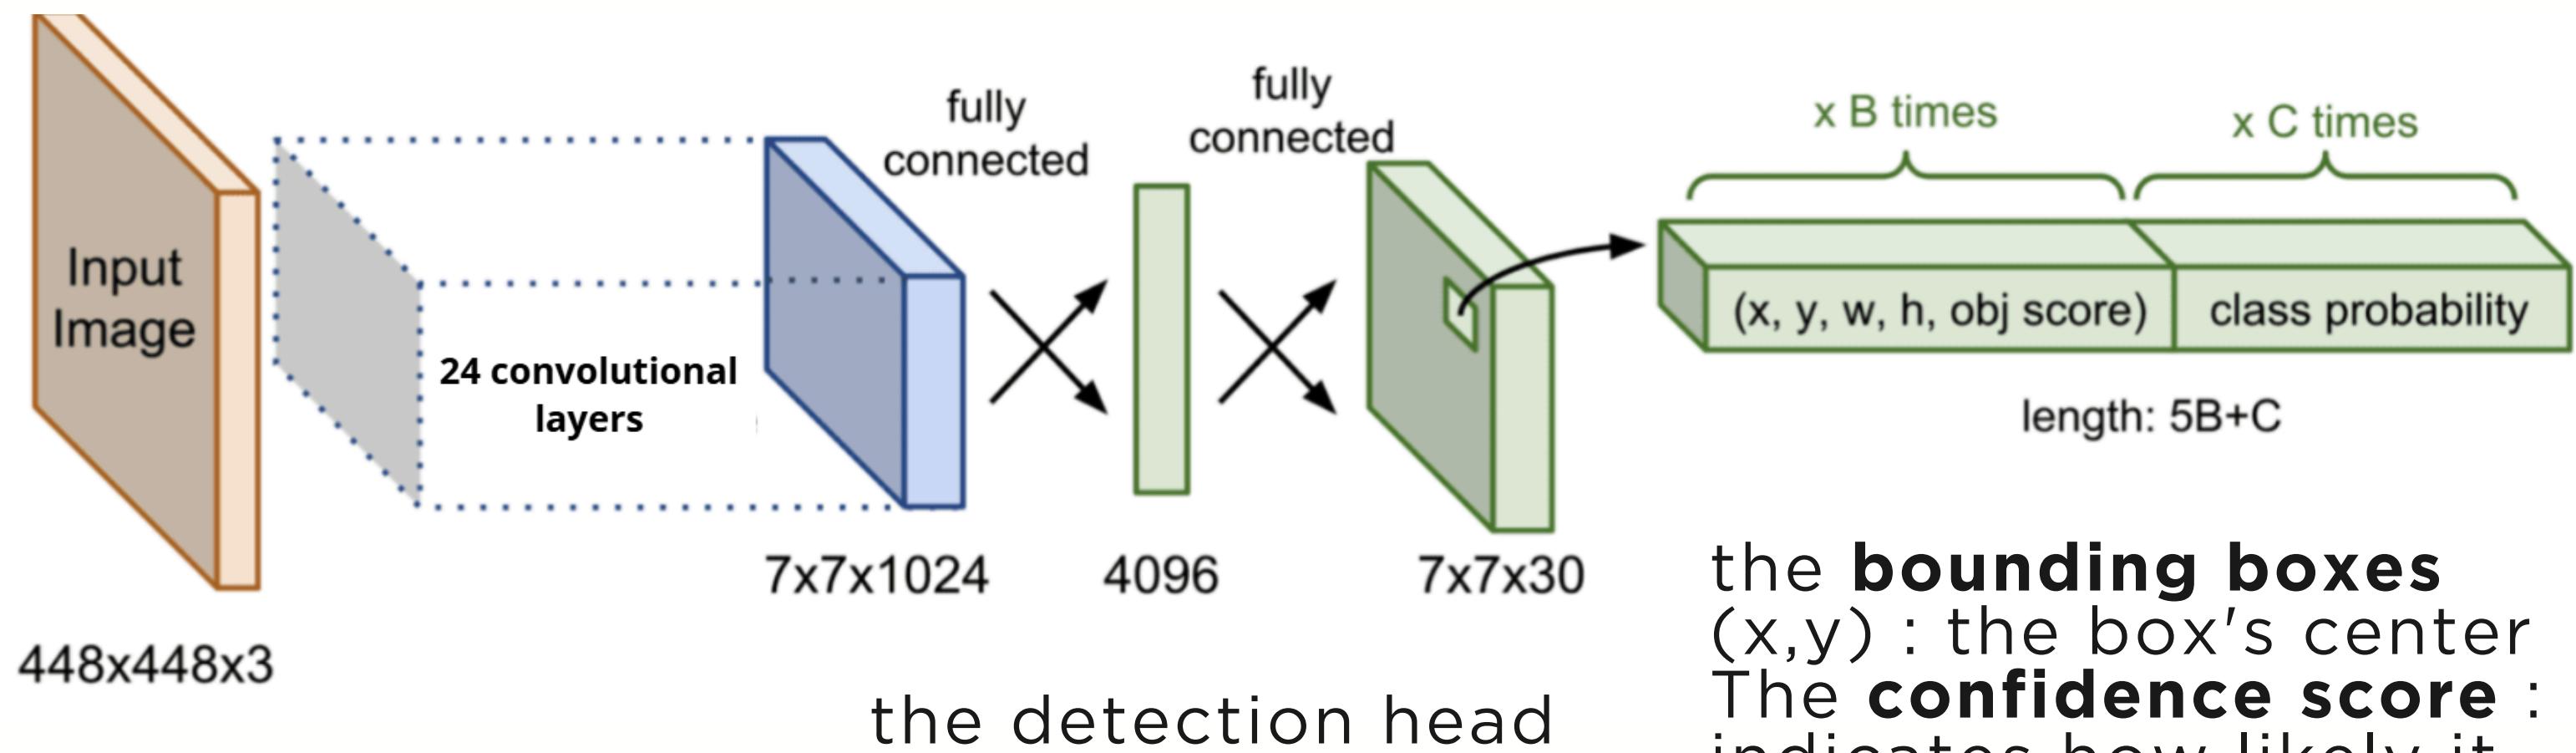


# Timeline of YOLO Versions



- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

# YOLOv1

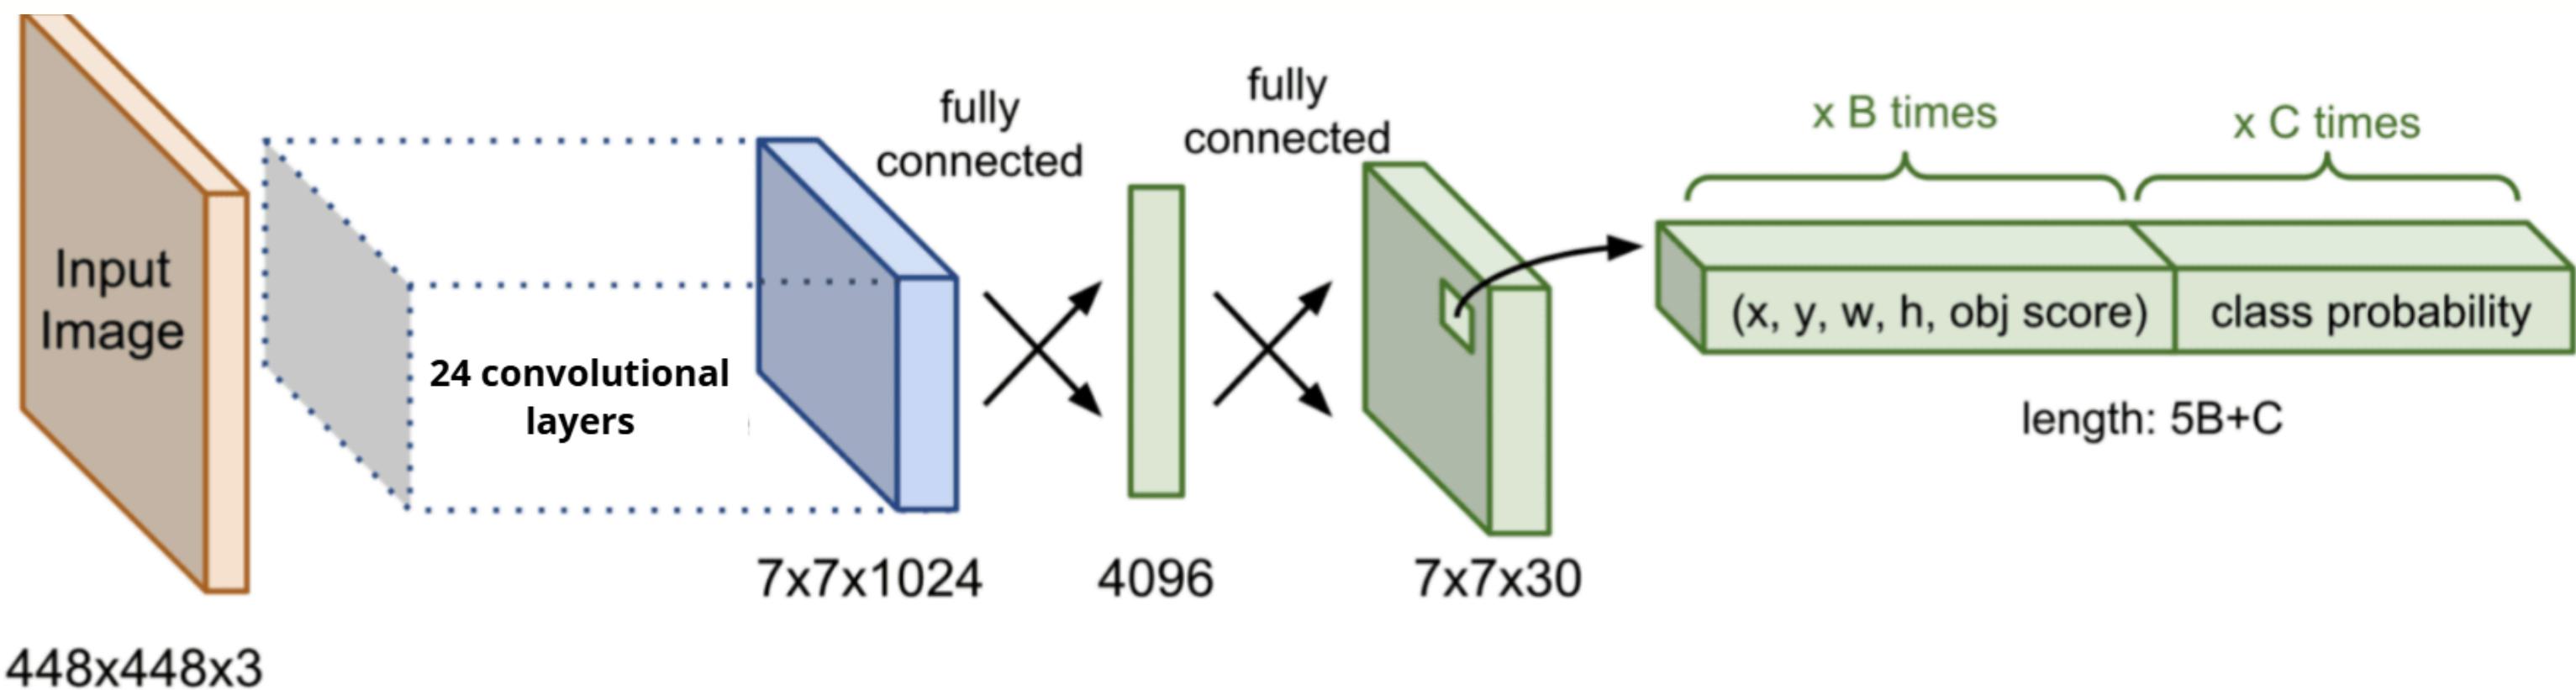


**the backbone** : the core  
feature extraction network,  
inspired by GoogleNet's  
architecture.

the **bounding boxes**  
 $(x,y)$  : the box's center  
The **confidence score** :  
indicates how likely it  
contains an object, as well  
as its accuracy  
the **Class Probability**.

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

# YOLOv1



To avoid **overfitting**:

- dropout rate of 0.5
- extensive augmentation.

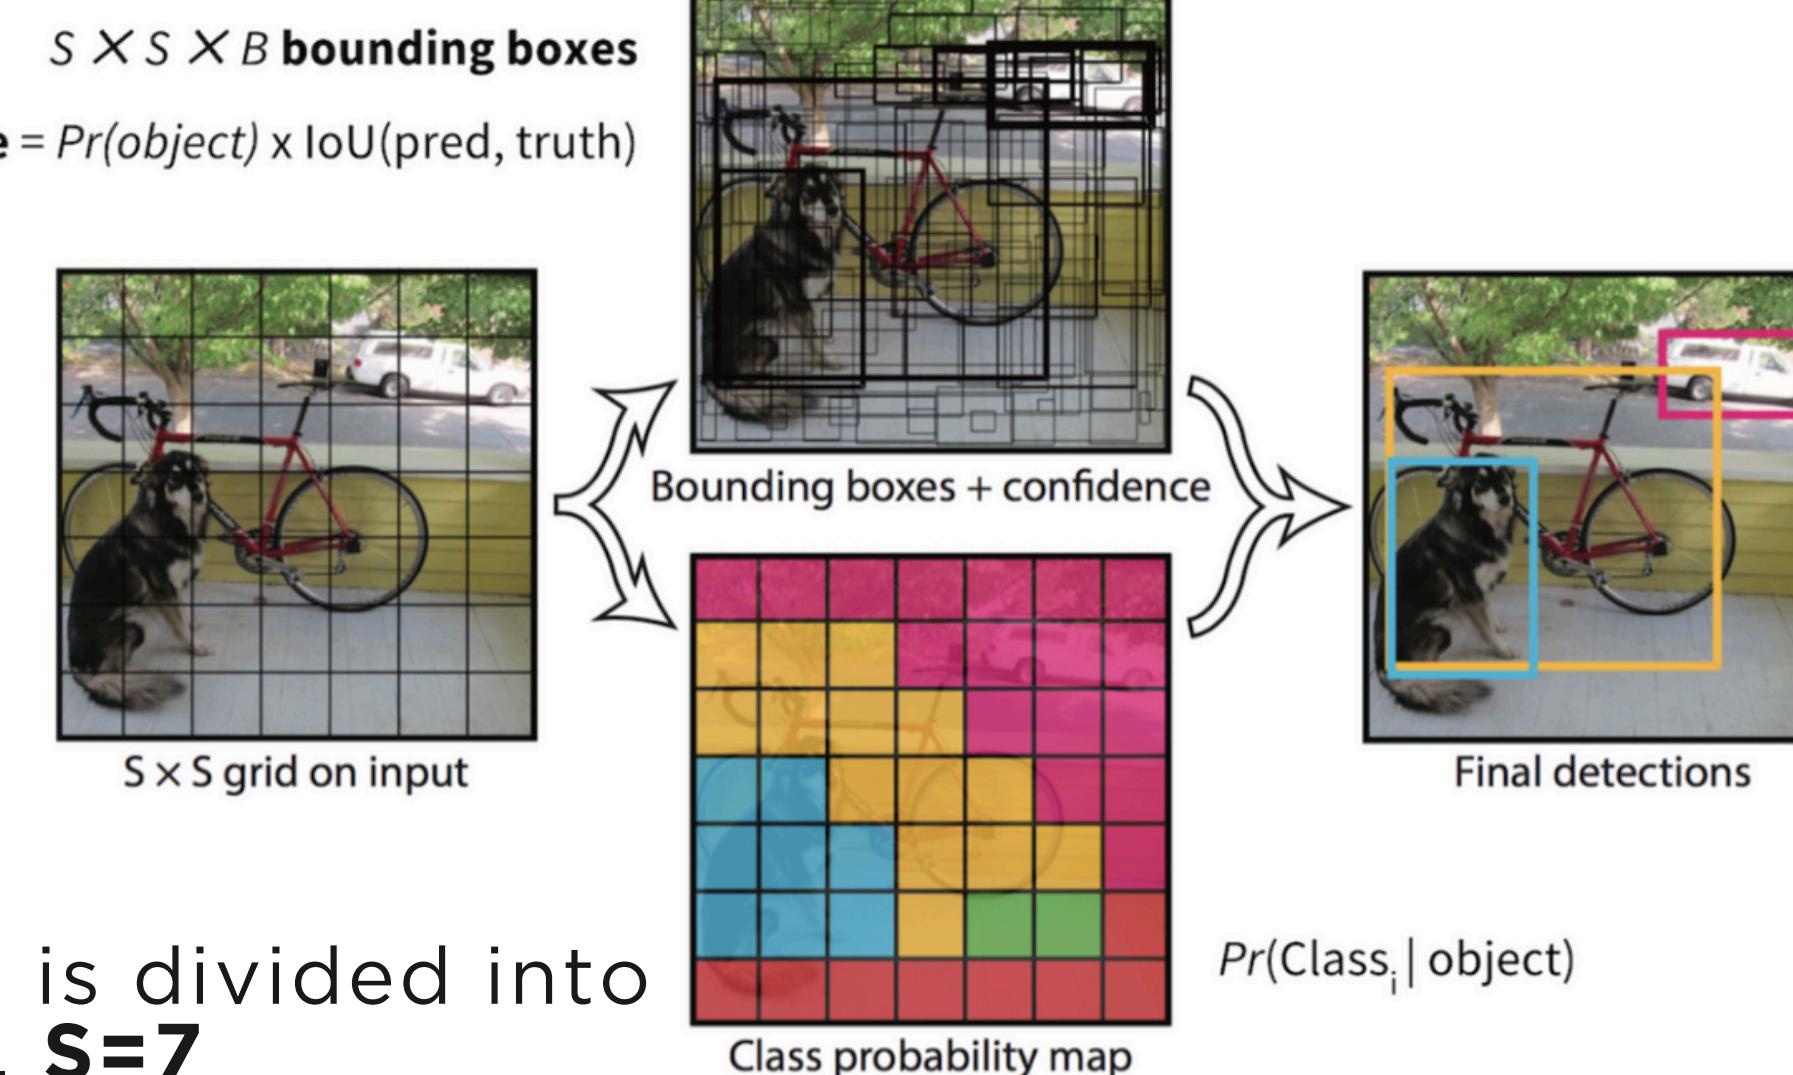
three parts to the **loss function**:

- localization loss,
- confidence loss,
- classification loss.

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

# YOLOv1

Each grid cell predicts 2 bounding boxes and confidence scores. **B=2**



The image is divided into a  $7 \times 7$  grid. **S=7**

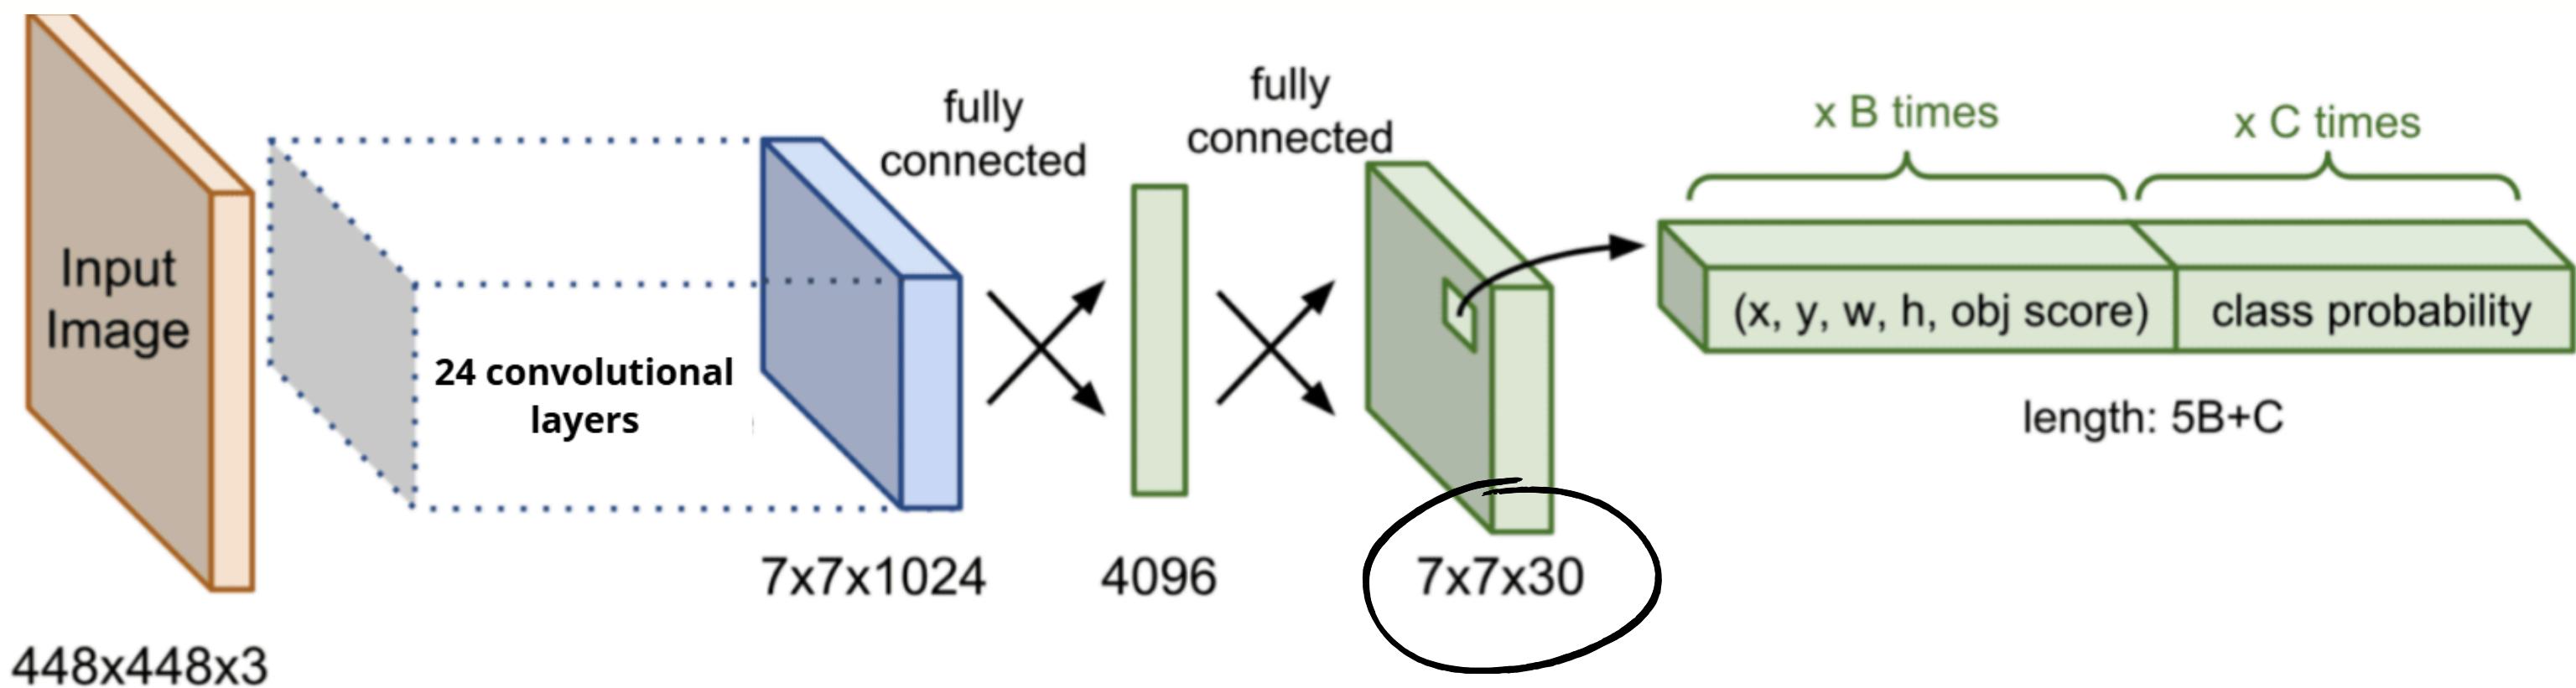
Each grid cell predicts the probability of each class. **C=20**

**Total Predictions:**  
2 bounding boxes  
 $\times 5$  values ( $x, y, w, h, c$ )  
+ 20 class scores = 30 values



- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

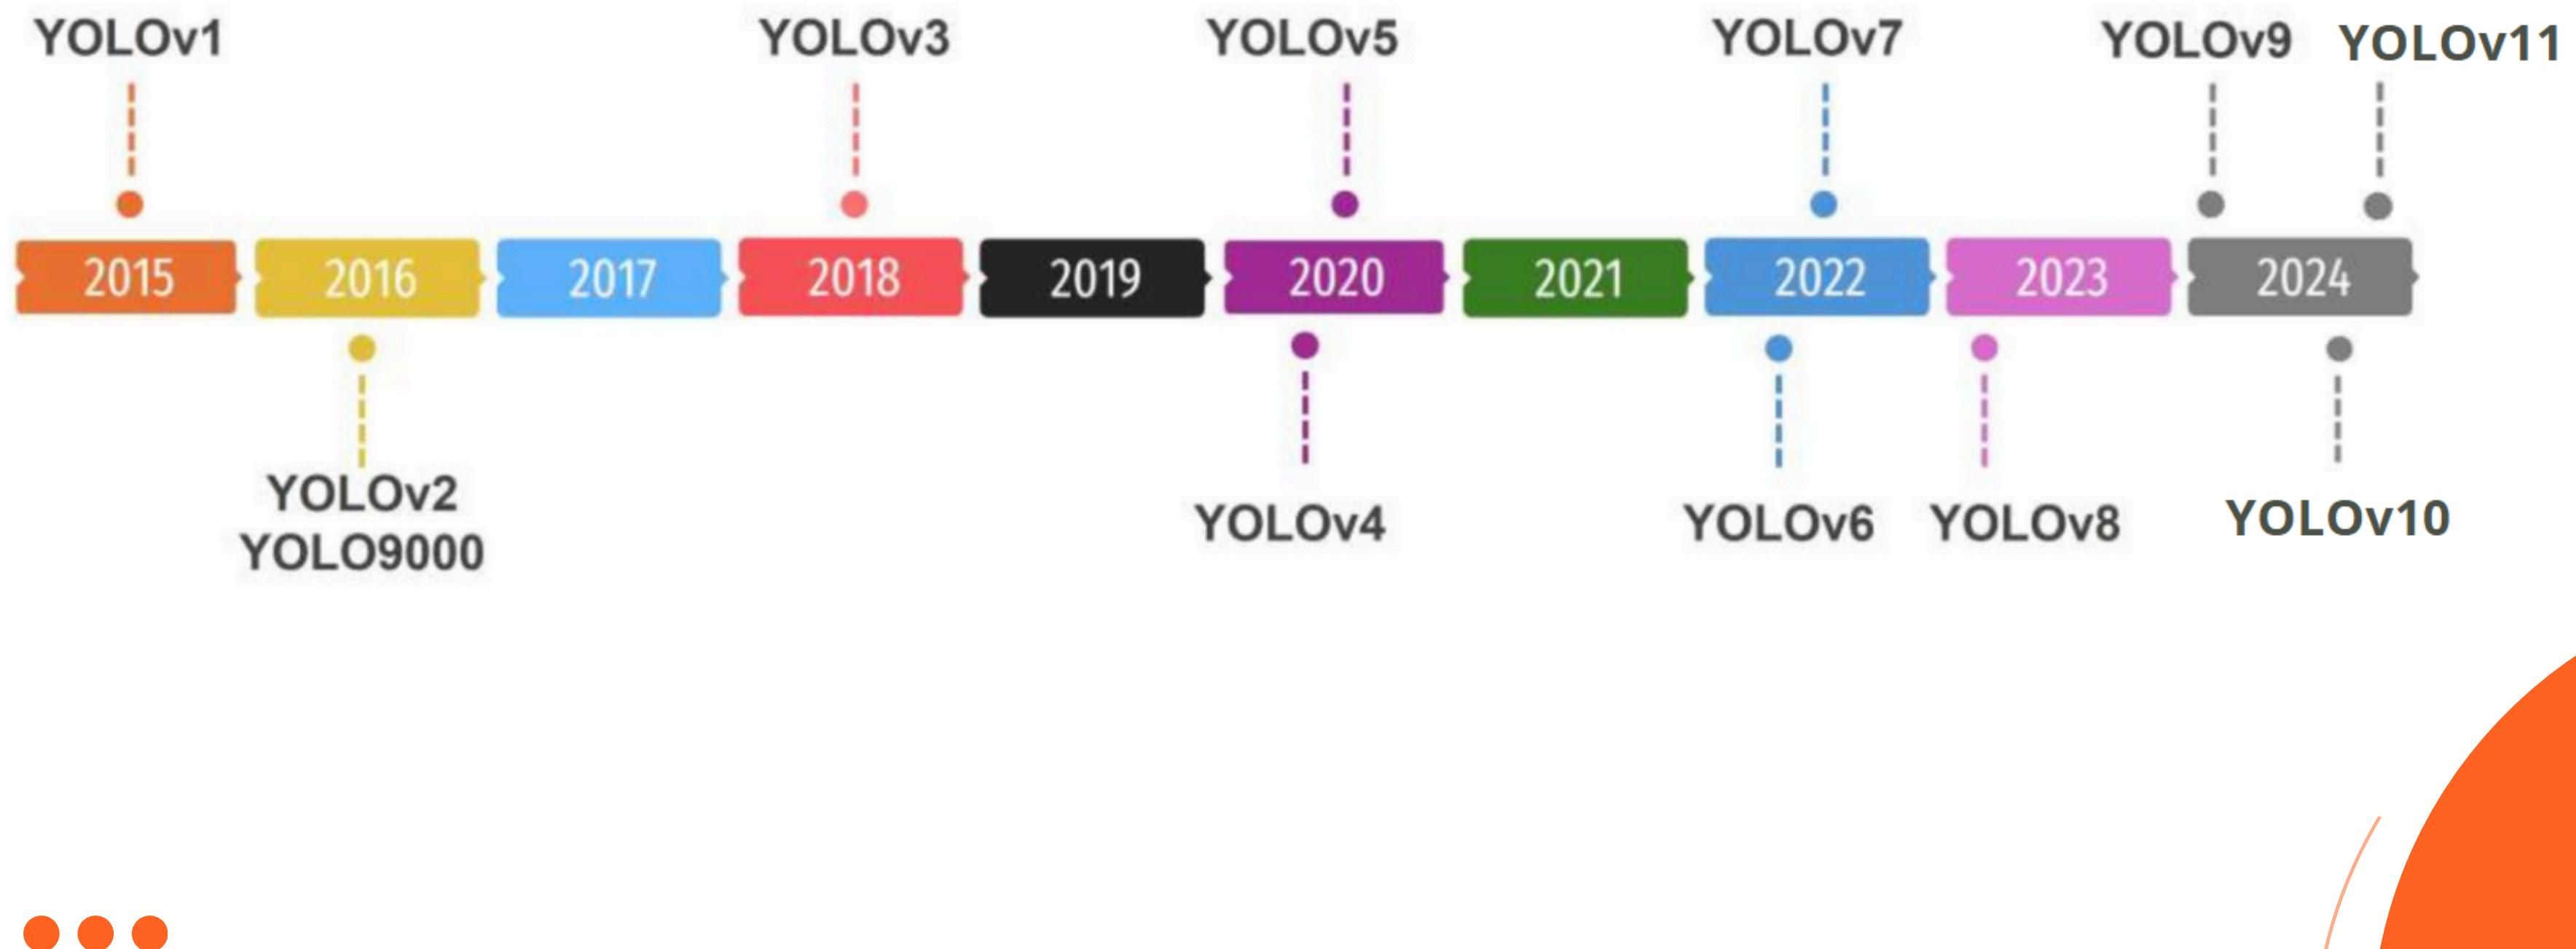
# YOLOv1



- Speed
- Simplicity

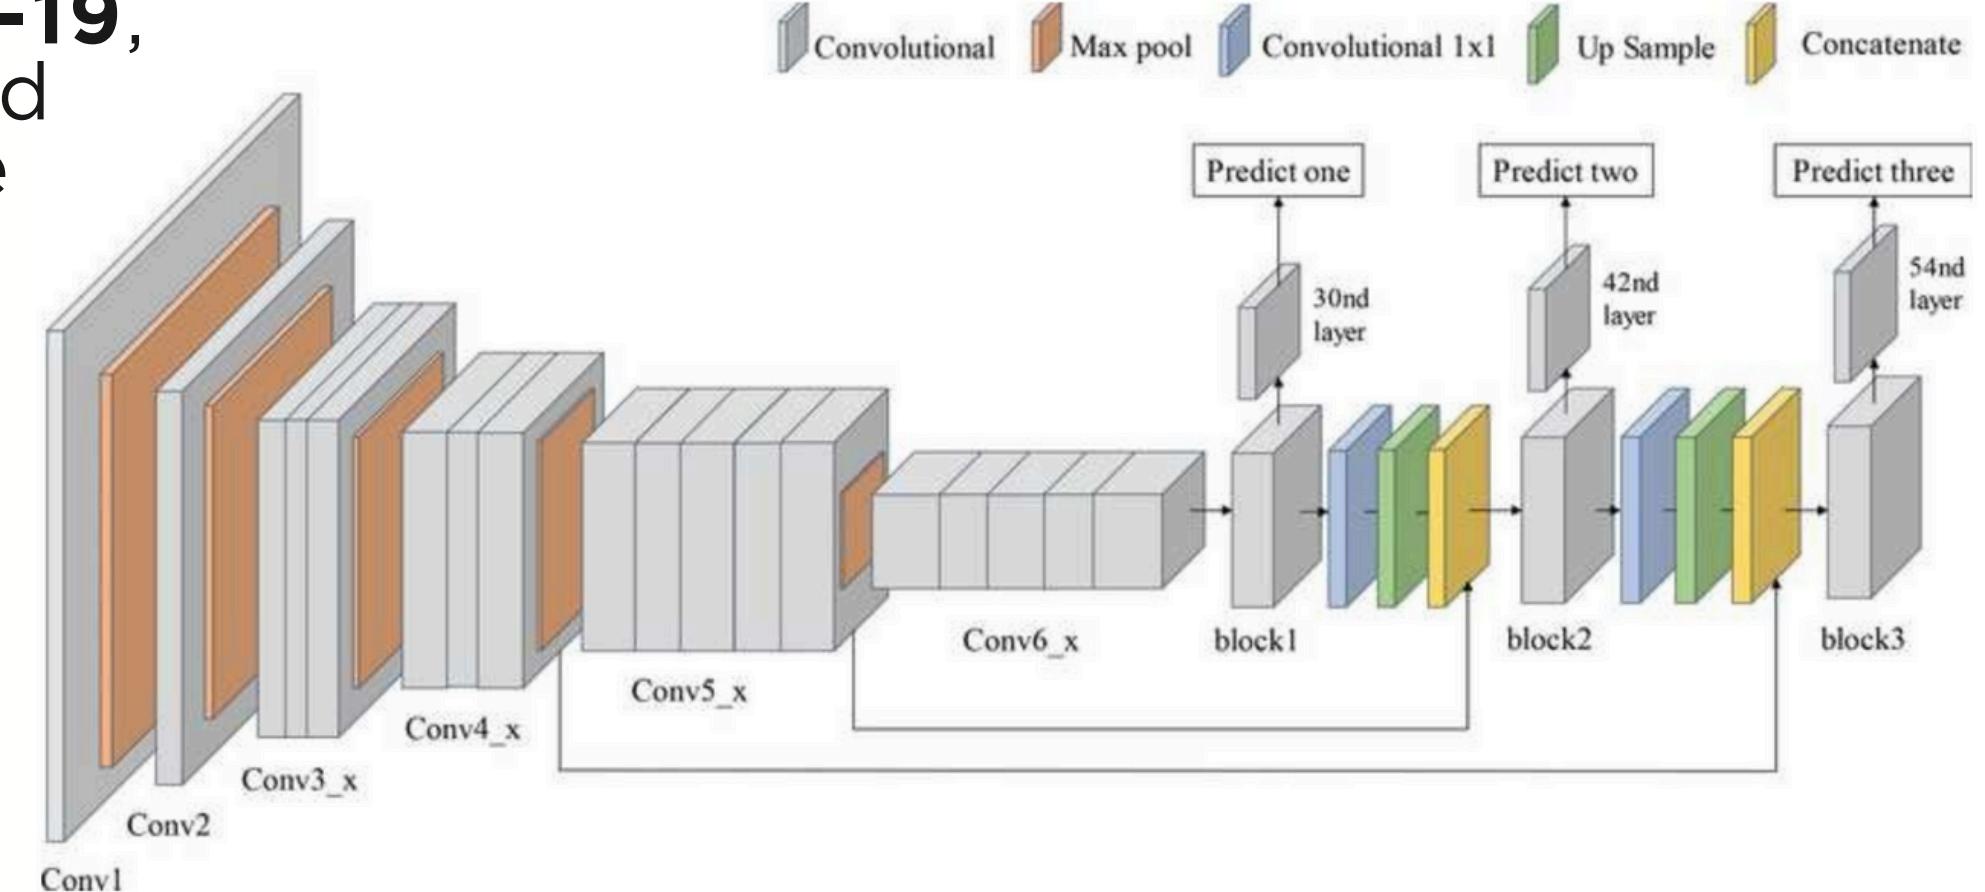
- Low Localization Accuracy: only 2 bounding boxes.
- Rough Predictions: The  $7 \times 7$  grid size.
- Low Recall: particularly for smaller objects.

# Timeline of YOLO Versions



# YOLOv2

- Introduces **Darknet-19**,  
a more powerful and  
deeper architecture



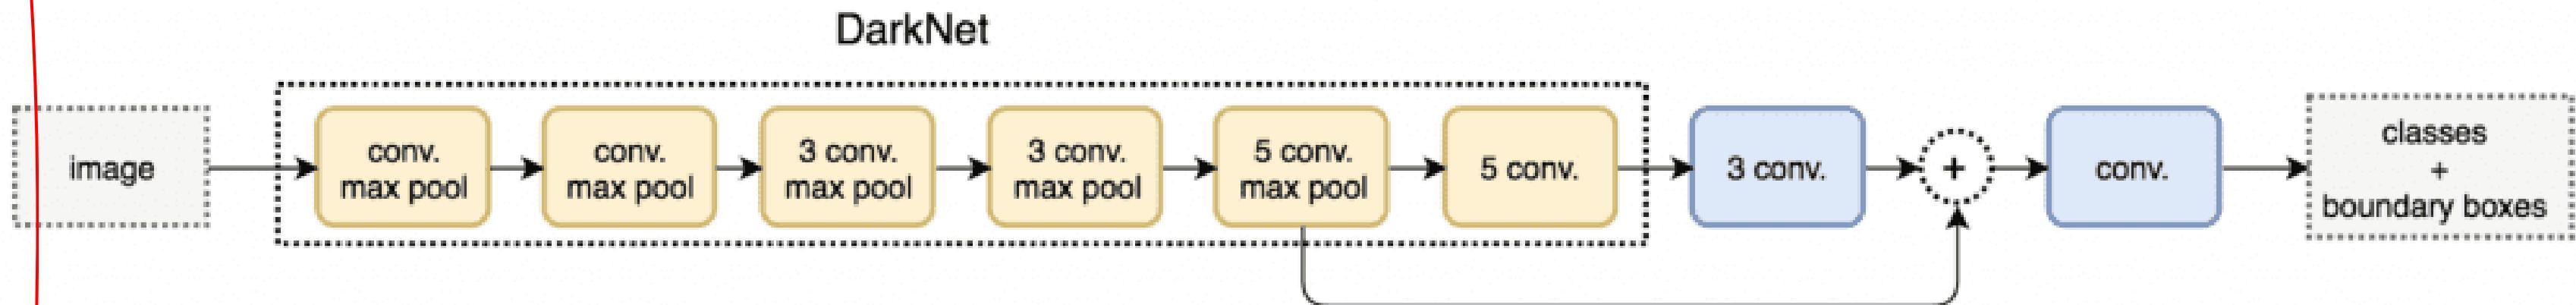
**passthrough layer** that combines low-level, high-resolution features  
from earlier layers with high-level, low-resolution features.

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11



# YOLOv2

- Introduces **Darknet-19**, with max-pooling layers.

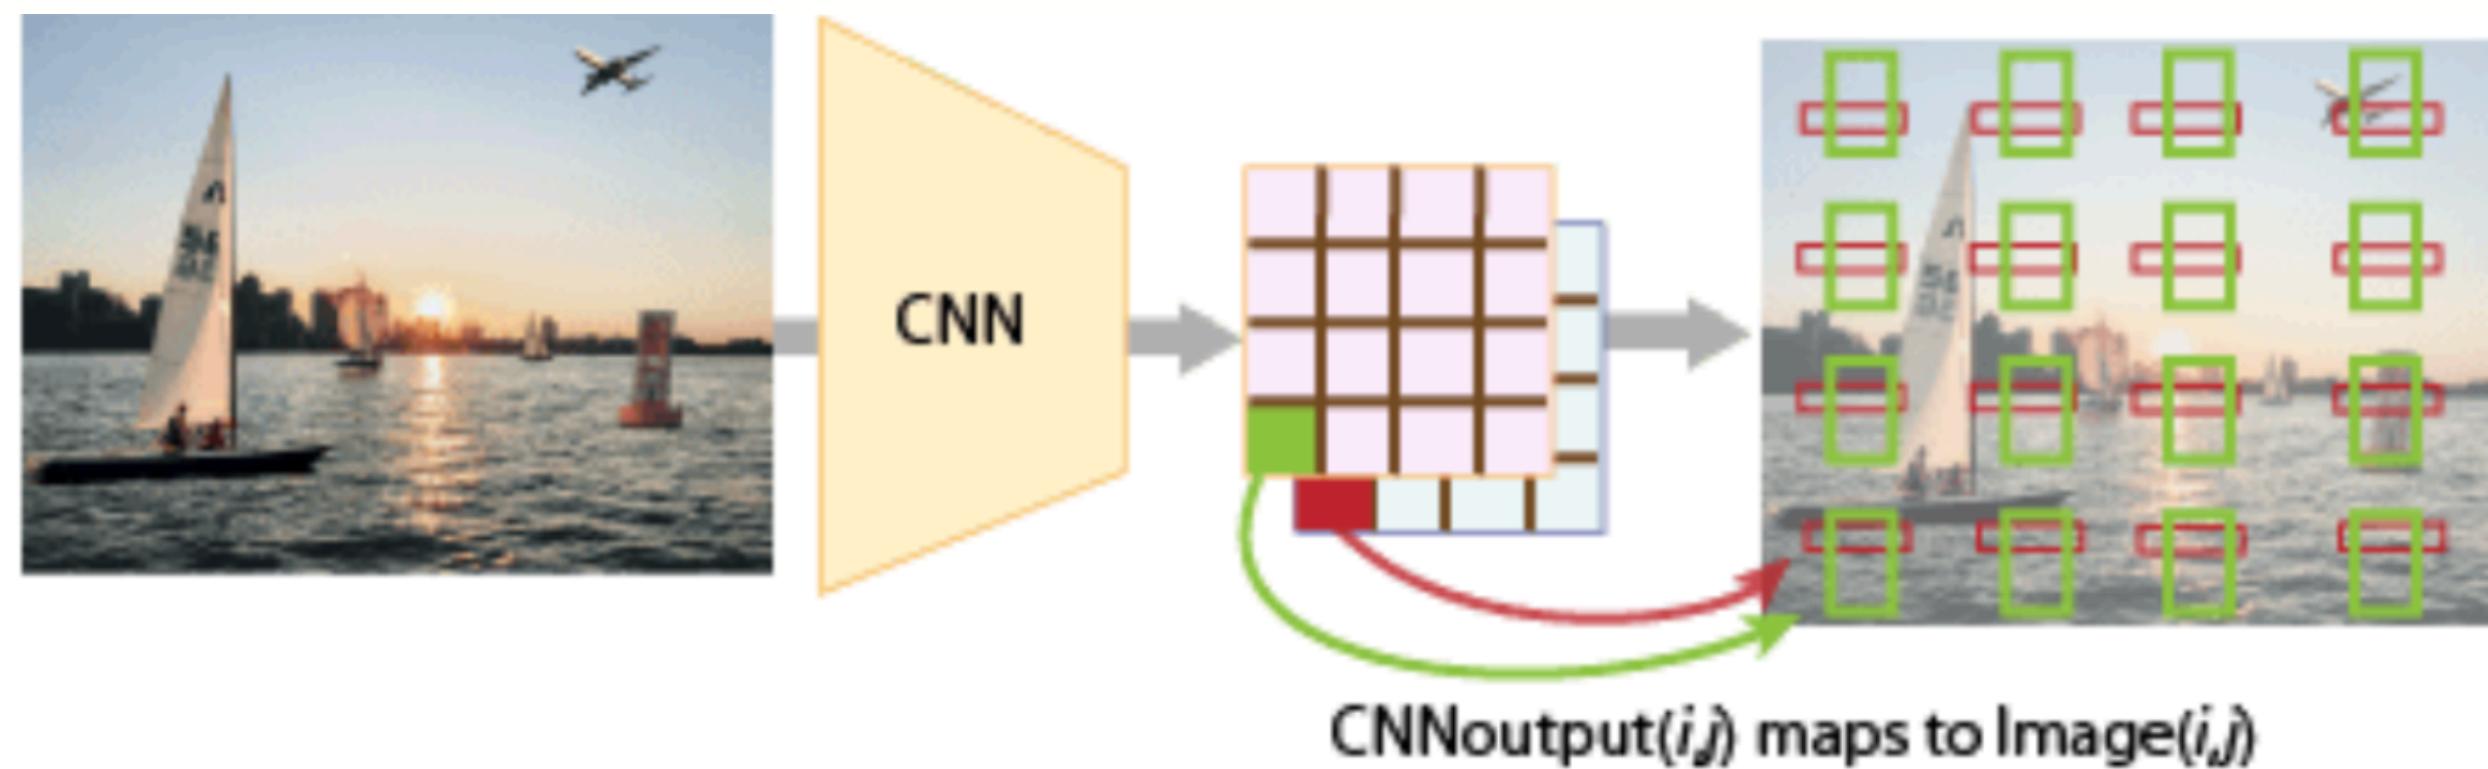


- allowing different image resolutions : ranging from 320x320 to 608x608 pixels
- Implements **batch normalization** across all convolutional layers

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

# YOLOv2

- Introduction to **Anchor Boxes**: predefined bounding boxes.
- “starting points” for predicting object locations :

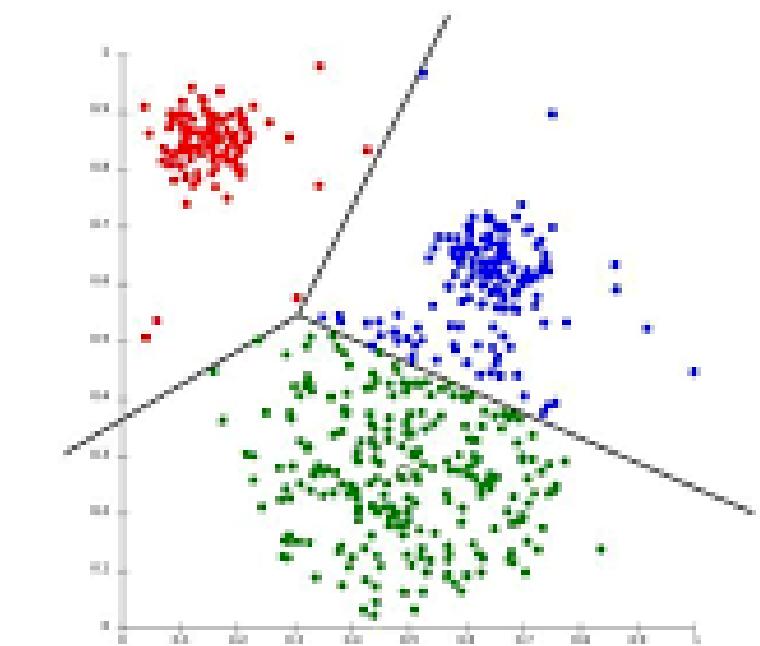


Instead of predicting the bounding boxes directly from scratch, **the model predicts adjustments** to these anchor boxes to better fit objects in the image.

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

# YOLOv2

- the authors proposed using **k-means clustering** on all the bounding boxes of the dataset to Set Anchor Box Dimensions.



Anchor boxes reduce the complexity of the model by giving it reference bounding boxes to work from rather than predicting each box completely from scratch.



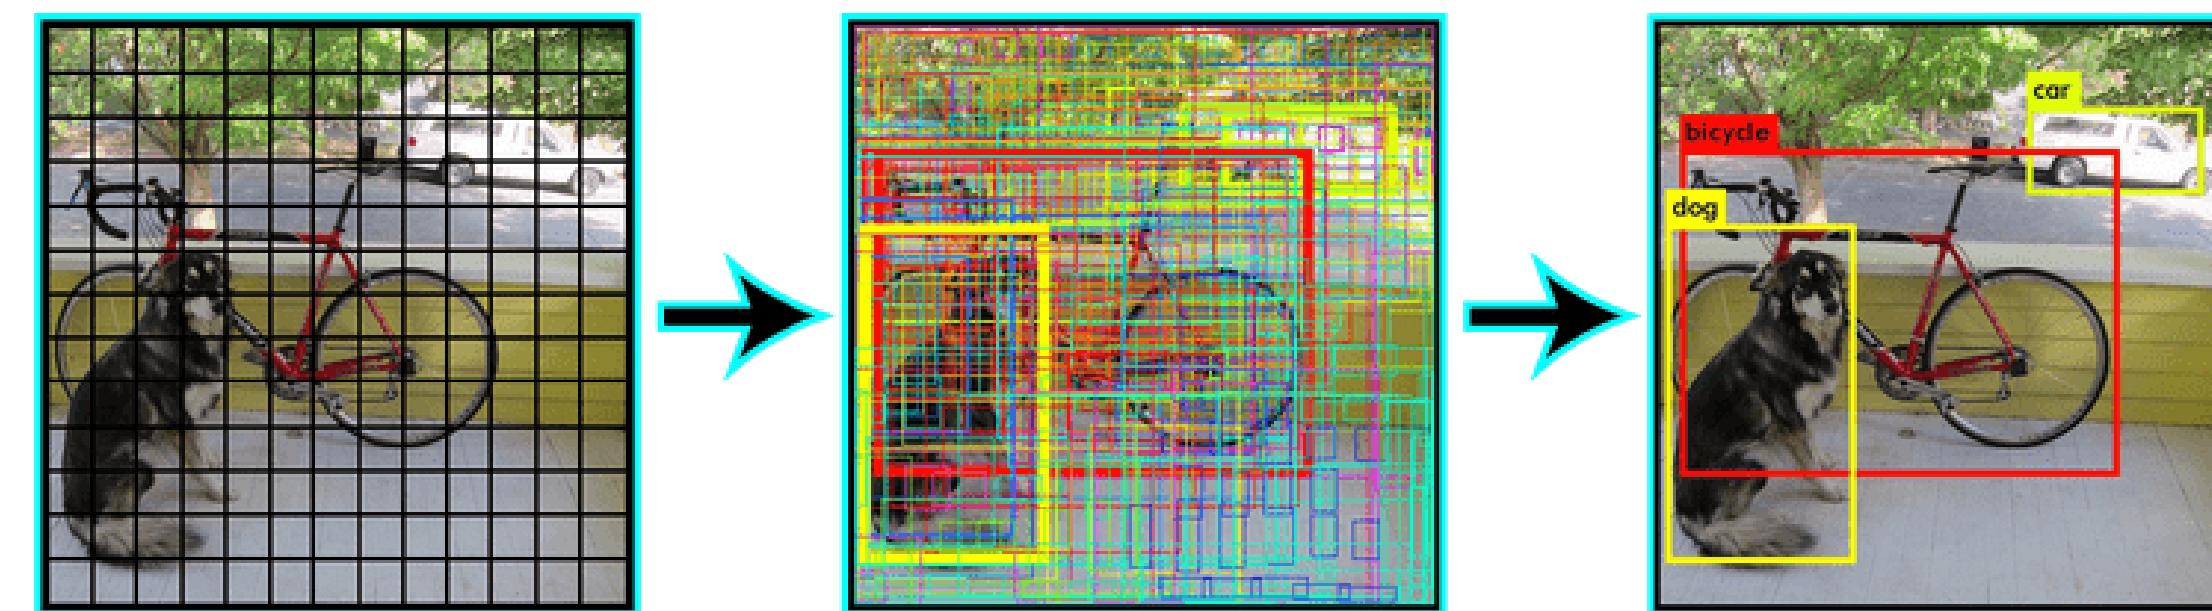
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11



# YOLOv2

Each anchor box predicts (5 values):

- Adjustments to its position and size to match the detected object.
- A confidence score (probability that an object is within the box).
- Class probabilities for each object type.



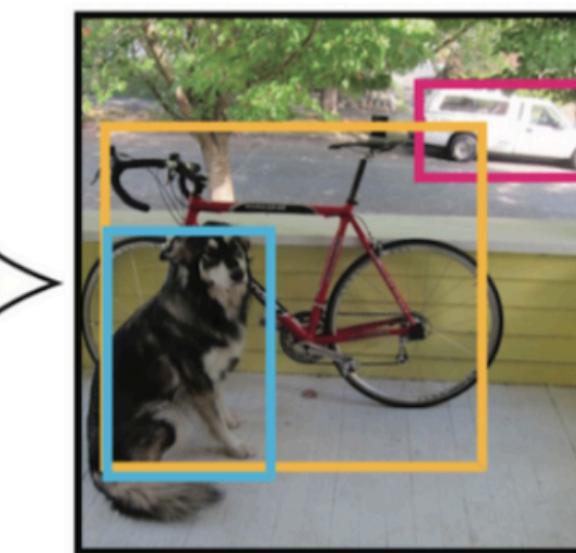
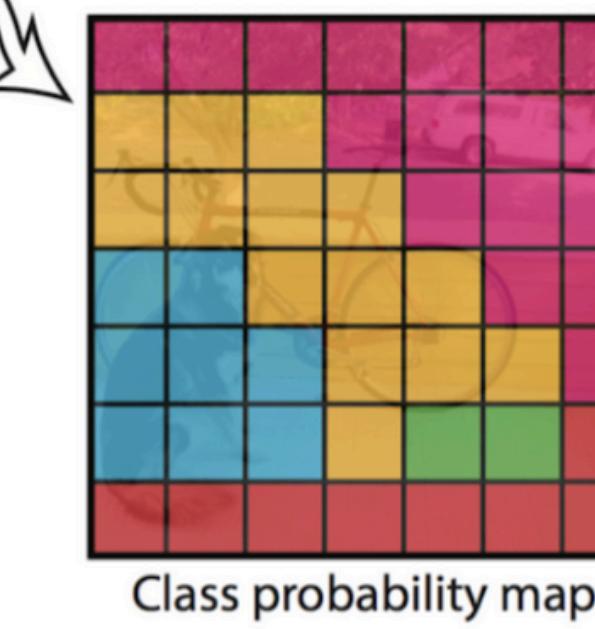
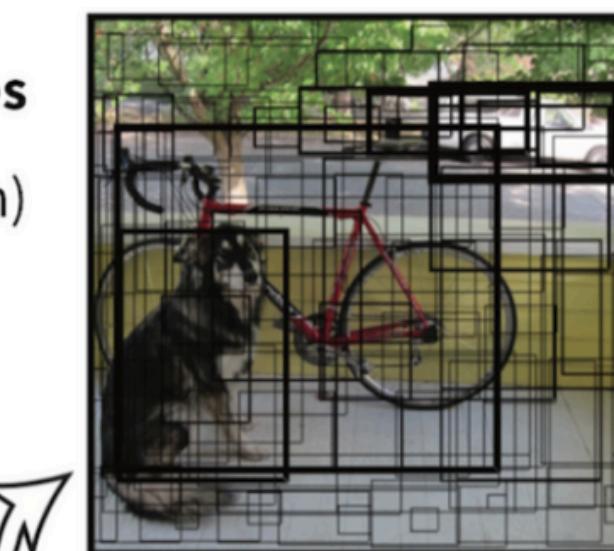
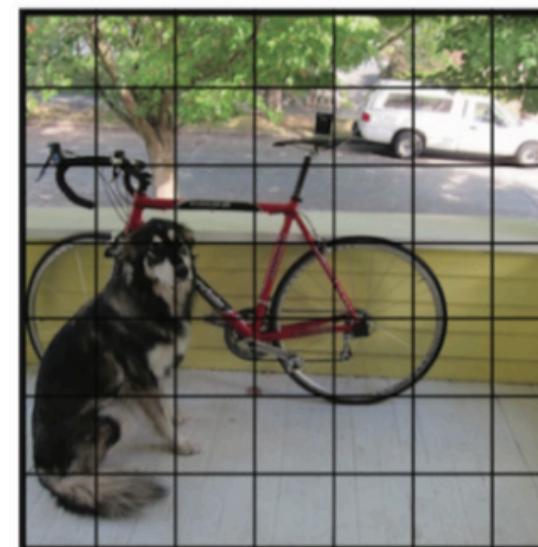
Ability to Detect Multiple Objects in One Cell

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

# YOLOv1

$S \times S \times B$  bounding boxes

$$\text{confidence} = Pr(\text{object}) \times \text{IoU}(\text{pred, truth})$$



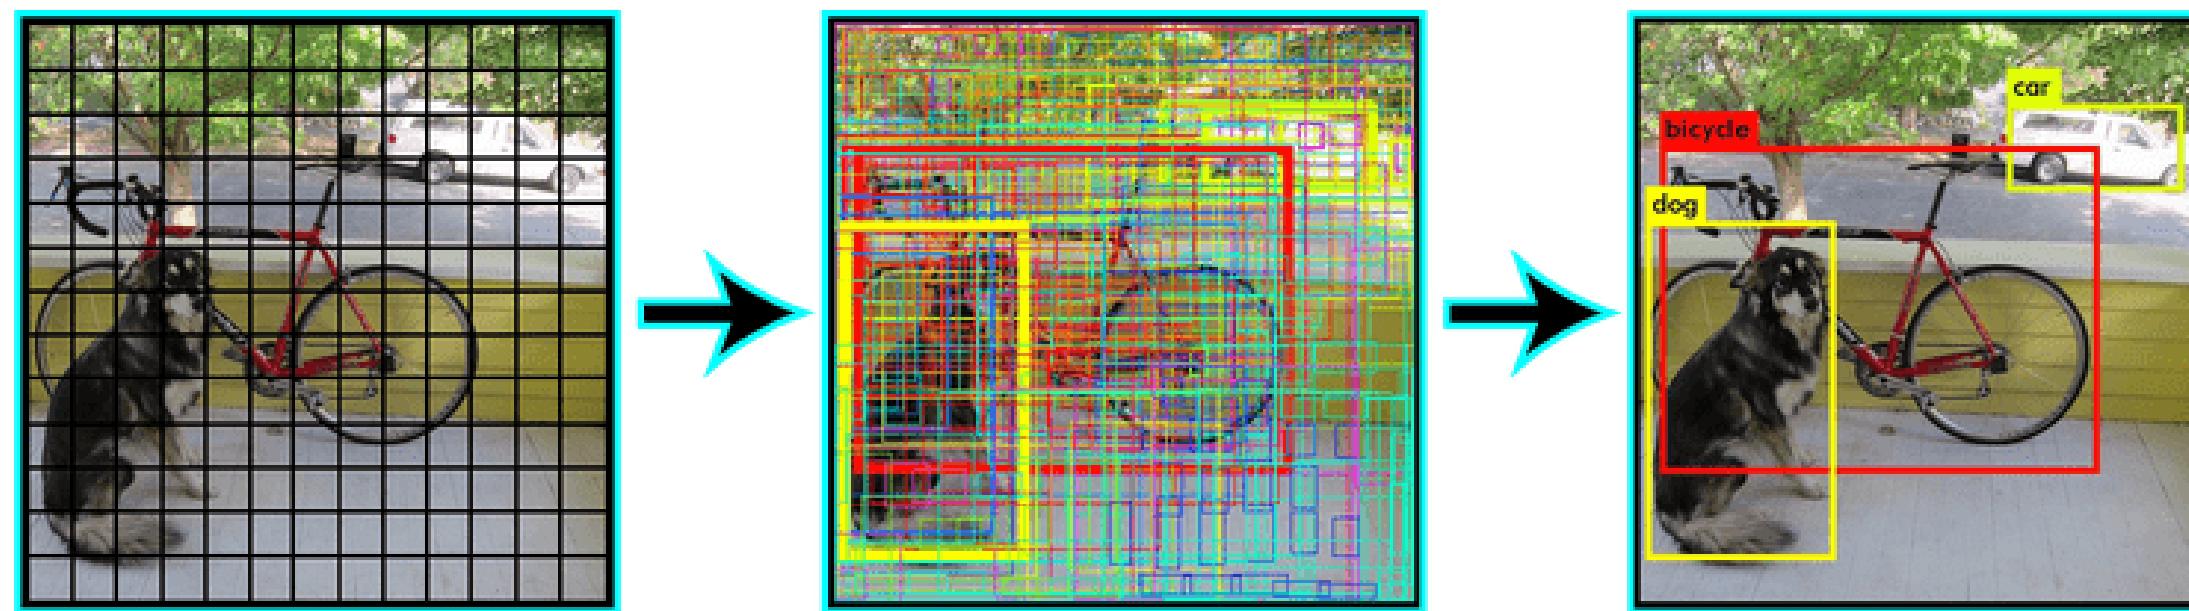
$$Pr(\text{Class}_i | \text{object})$$

detect 1 Objects for each Cell



- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

# YOLOv2



The image is divided into a **13x13** grid.

Each grid cell predicts **5 anchor boxes** and confidence scores.

- This results in a **13x13x425 output** tensor for the entire image

Each anchor box predicts the probability of each class. C=80

**Total Predictions:**  
5 anchor boxes ×  
(5 values (x,y,w,h,c)  
+ 80 class scores) =  
425 values

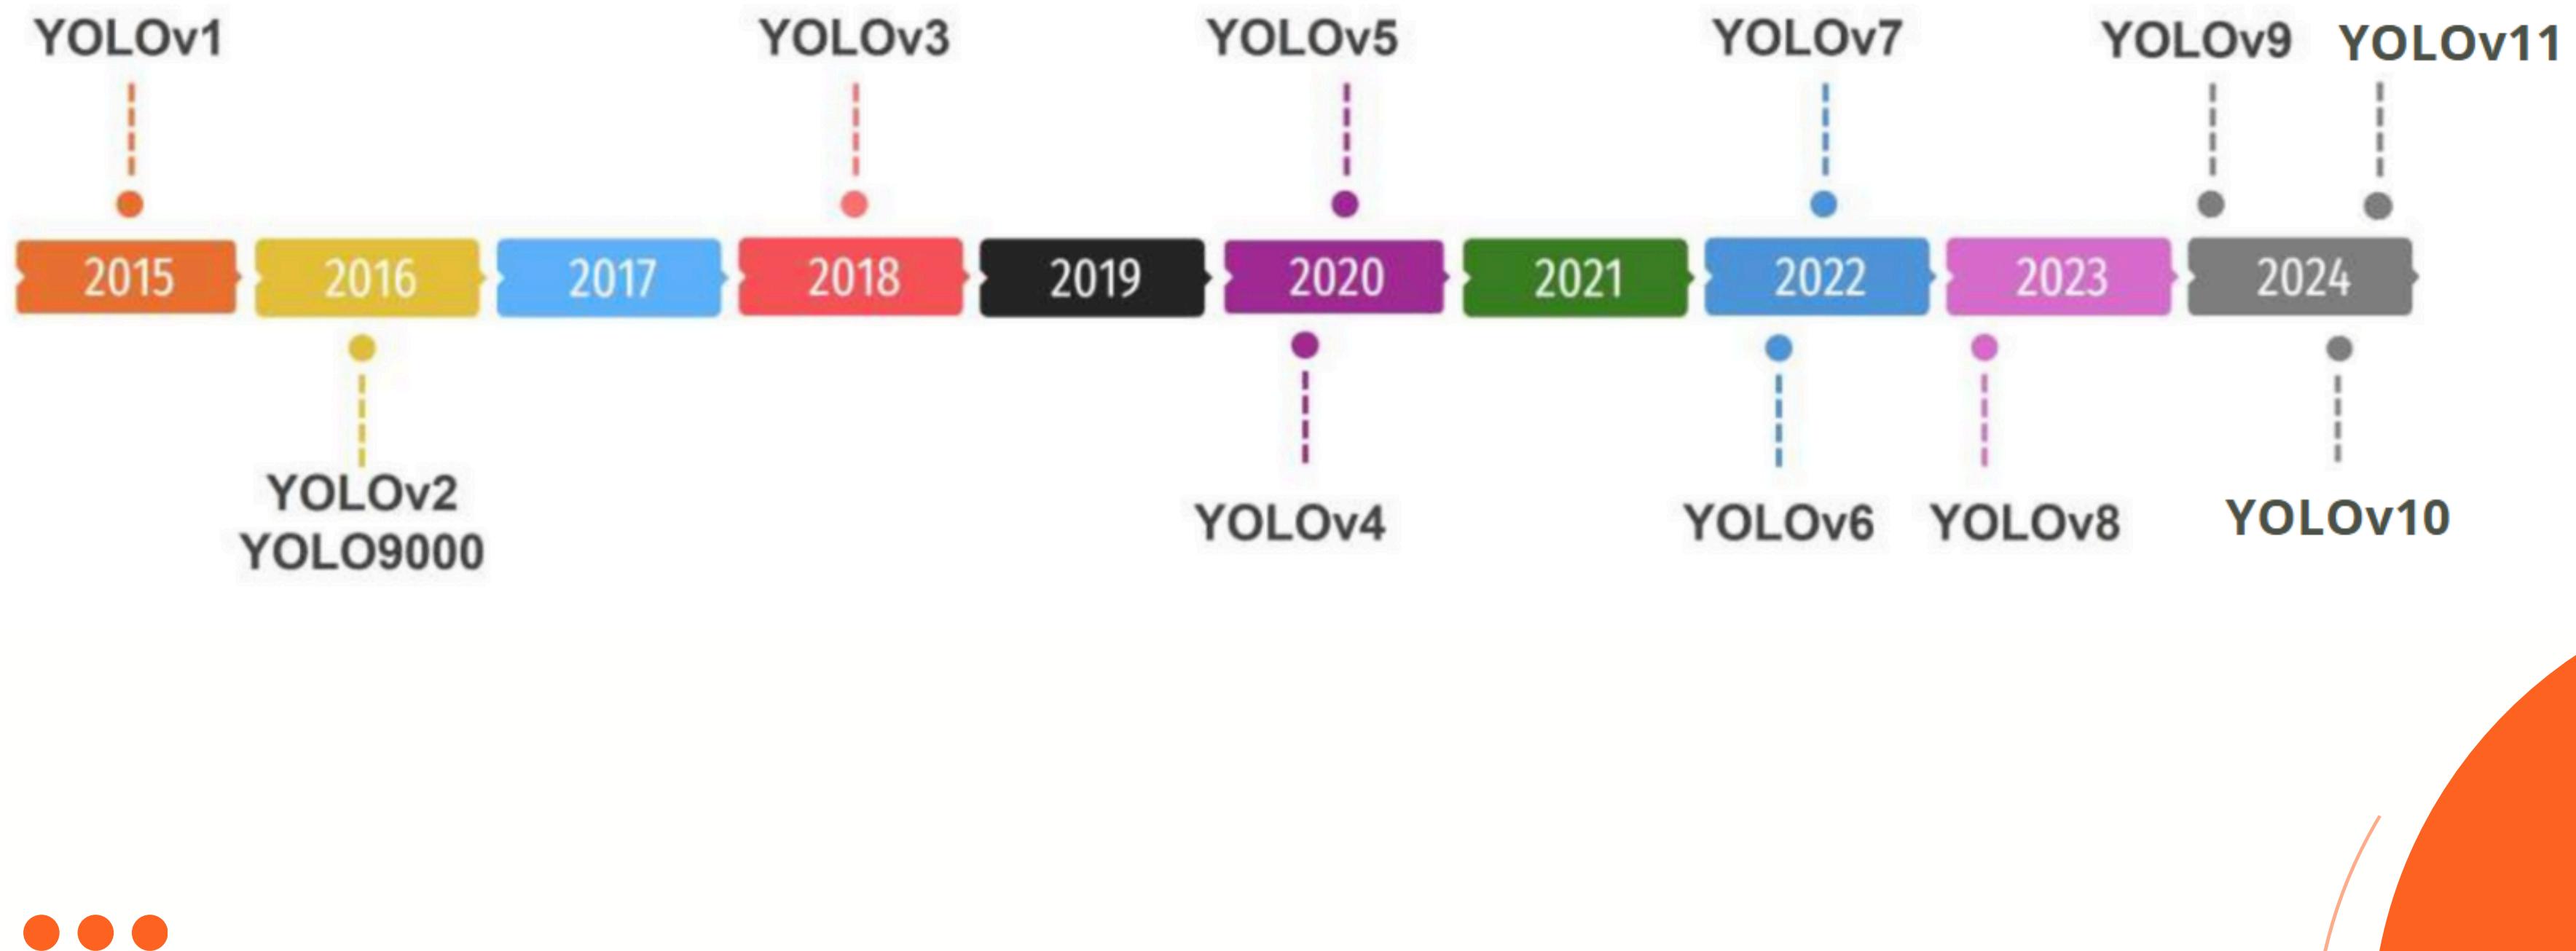


- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

<b>Aspect</b>	<b>YOLOv1</b>	<b>YOLOv2</b>
Backbone Network	Custom CNN (24 conv + 2 FC layers)	Darknet-19 (19 conv + 5 max-pooling layers)
Bounding Box Prediction	Direct prediction (2 boxes per grid cell)	Anchor boxes (5 per grid cell, k-means clustered)
Input Resolution	Fixed at 448x448	Multi-scale (320x320 to 608x608)
Feature Fusion	Single scale features	Passthrough layer for fine-grained features
Pre-training	Trained from scratch	Uses ImageNet pre-trained weights
Performance	~45 FPS, lower mAP	~67 FPS, higher mAP
Additional Features	Basic object detection	YOLO9000 (9000+ classes), improved detection



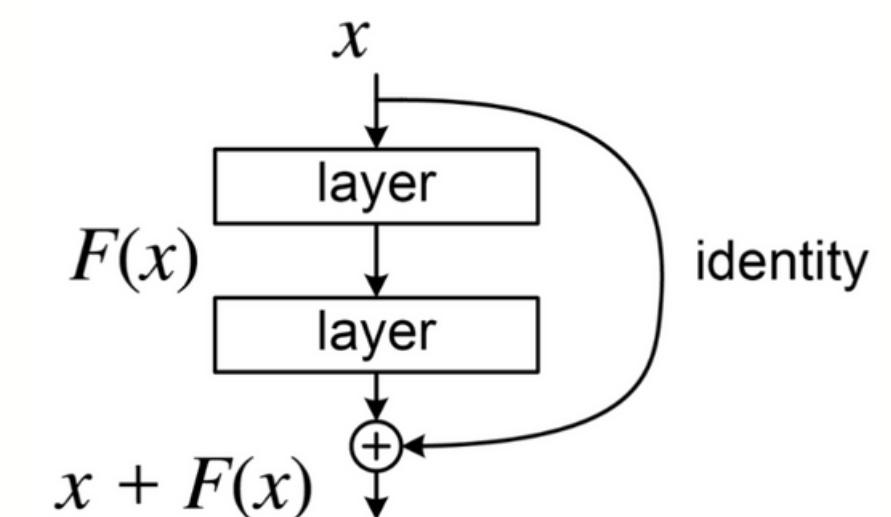
# Timeline of YOLO Versions



# YOLOv3

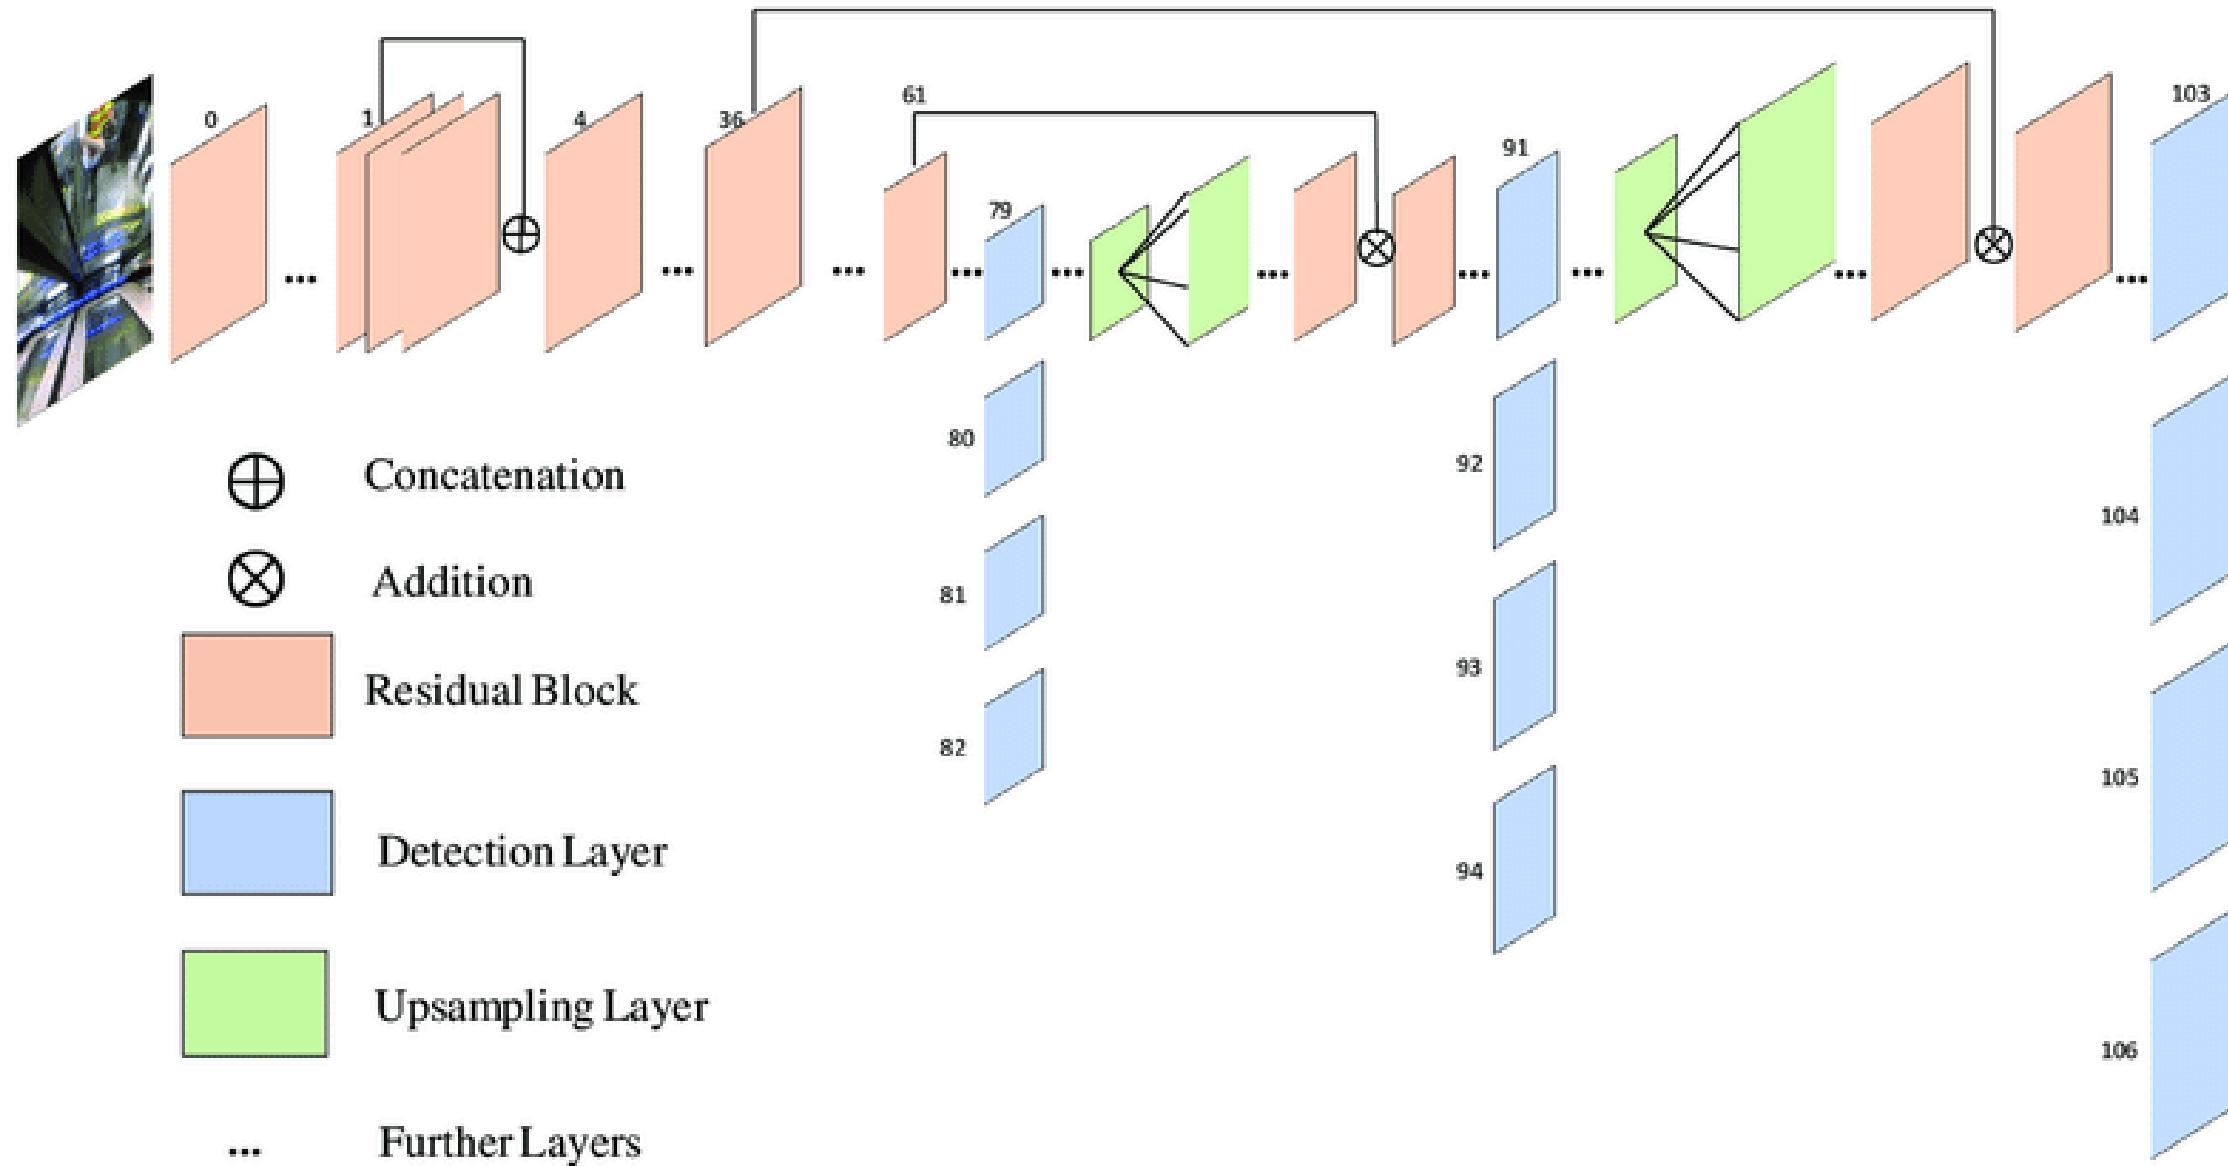
- Upgraded to **Darknet-53** with 53 convolutional layers, Incorporating **residual blocks** throughout the network,

Residual blocks solve the problem of vanishing or exploding gradients in very deep networks.



- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

# YOLOv3



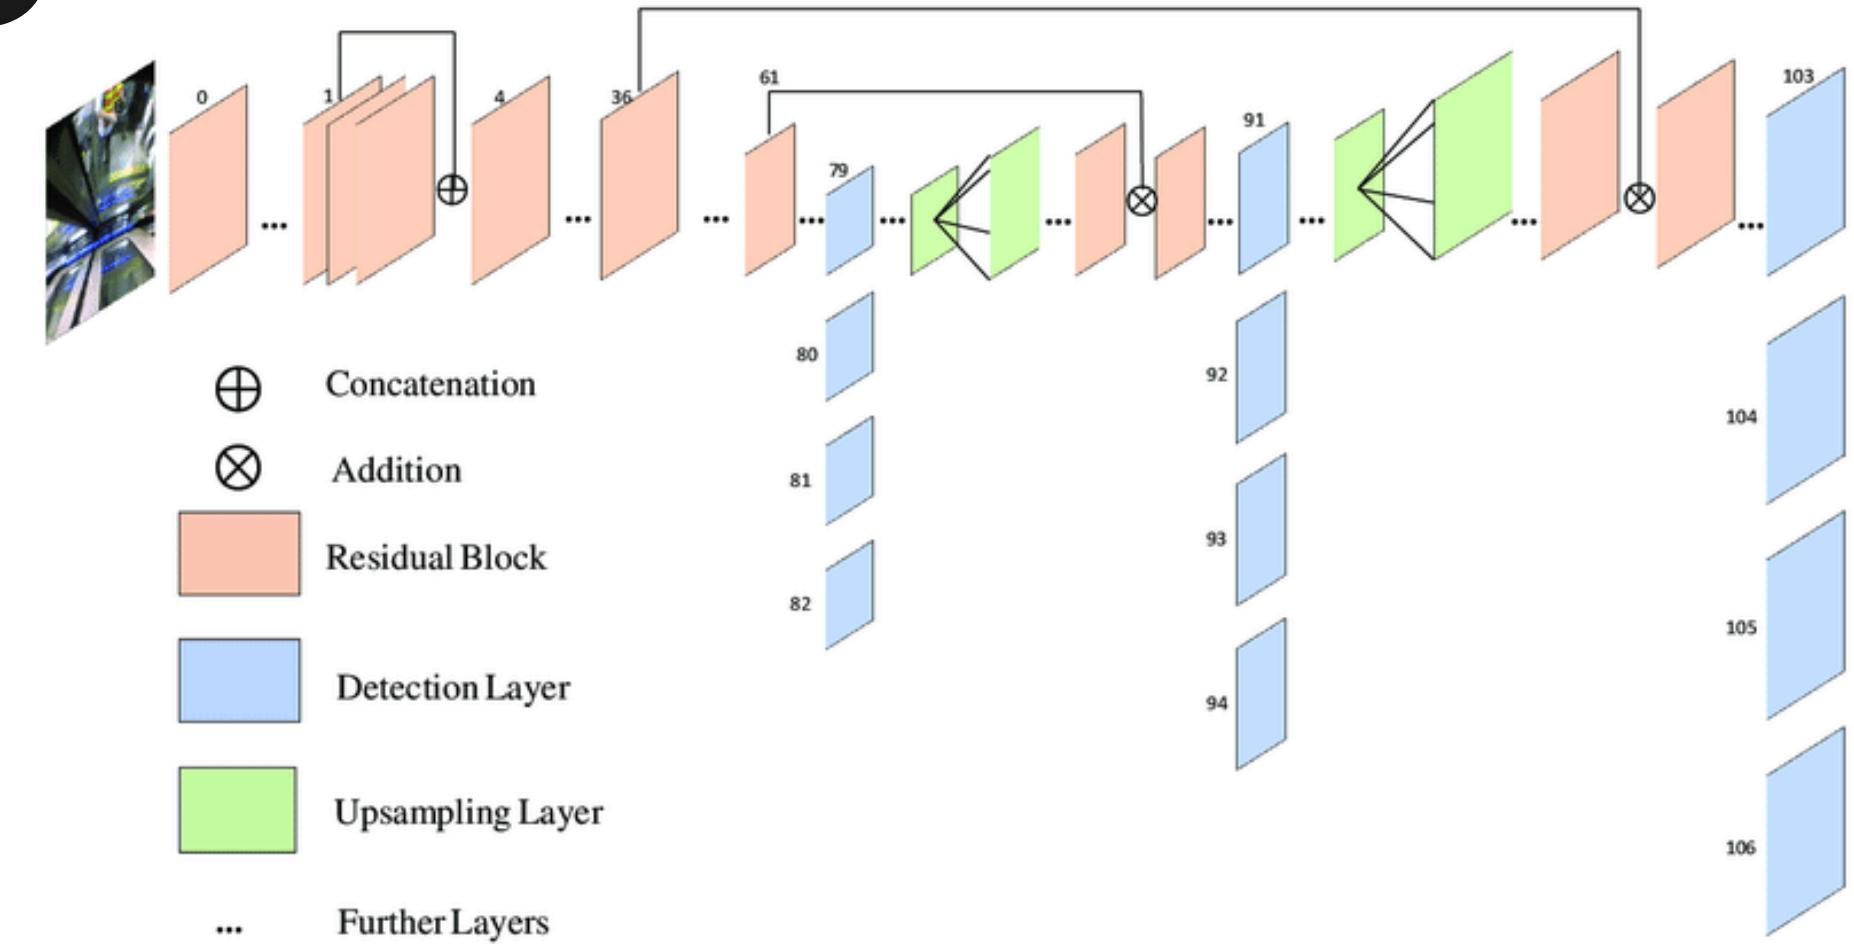
- Introduces **multi-scale prediction** by making predictions at three different scales. (13x13, 26x26, 52x52)



- 1
- 2
- 3**
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

# YOLOv3

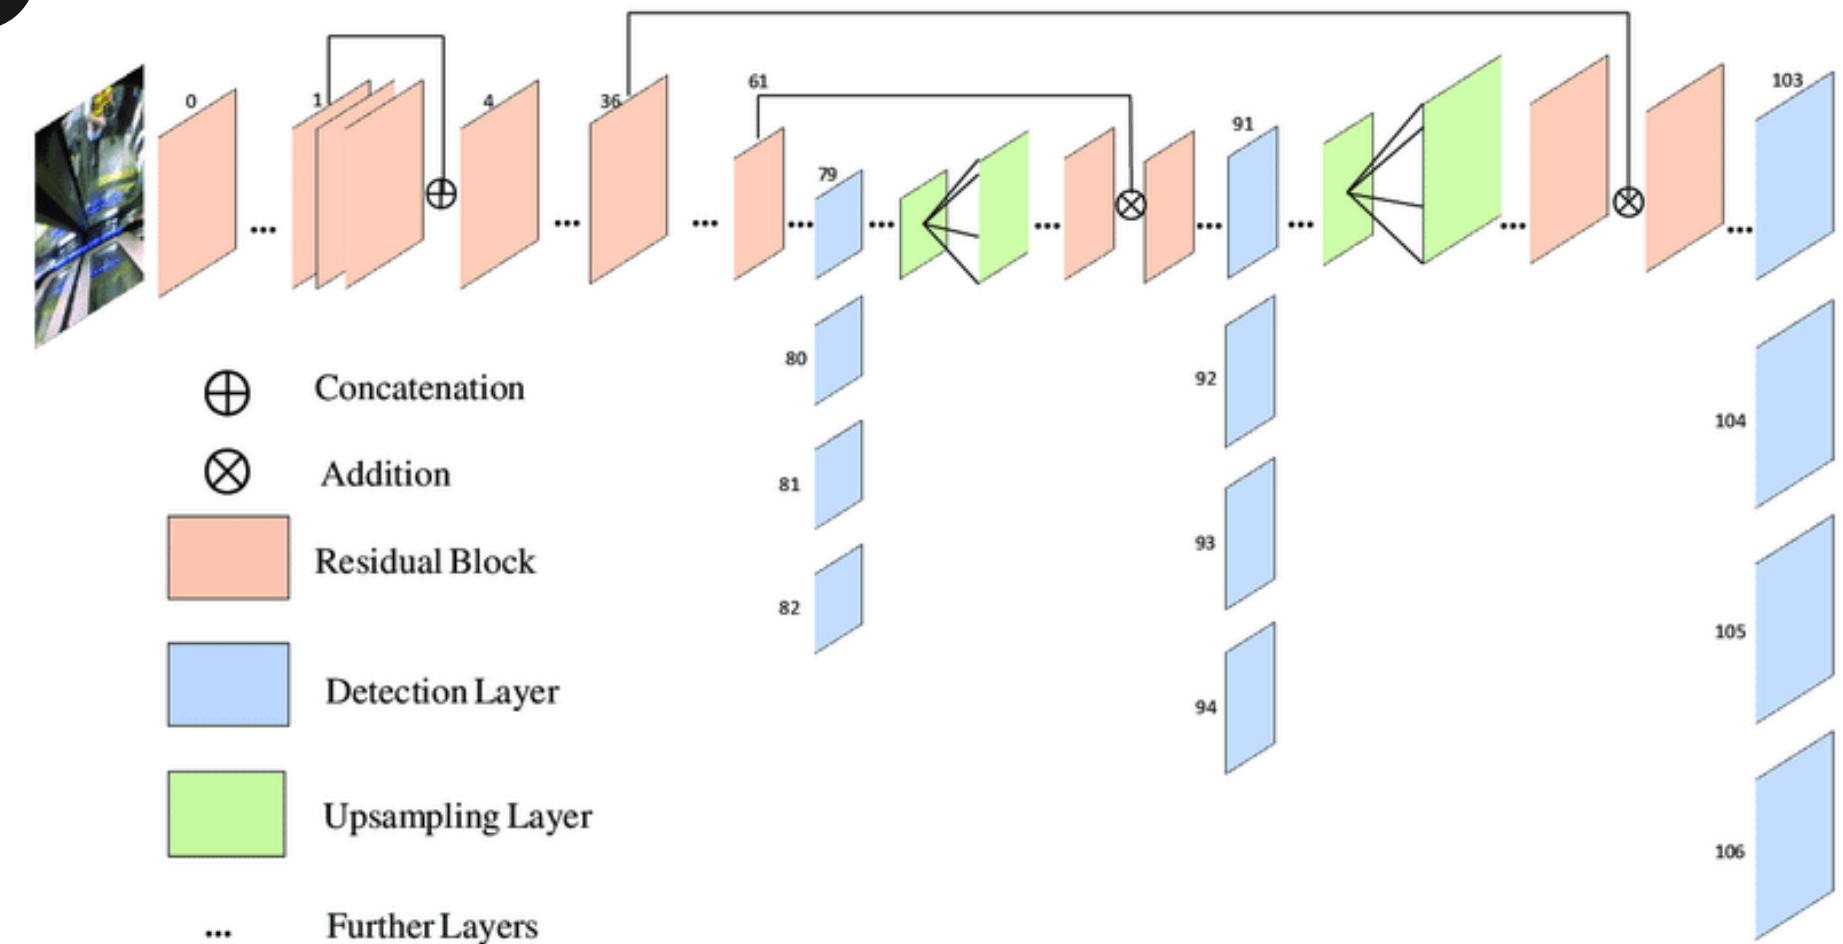
- Each detection layer is assigned **3 anchor boxes**, (9 anchor boxes divided across the three different scales.)
- **Total Predictions:** 3 anchor boxes  $\times$  ( 5 values + 80 class scores) = 255 values for each scale
- This results in a **13×13×255 | 26×26×255 | 52×52×255 output** tensor for the entire image.



- 1
- 2
- 3**
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

# YOLOv3

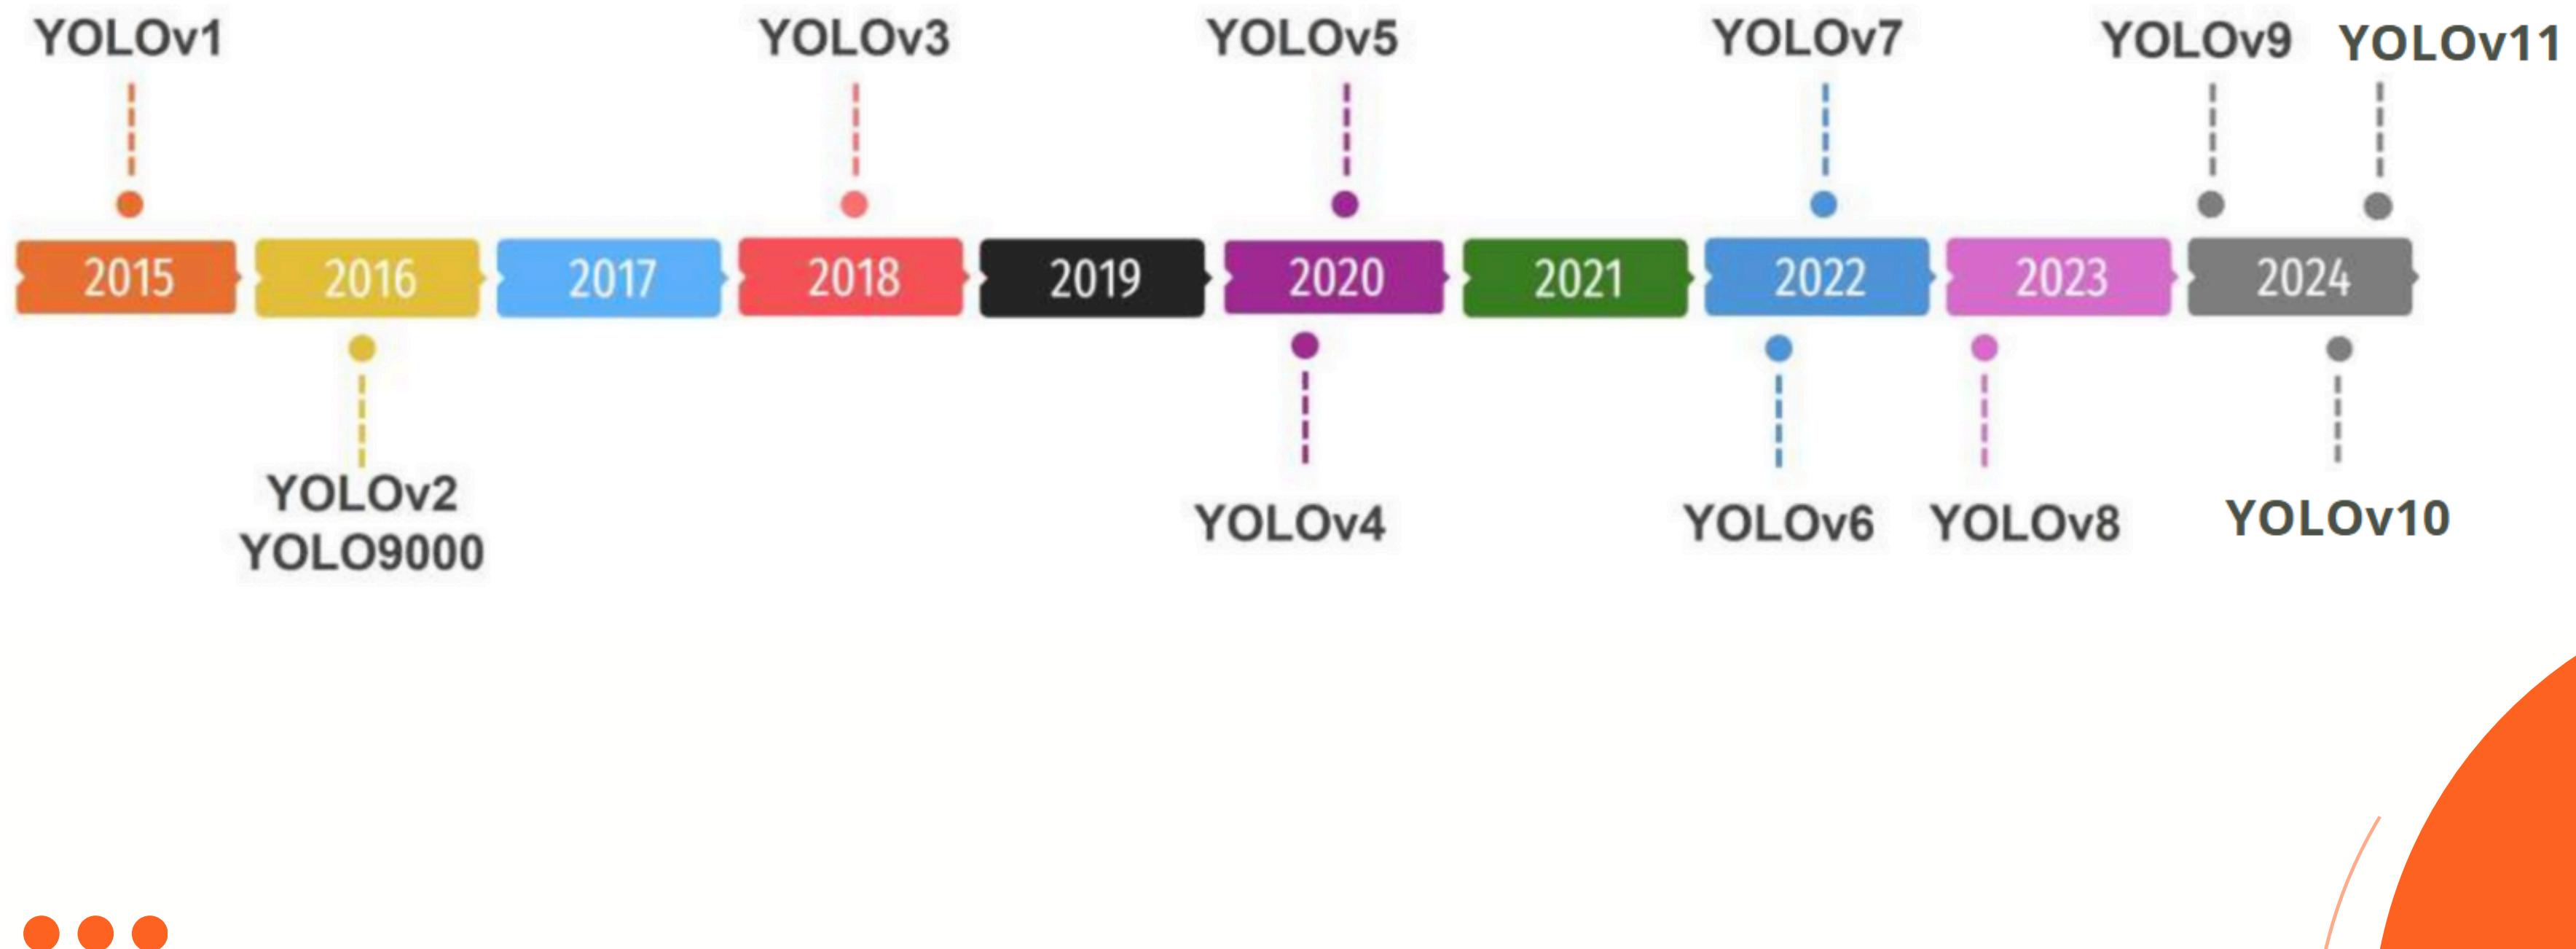
- Each detection layer is assigned **3 anchor boxes**, (9 anchor boxes divided across the three different scales.)



Replaces the softmax with independent logistic classifiers for each class, allowing **multi-label classification** (an object can belong to more than one class)

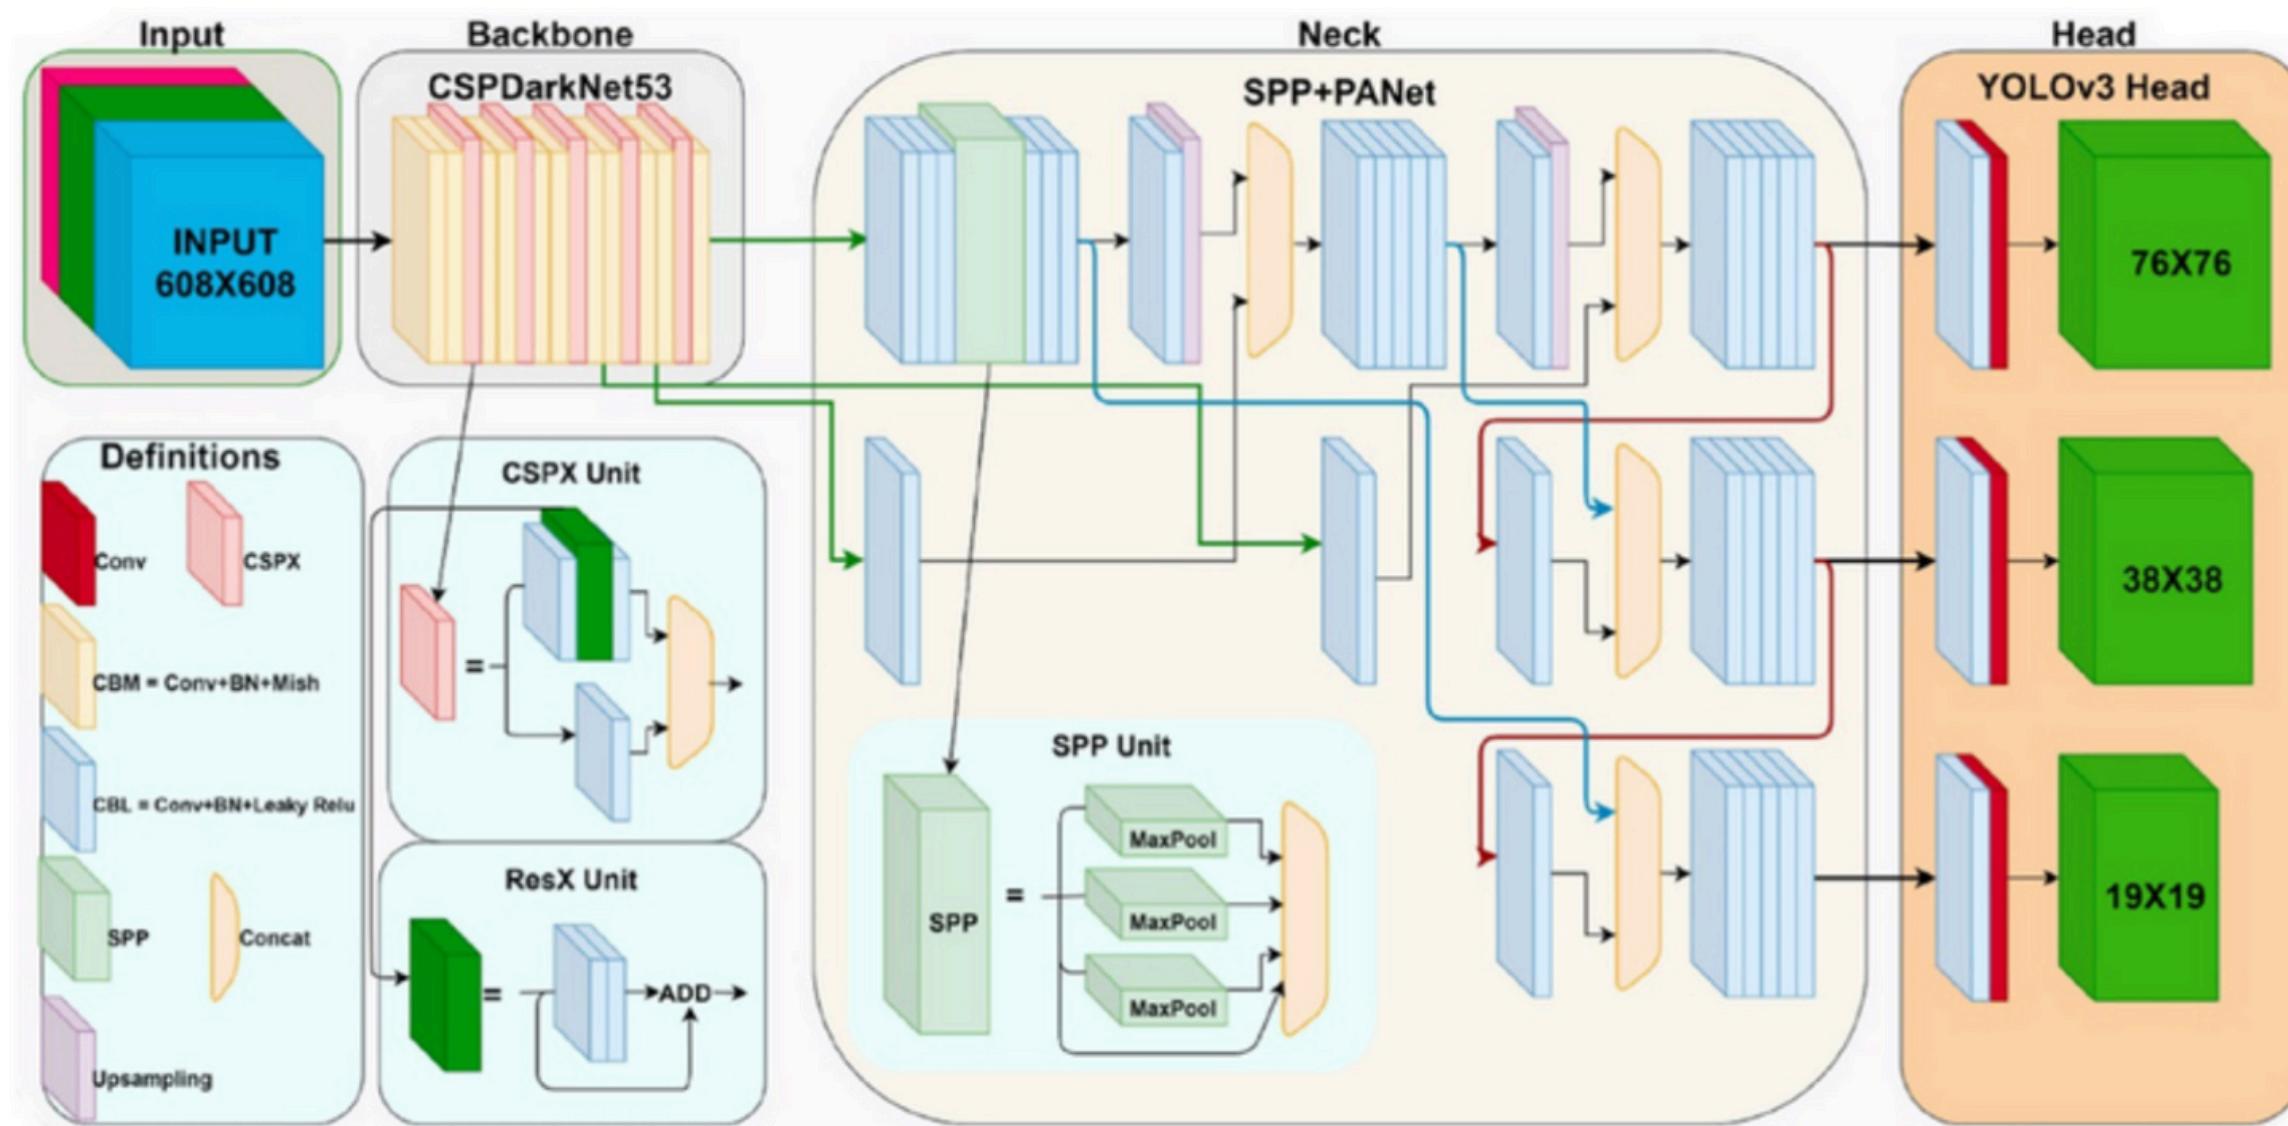
	<b>Aspect</b>	<b>YOLOv2</b>	<b>YOLOv3</b>
1	Backbone Network	Darknet-19	Darknet-53 with residual connections
2	Detection Scale	Single-scale	Multi-scale (3 scales for better accuracy)
3	Anchor Boxes	5 anchor boxes, single detection layer	9 anchor boxes, 3 scales (3 per scale)
4	Classification	Softmax (single-label classification)	Independent logistic classifiers (multi-label)
5	Objectness Score	Single sigmoid for object confidence	Refined logistic regression for better localization
6	Residual Blocks	Not used	Used, improving gradient flow and accuracy
7	Small Object Detection	Limited	Improved through multi-scale prediction
8	FPS (Frames per Second)	~45-67 FPS	~30-45 FPS
9	Overall Accuracy	Good for large objects	Improved accuracy, especially on small/multi-scale objects

# Timeline of YOLO Versions



# YOLOv4

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

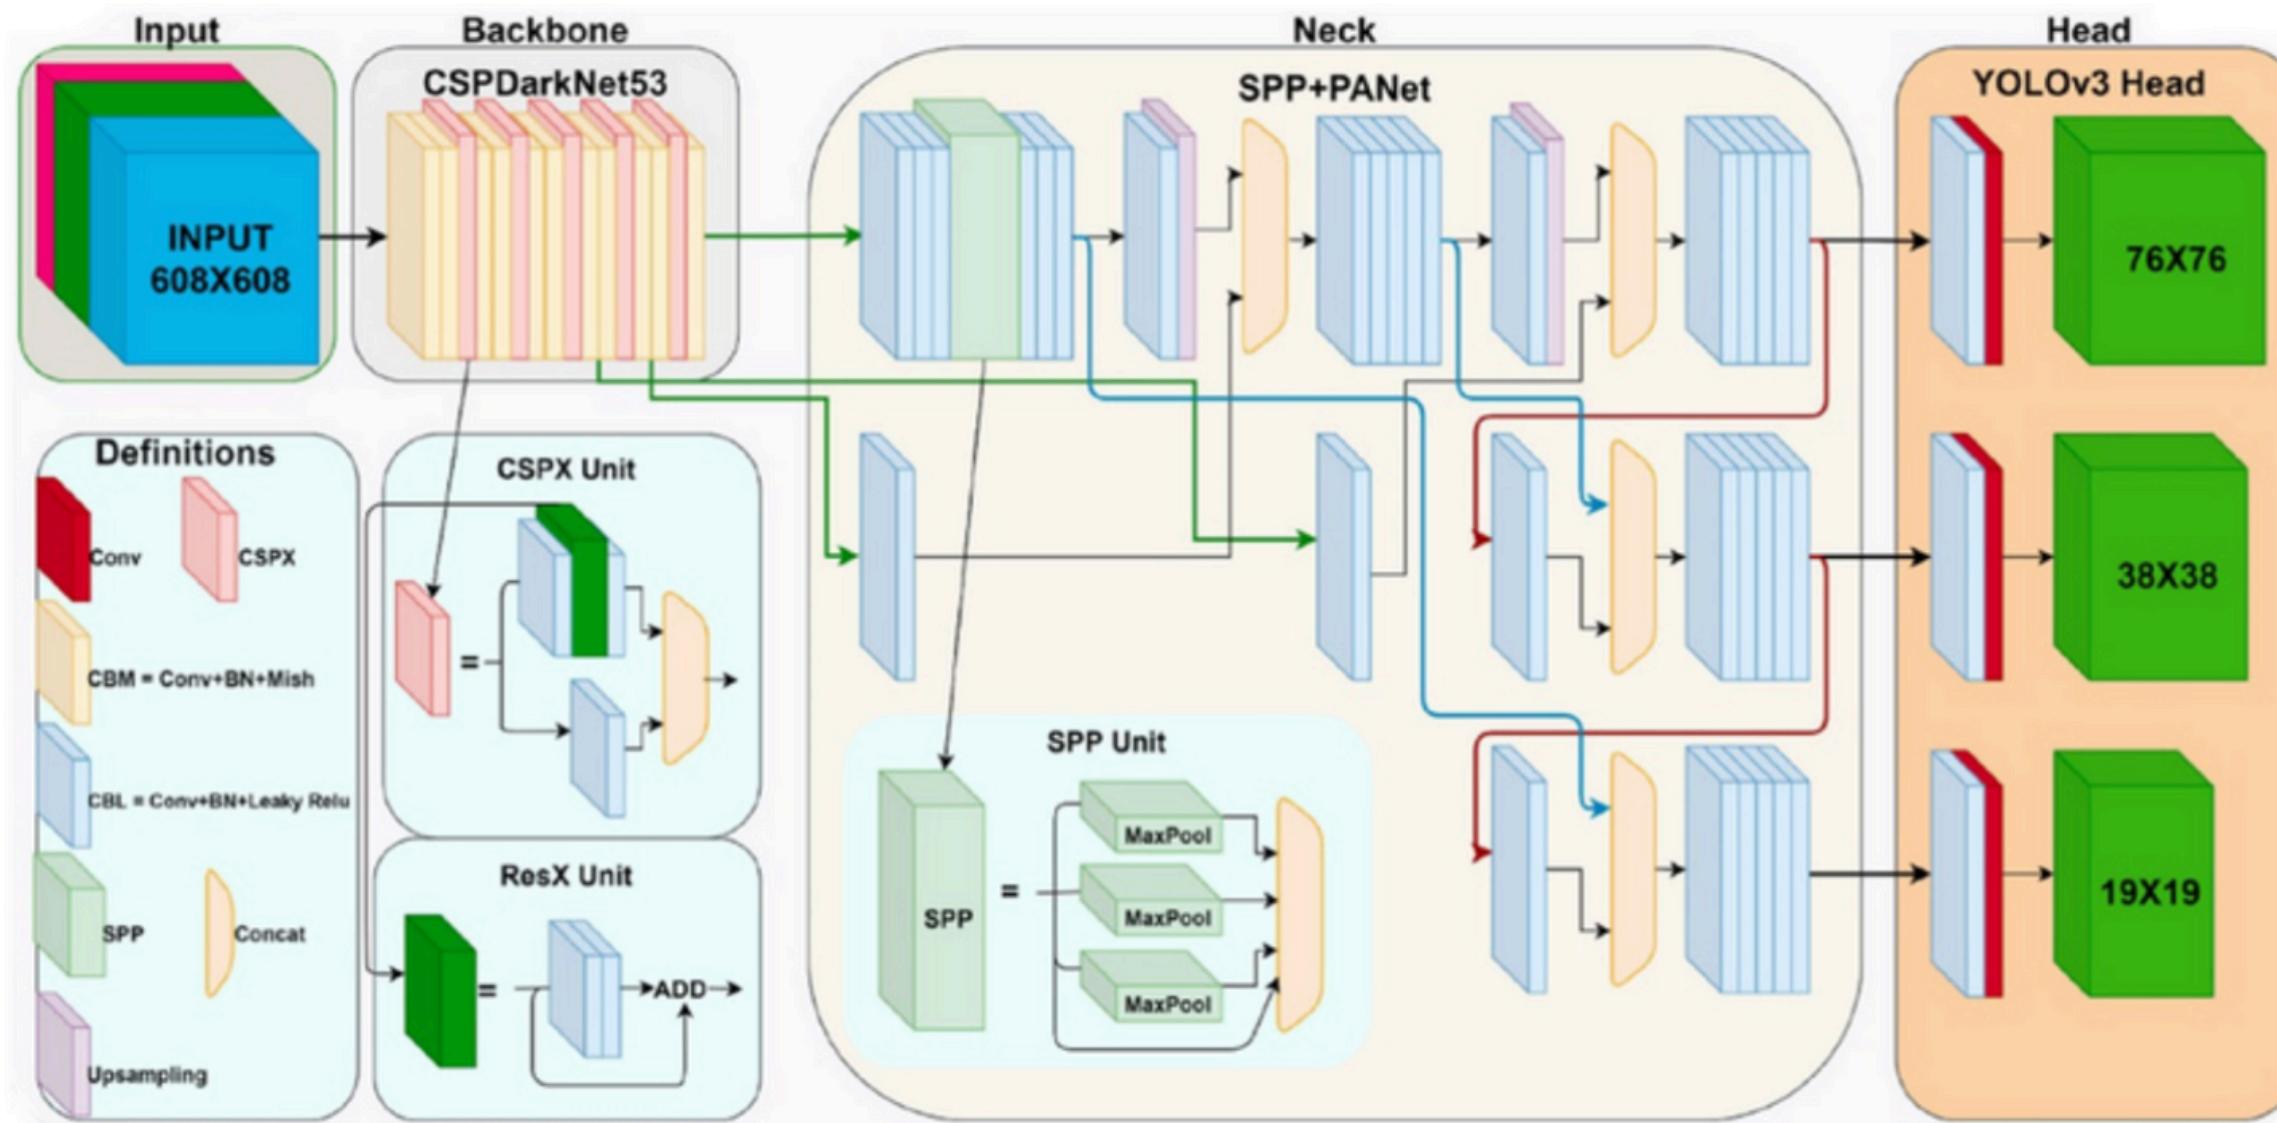


- Uses **CSPDarknet-53** as the backbone, an improved version of Darknet-53.



# YOLOv4

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11



from YOLOv3 and YOLOv4, the authors introduced **the Neck**, placed between the backbone and the head.

- The use of this block is to collect features of different stages from the backbone and pass them to the head.

# YOLOv4

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11

In the paper, they introduced two concepts:  
**bag of freebies and bag of specials.**

The concept of “**Bag of Freebies**” refers to techniques that modify the training approach or **raise the training cost without affecting the cost during inference**.

Data augmentation is a typical example of such a method.

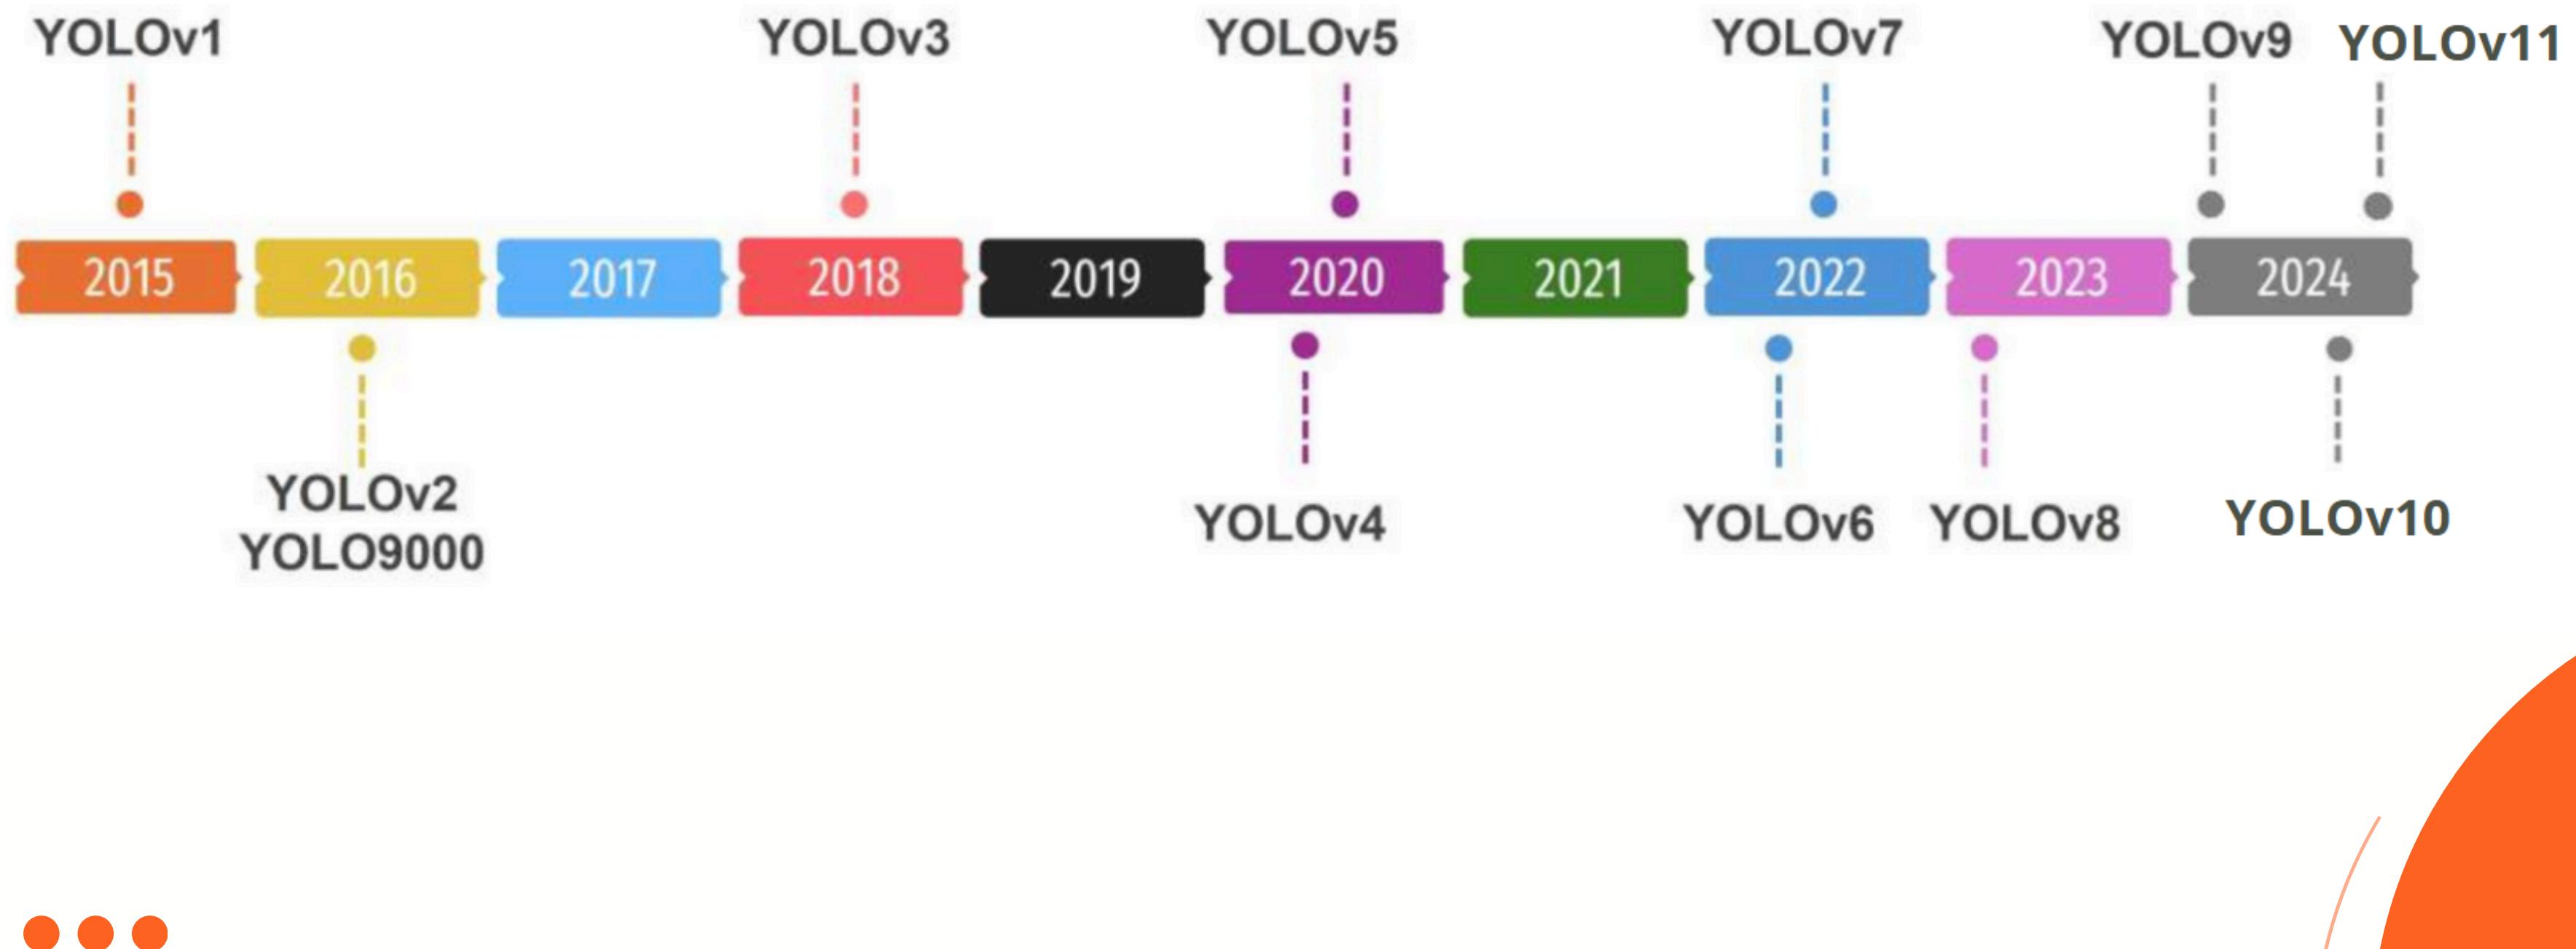
In contrast, “**Bag of Specials**” includes methods that marginally **increase the cost at inference time** but substantially **enhance accuracy**.

Specific examples like SPP (Spatial Pyramid Pooling) and PANet (Path Aggregation Network)

	<b>YOLOv3</b>	<b>YOLOv4</b>
1		
2	Backbone Network	Darknet-53
3		CSPDarknet-53
4	Neck	None
5	Bag of Freebies	None
6	Bag of Specials	None
7	Activation Function	Leaky ReLU
8		Mish
9	Speed (FPS)	~45 FPS
10		~45 FPS (but higher mAP)
11	Accuracy (COCO mAP)	~33%
		~43%

mean Average Precision (mAP) on COCO Dataset

# Timeline of YOLO Versions



# YOLOv5

While YOLOv1 to 4 was developed by the original YOLO community researchers with **C-based framework**, YOLOv5 was created by Ultralytics, implemented in **PyTorch**,

This makes YOLOv5 more accessible and easy to deploy and adapt in various machine learning pipelines.

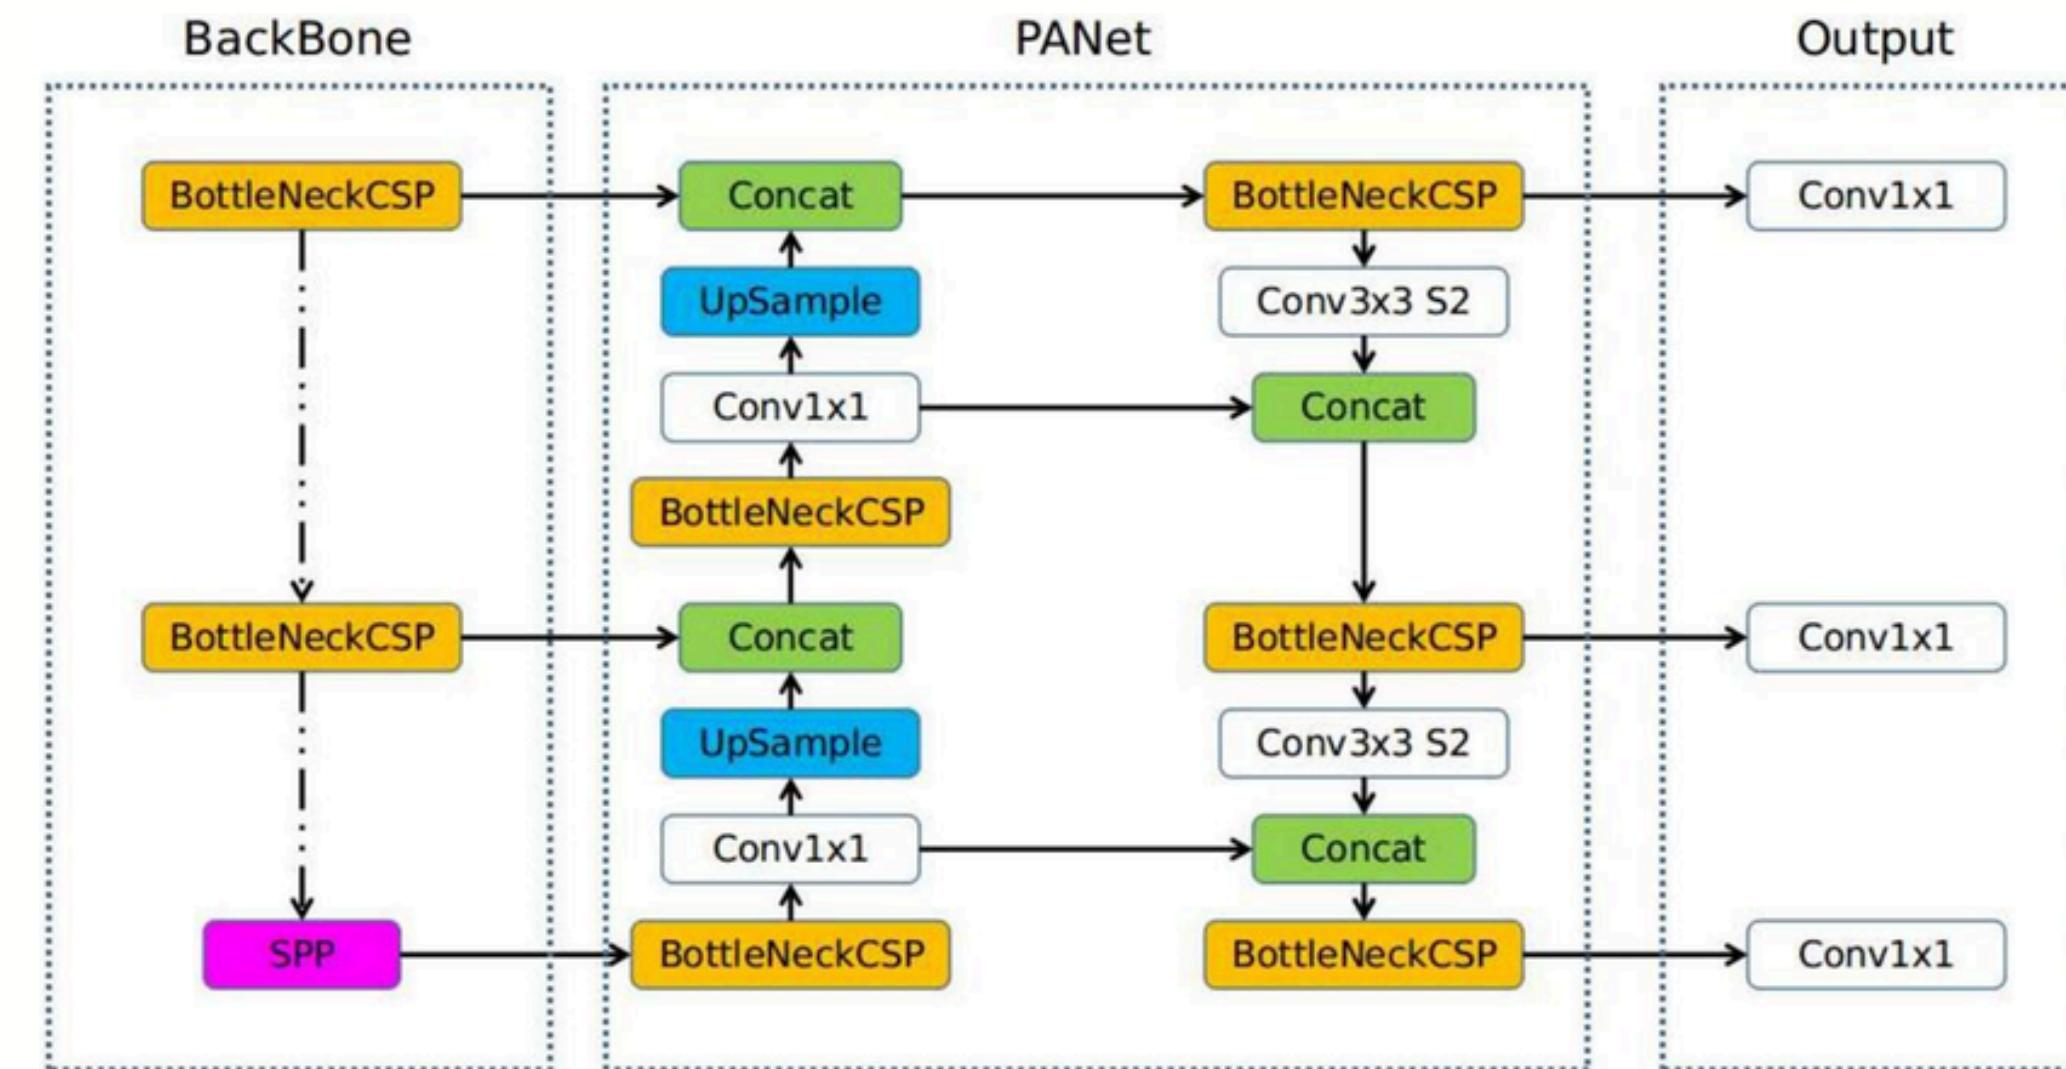
Uses a simpler, more efficient backbone and neck optimized directly for PyTorch

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11



# YOLOv5

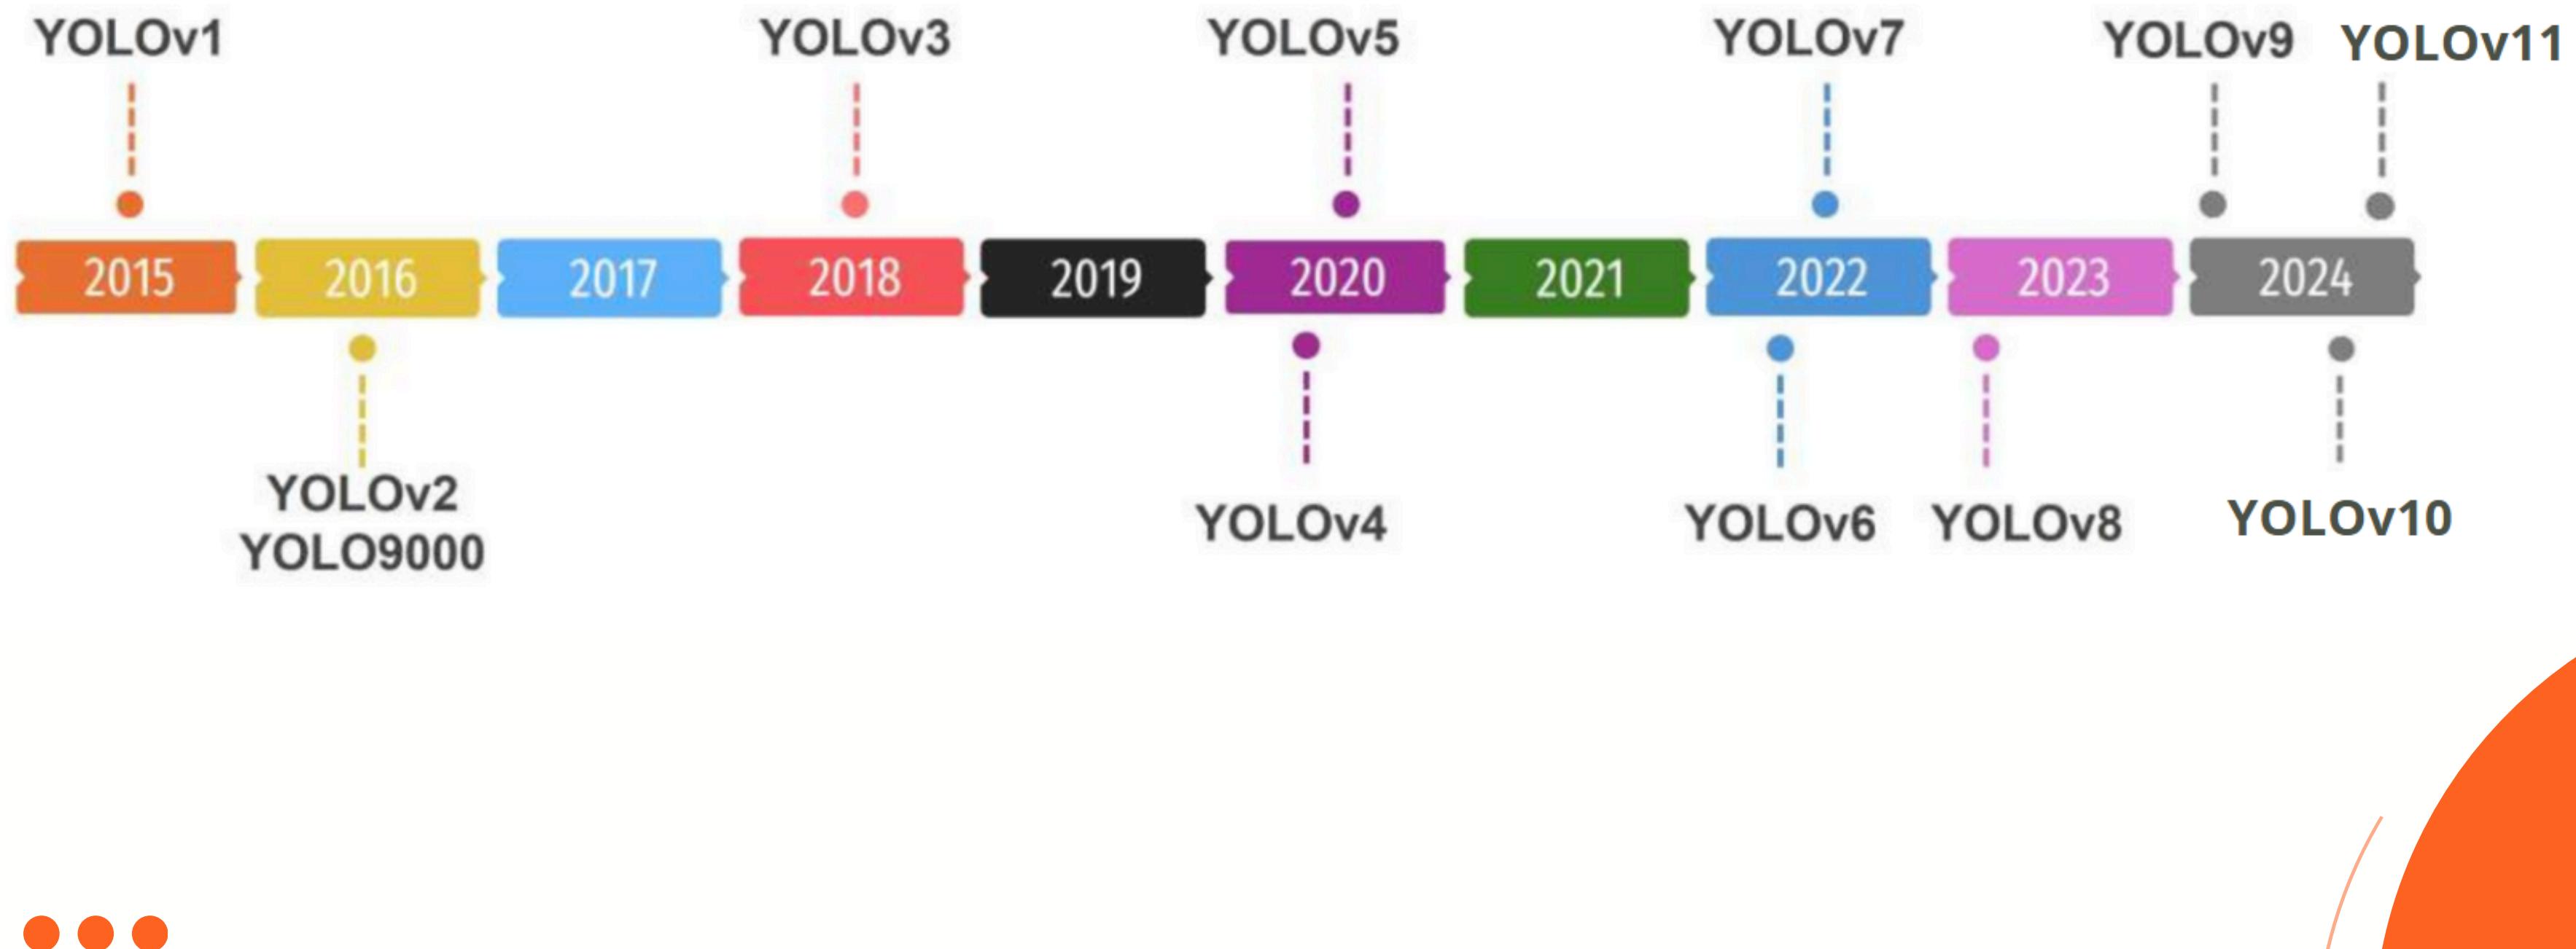
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11



it provides four model sizes (YOLOv5s, YOLOv5m, YOLOv5l, YOLOv5x)  
to balance speed and accuracy.

	<b>Aspect</b>	<b>YOLOv4</b>	<b>YOLOv5</b>
1	Framework	Darknet	PyTorch
2	Backbone	CSPDarknet53	Custom backbone optimized for PyTorch
3	Neck	PANet + SPP	PANet with optimizations
4	Head	Three output layers for three scales	Three output layers for three scales
5	Model Sizes	Single size (manually configurable)	YOLOv5s, YOLOv5m, YOLOv5l, YOLOv5x
6	Training Complexity	Complex (requires Darknet with advanced tuning)	Simple (end-to-end in PyTorch)
7	Inference Speed	Slower than YOLOv3 but improved mAP	Faster, highly optimized for edge devices
8	Performance (COCO mAP)	~43%	~45% (varies by model size)
9	Deployment Flexibility	Limited by Darknet	High, with native PyTorch and ONNX support

# Timeline of YOLO Versions



# YOLOv7

Introduced in July 2022, YOLOv7 marked a significant leap forward from its predecessors, showcasing improved accuracy and speed enhancements ranging from 5 FPS to 160 FPS.

The YOLOv7 model introduces novel modifications such as E-ELAN, Model Scaling, Planned re-parameterized convolution, Coarse for auxiliary, and penalty for lead loss.

1

2

3

4

5

6

7

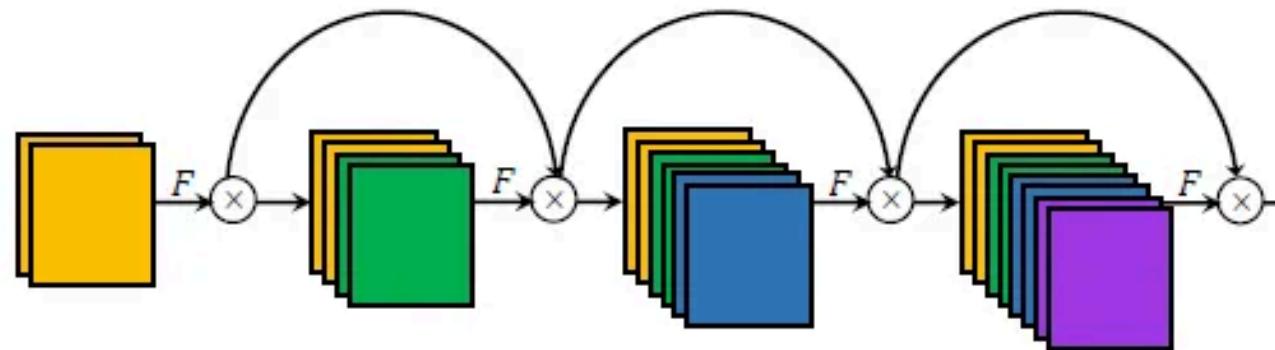
8

9

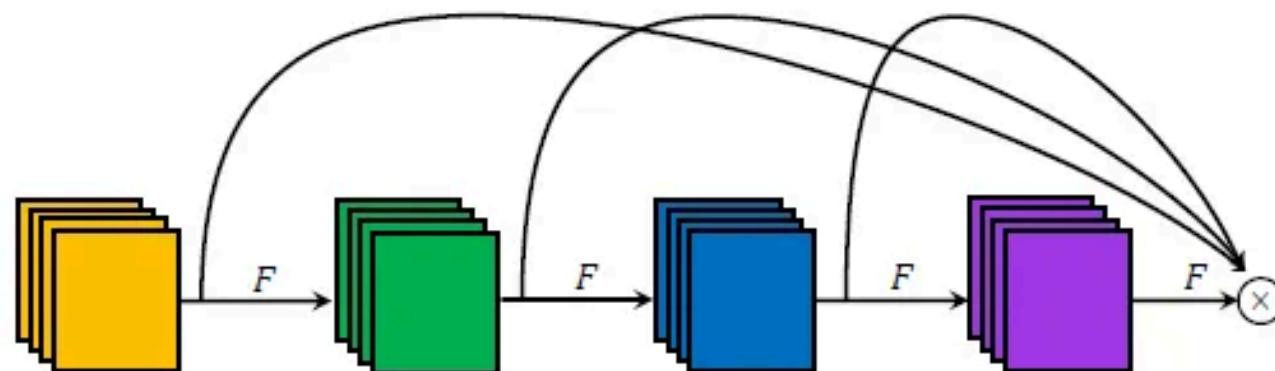
...

# YOLOv7

## Extended-ELAN (E-ELAN):



(a) Dense Aggregation (DenseNet)



(b) One-Shot Aggregation (VoVNet)

E-ELAN is an Extended efficient layer aggregation network, a variant of ELAN. ELAN is inspired by VoVNet and CSPNet. We already know about CSPNet, and VoVNet is nothing but an Object Detection Network composed of cascaded OSA modules.

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

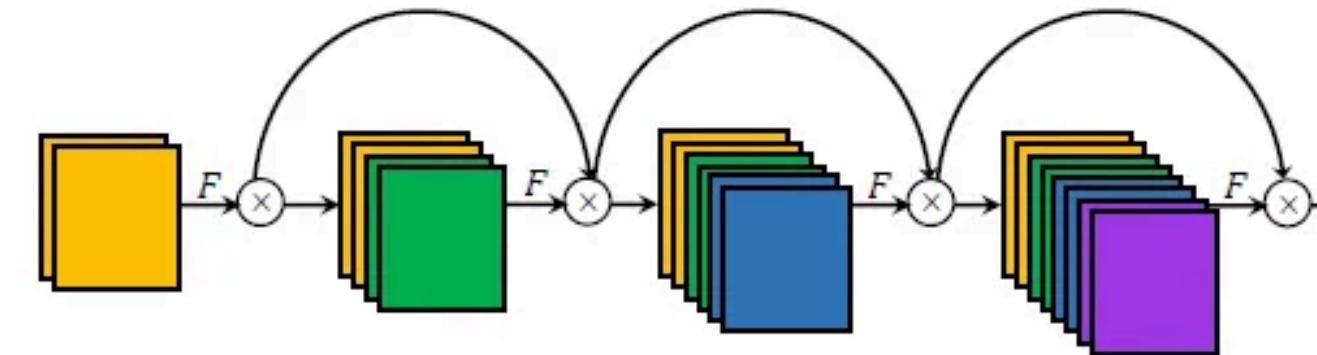


# YOLOv7

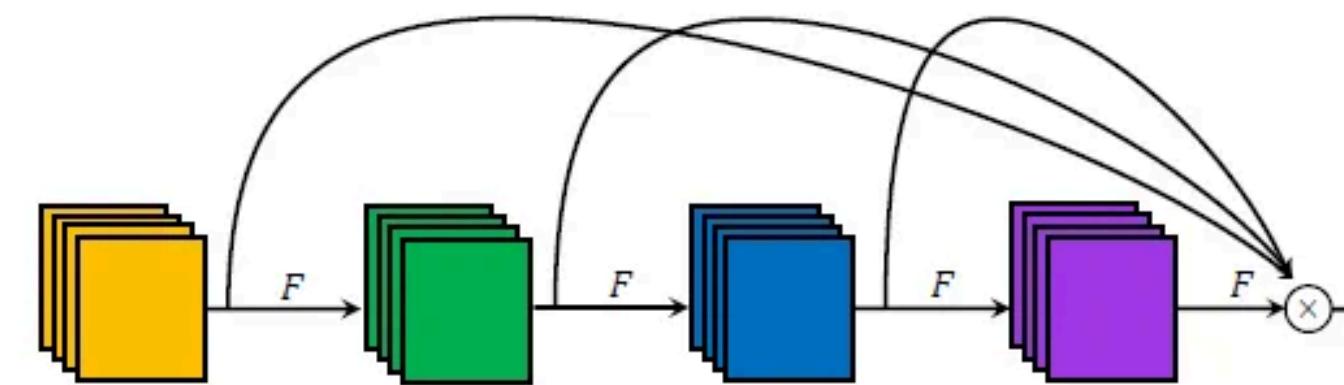
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

OSA means One-shot Aggregation. It is more efficient than Denses Block in DenseNet and optimized for GPU computation. The below image clears the concept of OSA,

It is similar to DensesNet, but we concatenate the features after a few conv blocks here.



(a) Dense Aggregation (DenseNet)



(b) One-Shot Aggregation (VoVNet)

# YOLOv7

## **Compound Model Scaling:**

The authors used model scaling to adjust parameters to generate models of different scales. By different scales, it means a change in model parameters. It helps to meet the need for different inference speeds.

In this model, they used NAS (Network Architecture Search), which is a commonly used model scaling method, authors also showed that this method can be further improved using the compound model scaling approach.

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



# YOLOv7

## **Model Re-parameterization:**

Model re-parameterization means model weight averaging. There are two ways of doing model averaging,

- Averaging the weights of the same model trained on different folds of data.
- Averaging the model weights of different epochs.

This is a very popular ensemble technique. Pytorch has an implementation of the same named SWA(Stochastic Weight Averaging).

1

2

3

4

5

6

7

8

9



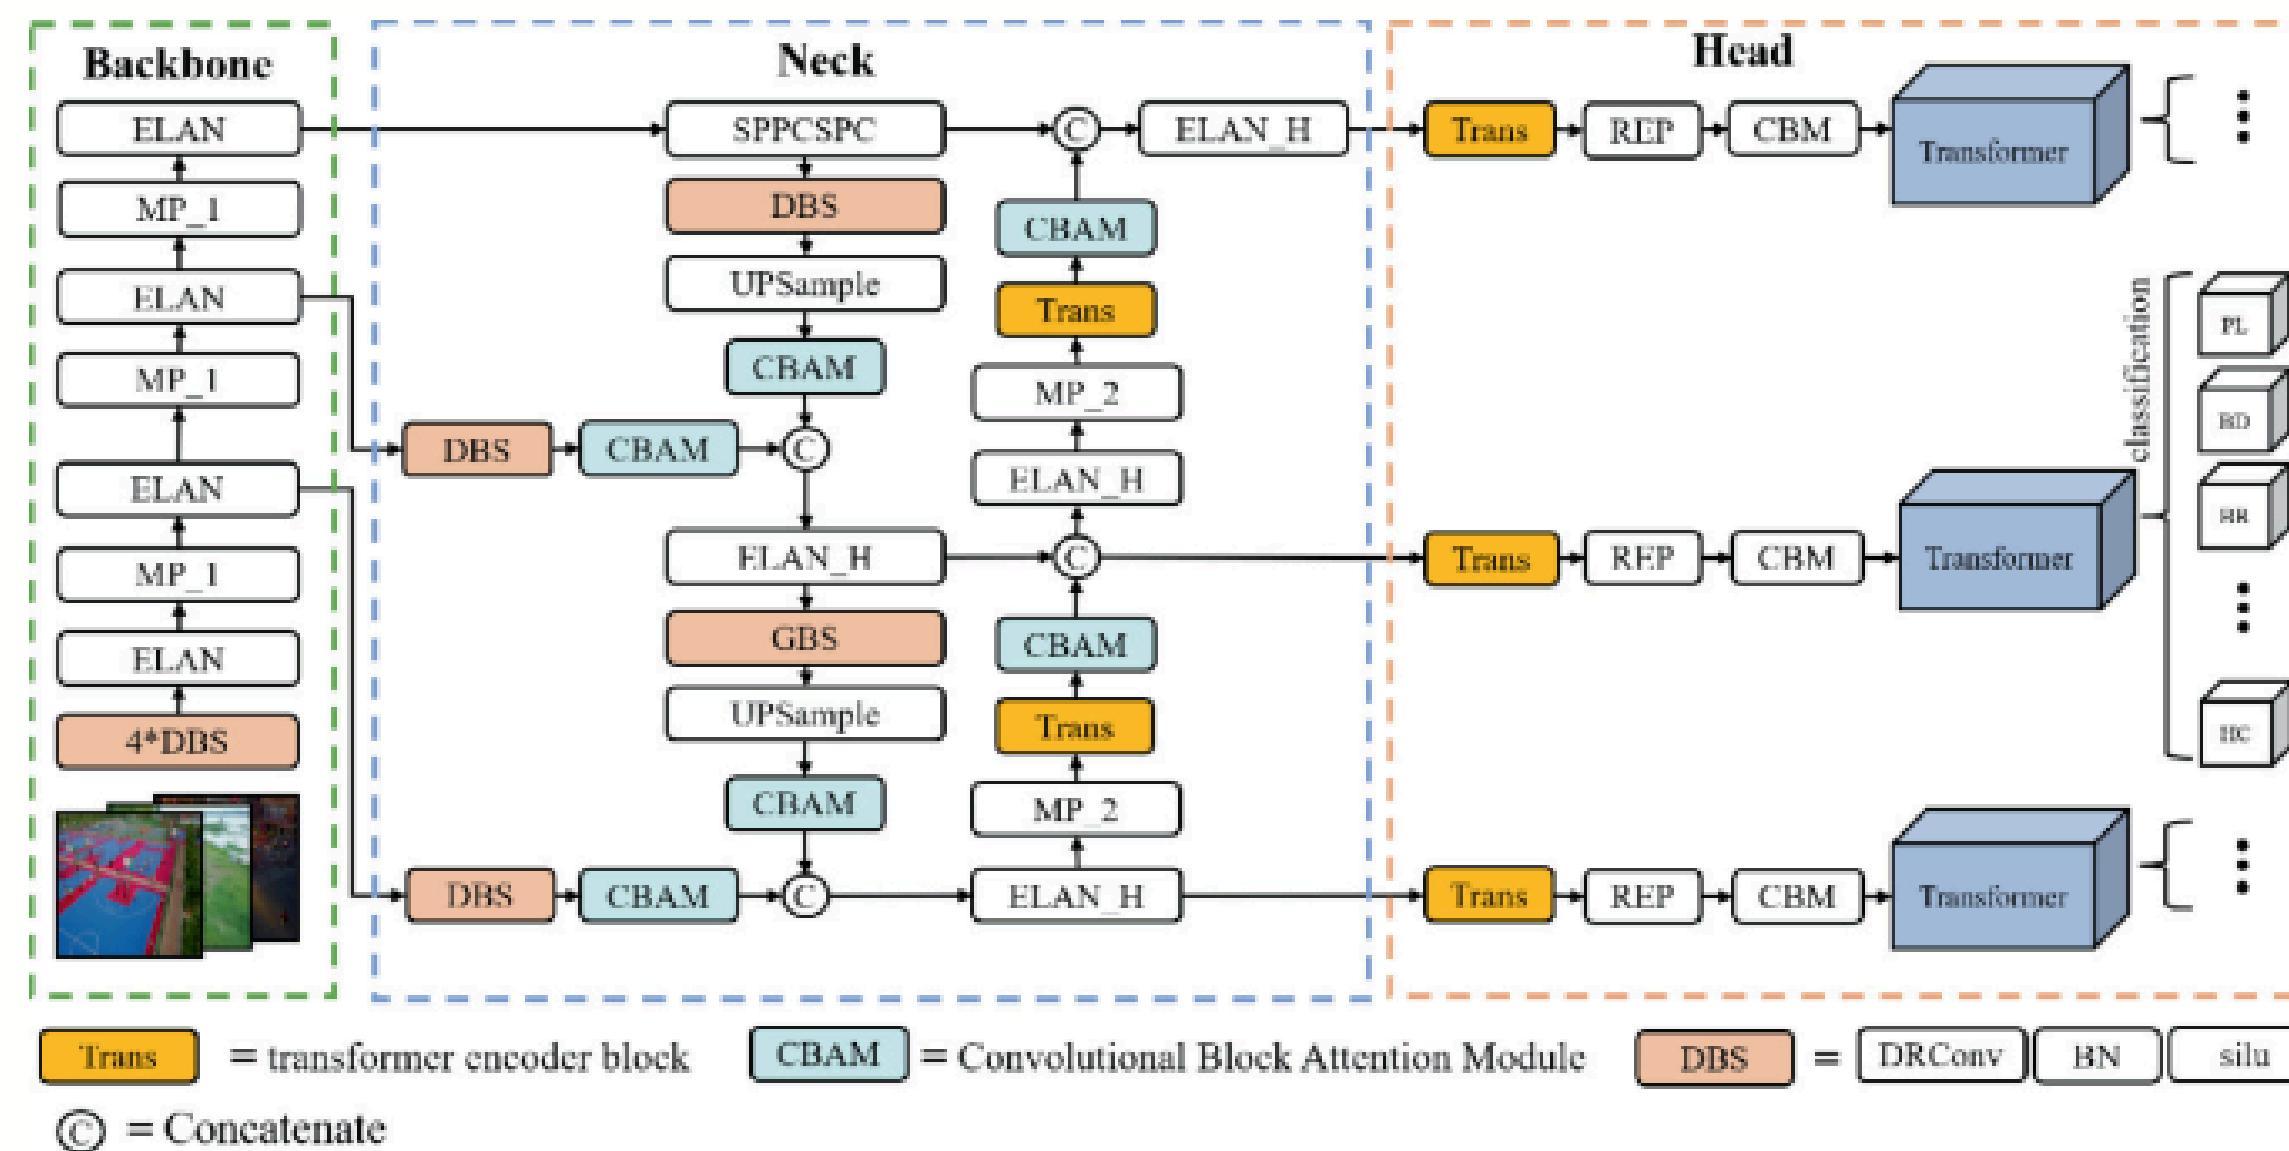
# YOLOv7



- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

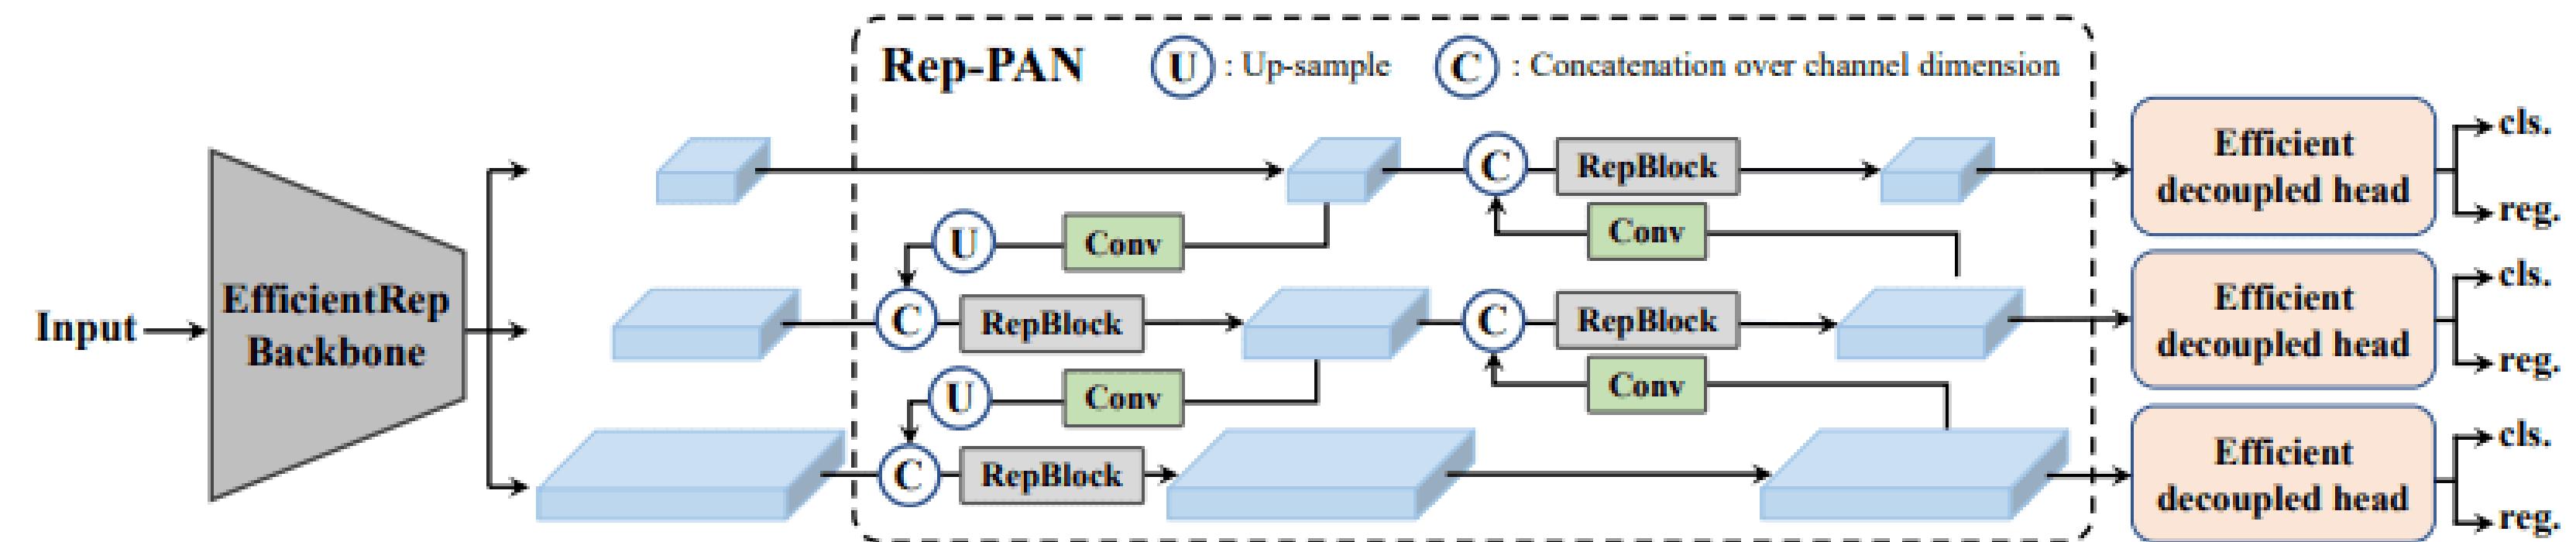
## Coarse for Auxiliary and Fine for Lead Loss:

As you can see, the model architecture has multiple heads predicting the same thing. The head responsible for the final output is the lead head, and the other heads assist in the model training. With the help of an assistant loss, the weights of the auxiliary heads are updated.



# YOLOv6

YOLOv6 is an anchor-free, decoupled head architecture with a unique backbone named EfficientRep.



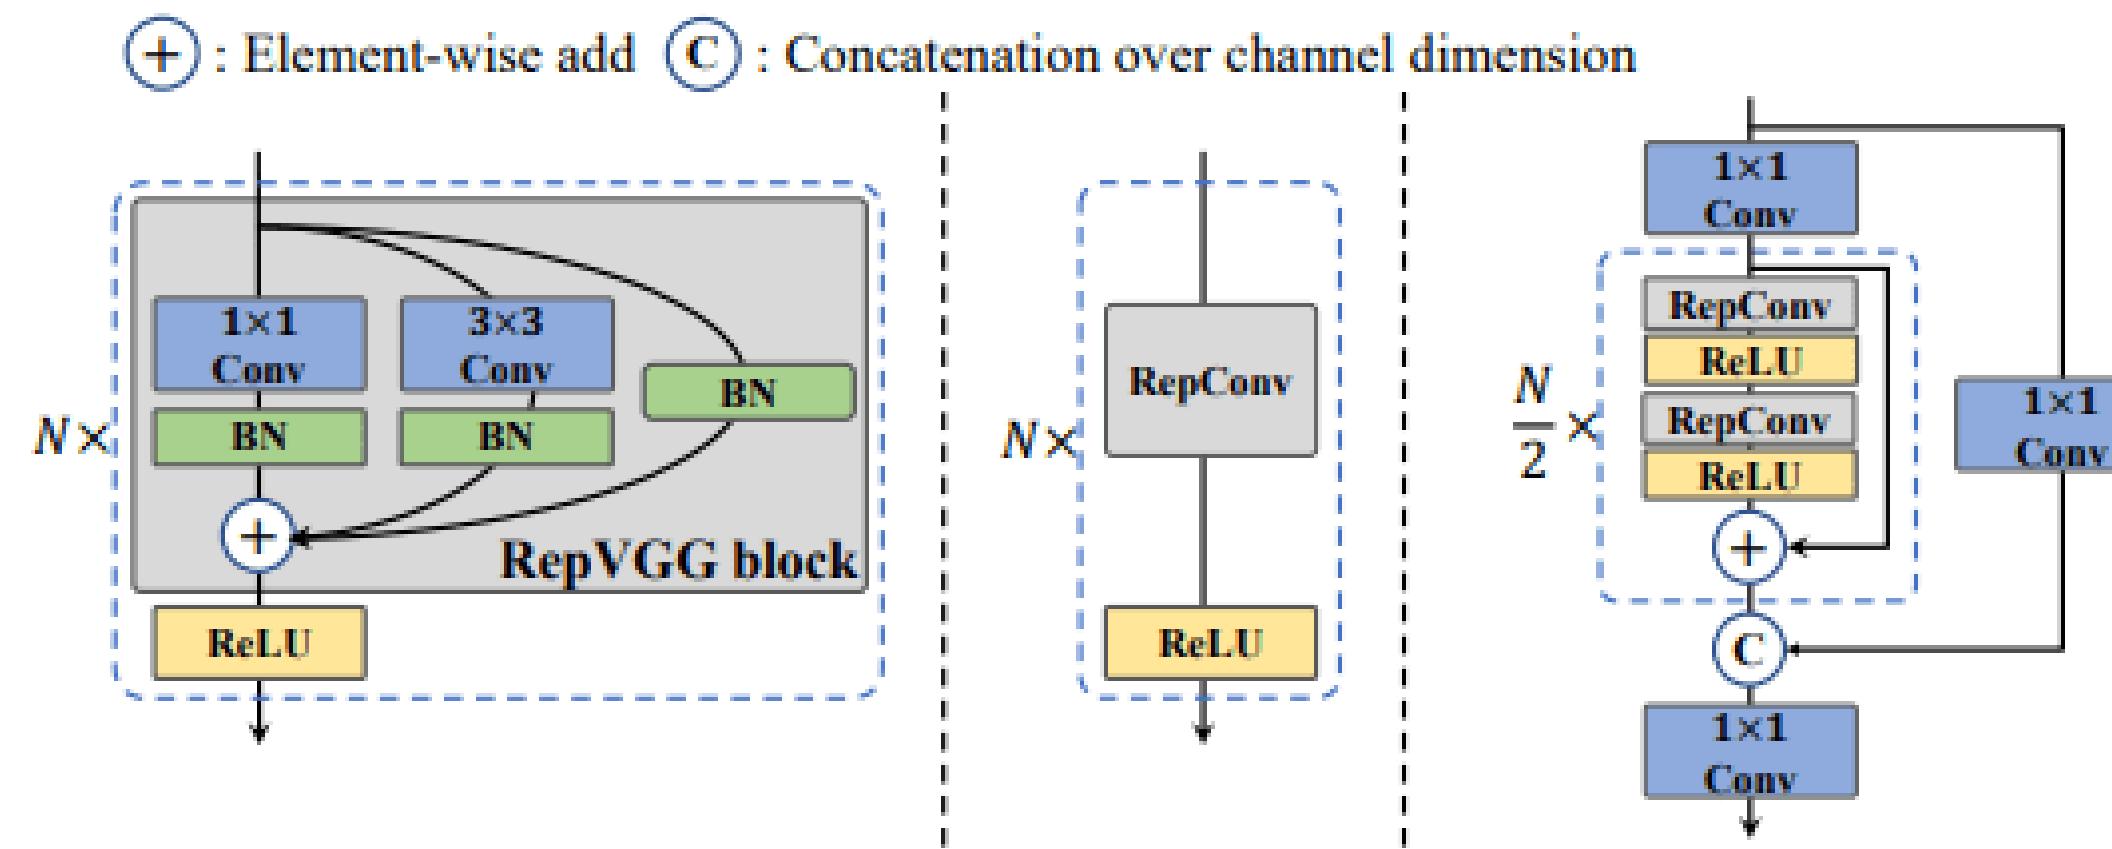
Instead of predicting the bounding boxes directly from scratch,  
**the model predicts adjustments** to these anchor boxes to better fit objects in the image.

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

# YOLOv6

## EfficientRep

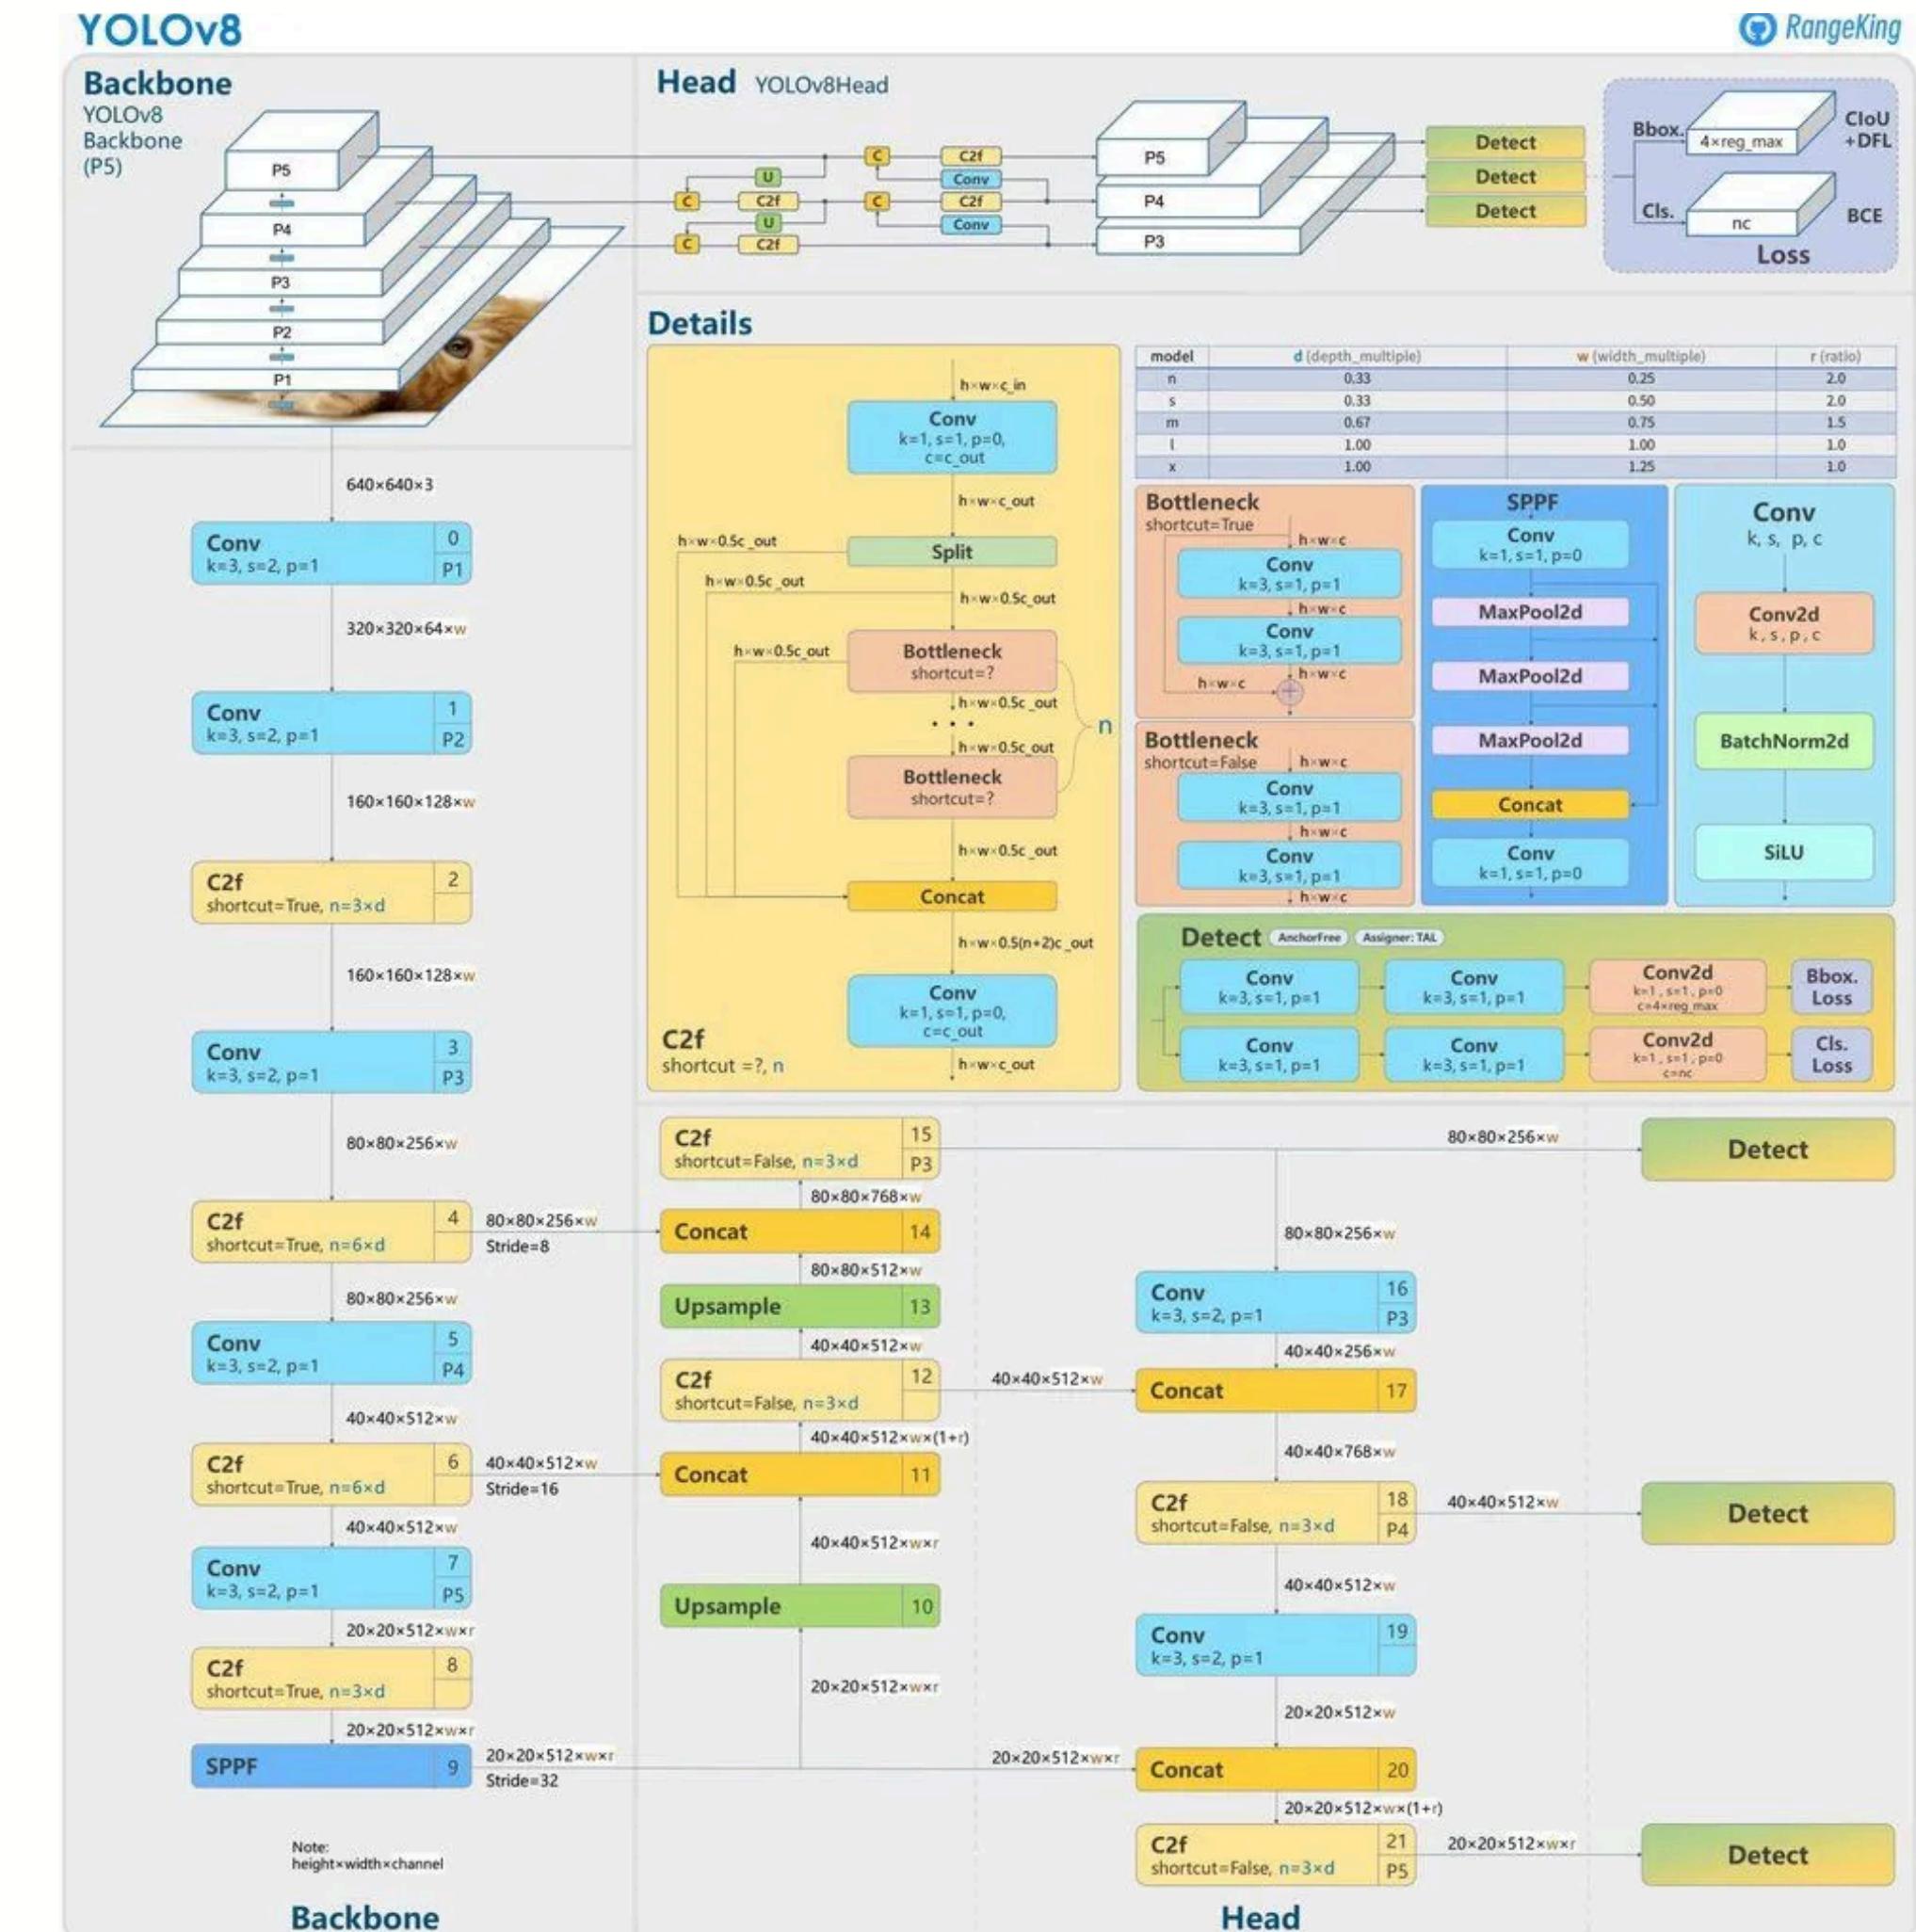
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



Compared with the CSP-Backbone used by YOLOv5, this backbone can efficiently utilize the computing power of hardware (such as GPU) and also has strong representation capabilities.

# YOLOv8

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



# YOLOv8

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

The YOLOv8 architecture follows the same architecture as YOLOv5, with a few slight adjustments, such as the use of the c2f module instead of CSPNet module, which is just a variant of CSPNet, (CSPNet followed by two convolutional networks).

They designed YOLOv8 anchor-free using YOLOv5 as a based model.

They also introduced Decoupled heads to independently process objectness, classification and bounding box prediction tasks. They used the sigmoid layer as the last layer for objectness score prediction and softmax for classification.

# YOLOv8

YOLOv8 uses CloU and DFL loss functions for bounding box loss and binary cross-entropy for classification loss. These losses have improved object detection performance, particularly when dealing with smaller objects. They also introduced the YOLOv8-Seg model for semantic segmentation by applying minimal changes to the original model.

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



# YOLOv9

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

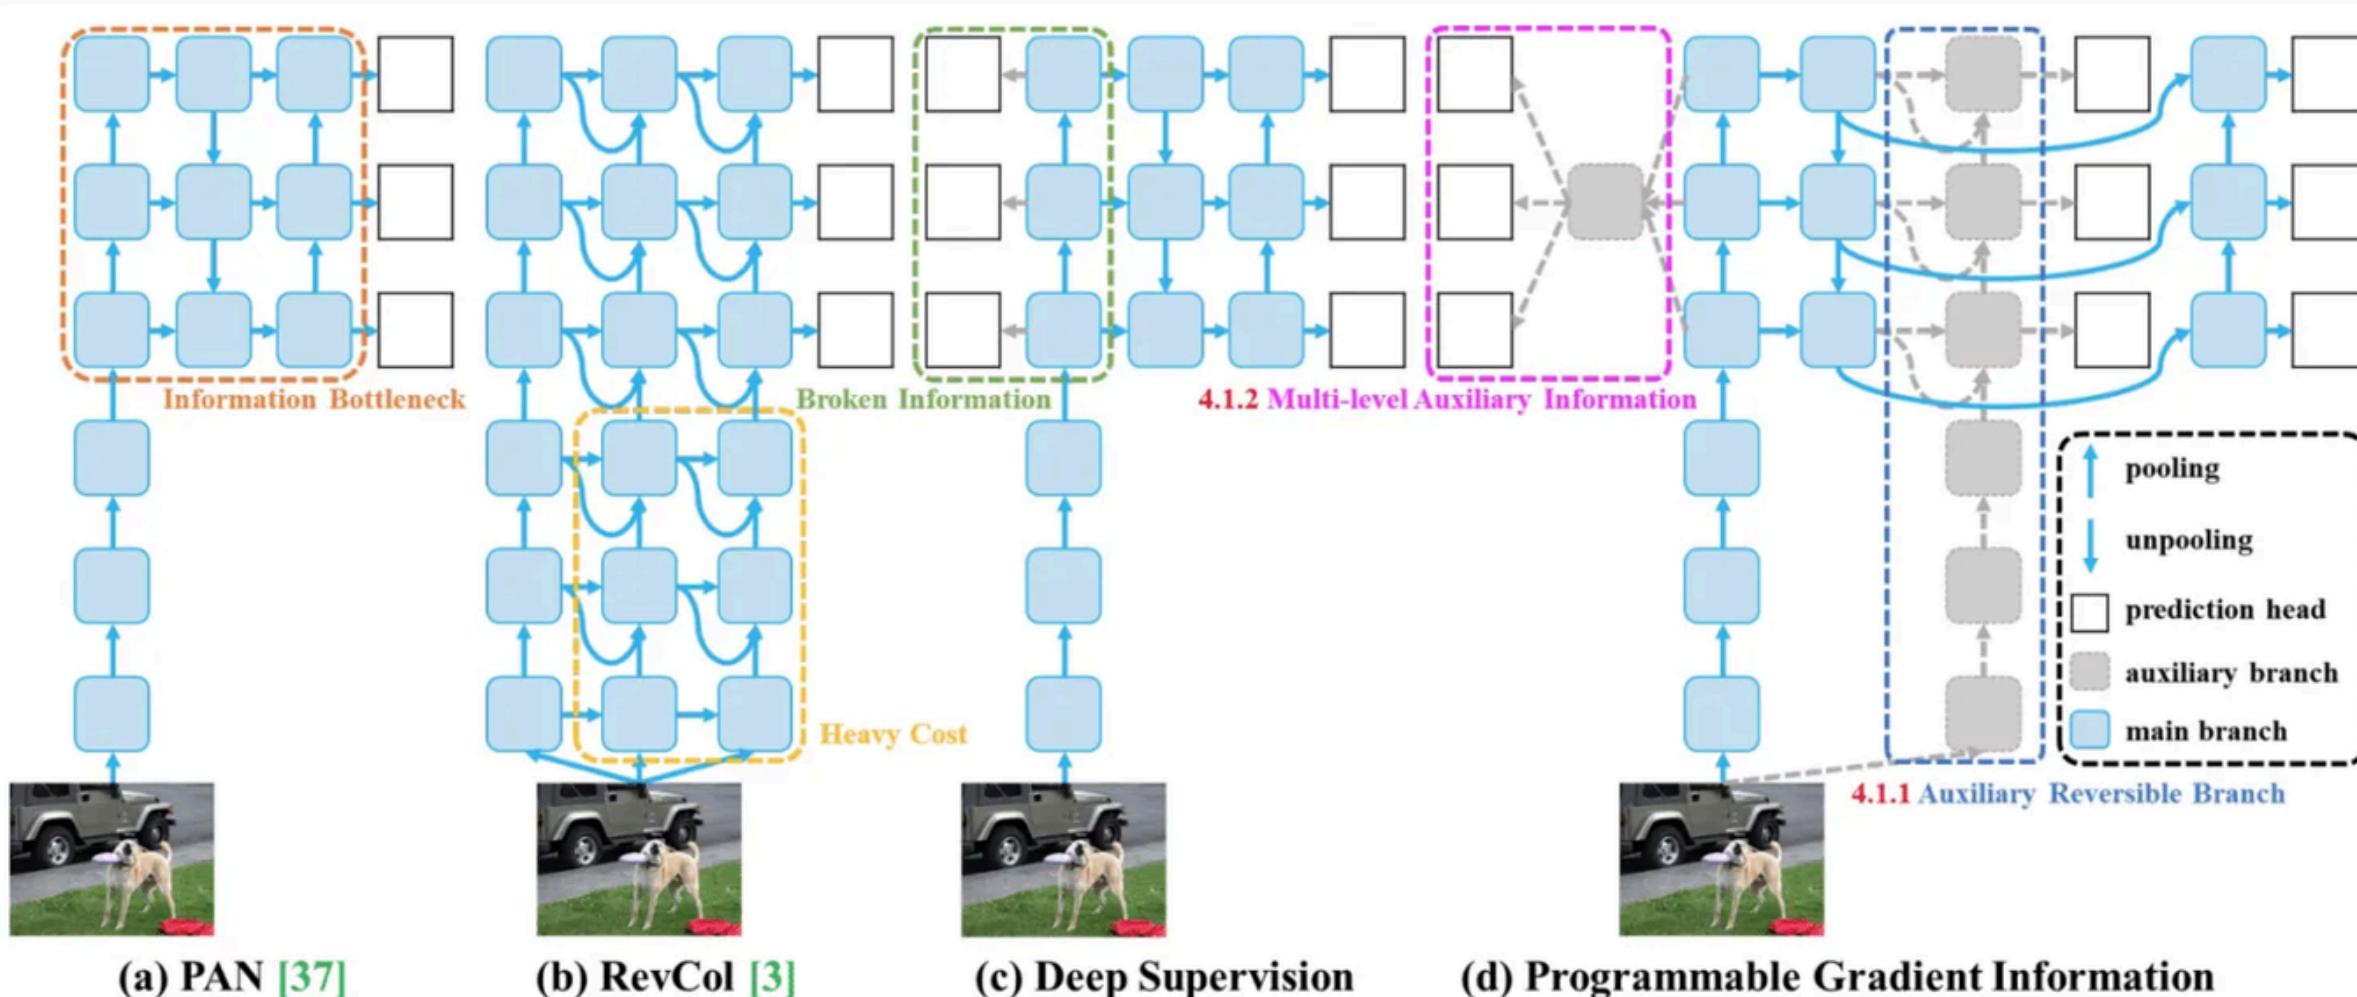


Figure 3. PGI and related network architectures and methods. (a) Path Aggregation Network (PAN) [37], (b) Reversible Columns (RevCol) [3], (c) conventional deep supervision, and (d) our proposed Programmable Gradient Information (PGI). PGI is mainly composed of three components: (1) main branch: architecture used for inference, (2) auxiliary reversible branch: generate reliable gradients to supply main branch for backward transmission, and (3) multi-level auxiliary information: control main branch learning plannable multi-level of semantic information.

YOLOv9 boasts two key innovations: the Programmable Gradient Information (PGI) framework and the Generalized Efficient Layer Aggregation Network (GELAN).

# YOLOv9

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

## The PGI framework:

PGI is a novel concept introduced in YOLOv9 to combat the information bottleneck problem, ensuring the preservation of essential data across deep network layers. This allows for the generation of reliable gradients, facilitating accurate model updates and improving the overall detection performance.



# YOLOv9

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

## The GELAN architecture:

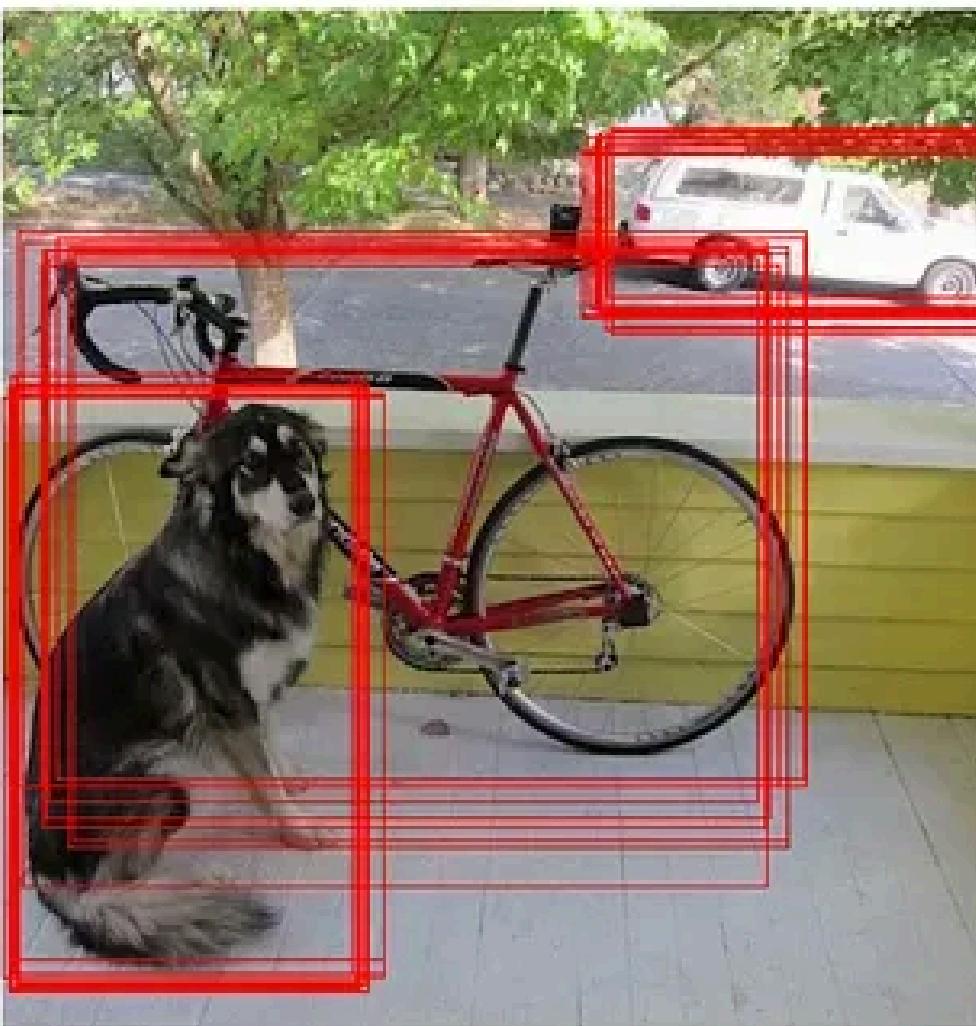
GELAN represents a strategic architectural advancement, enabling YOLOv9 to achieve superior parameter utilization and computational efficiency. Its design allows for flexible integration of various computational blocks, making YOLOv9 adaptable to a wide range of applications without sacrificing speed or accuracy.

By combining the above two frameworks (PGI and GELAN), YOLOv9 presents a significant advancement in lightweight object detection.

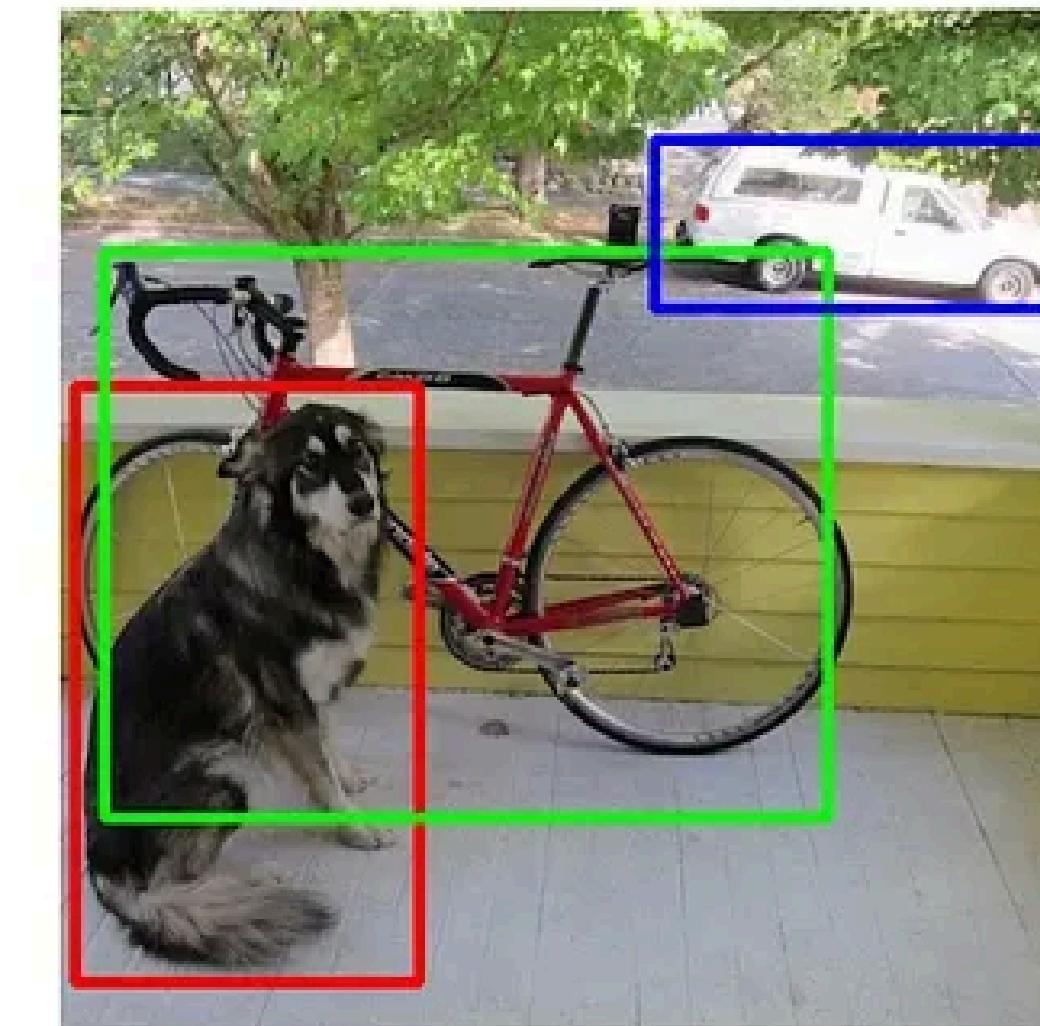
# YOLOv10

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10

Output Predictions



Non-Maximum Suppression (NMS)

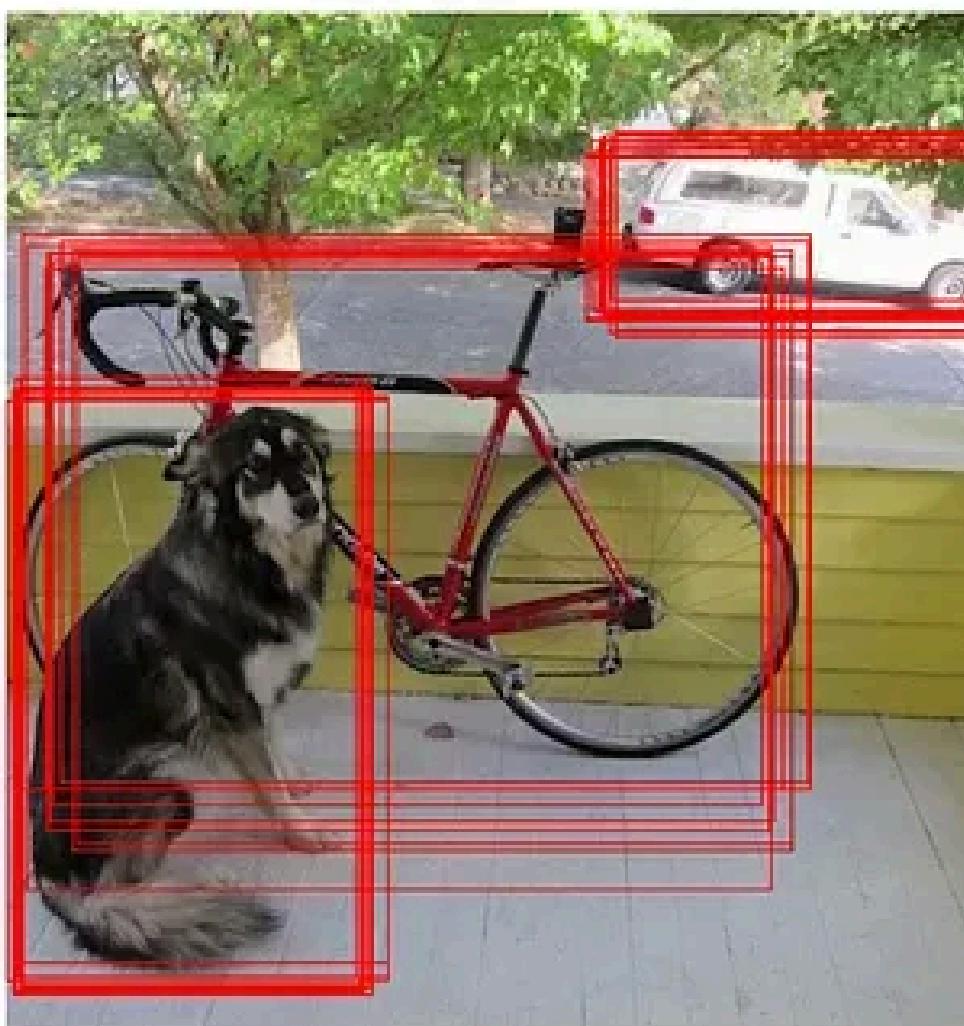


YOLOv10 distinguishes itself by completely eliminating the reliance on non-maximum suppression (NMS) during post-processing, which is a significant step forward in enhancing inference speed.

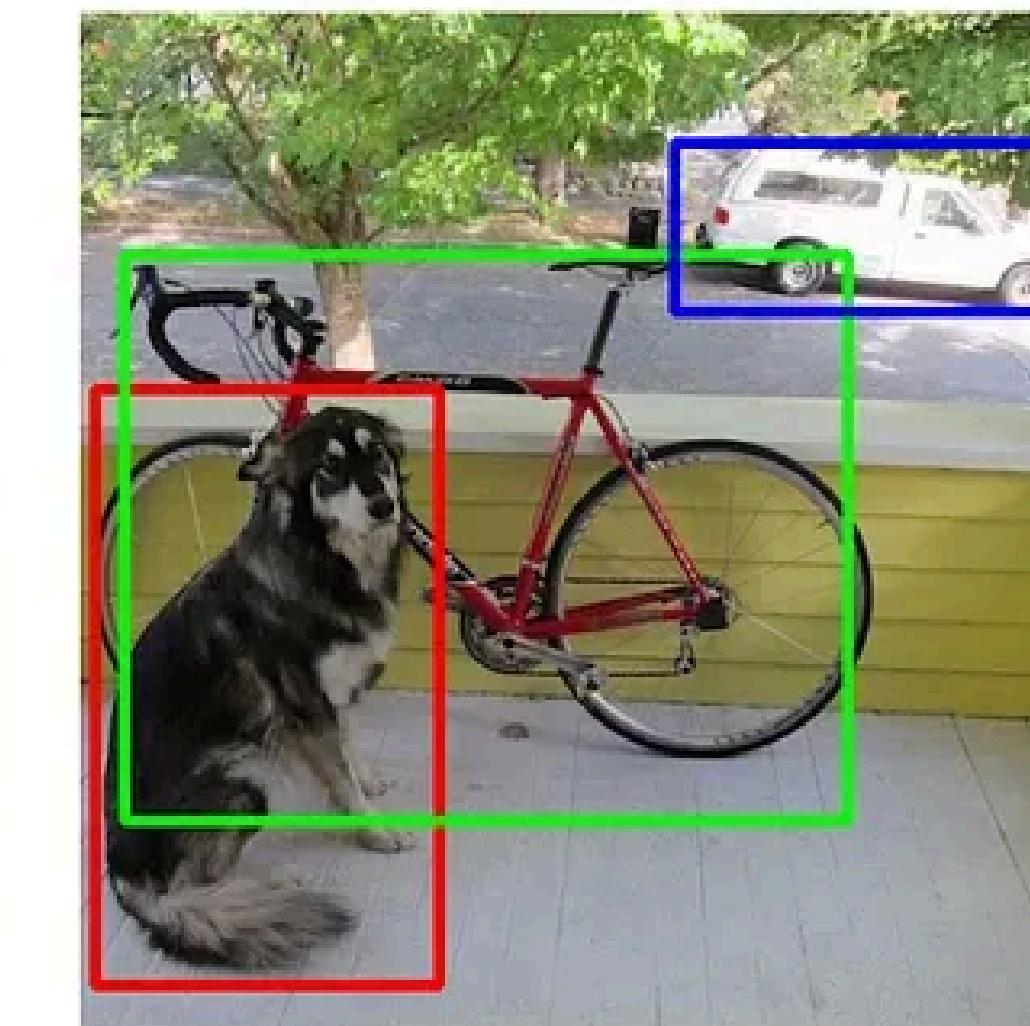
# YOLOv10

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10

Output Predictions



Non-Maximum Suppression (NMS)



This model adopts a novel NMS-free training approach using dual label assignments, allowing for a harmonious integration of accuracy and speed by ensuring that the model remains computationally efficient while still capturing essential detection features.



Version	Date	Contributions	Framework
v1	2015	One-shot object detector	Darknet
v2	2016	Multi-scale training, dimensional clustering	Darknet
v3	2018	SPP block, Darknet-53	Darknet
v4	2020	Mish-based activation, CSPDarknet-53 backbone	Darknet

1  
2  
3  
4  
5  
6  
7  
8  
9  
**10**

Three orange dots are located at the top right corner of the table.

	<b>Version</b>	<b>Date</b>	<b>Contributions</b>	<b>Framework</b>
1	v5	2020	Anchor-free detection, SWISH-based activation, PANet	PyTorch
2	v6	2022	Self-attention, anchor-free object detection	PyTorch
3	v7	2022	Transformers, E-ELAN reparameterization	PyTorch
4	v8	2023	GANs, anchor-free detections	PyTorch
5				
6				
7				
8				
9				
10				



	<b>Version</b>	<b>Date</b>	<b>Contributions</b>	<b>Framework</b>
1				
2				
3				
4				
5				
6				
7				
8				
9				
10	v9	2024	Programmable Gradient Information (PGI), Generalized Efficient Layer Aggregation Network (GELAN)	PyTorch
	v10	2024	NMS-free training approach, dual label assignments, holistic model design for enhanced accuracy and efficiency	PyTorch



# ***EMOTION DETECTION USING YOLO***



1

2

3

4

5

6

7

8

9

# Context

Emotion detection is key in **security, customer service, and mental health.**

Why YOLO:

- YOLO offers fast, efficient detection, ideal for real-time needs.
- Adapting YOLO enables instant emotional recognition in live video.

# Objective

Detect and classify people emotions in real time,  
achieve a certain level of accuracy

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



# Data Understanding

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

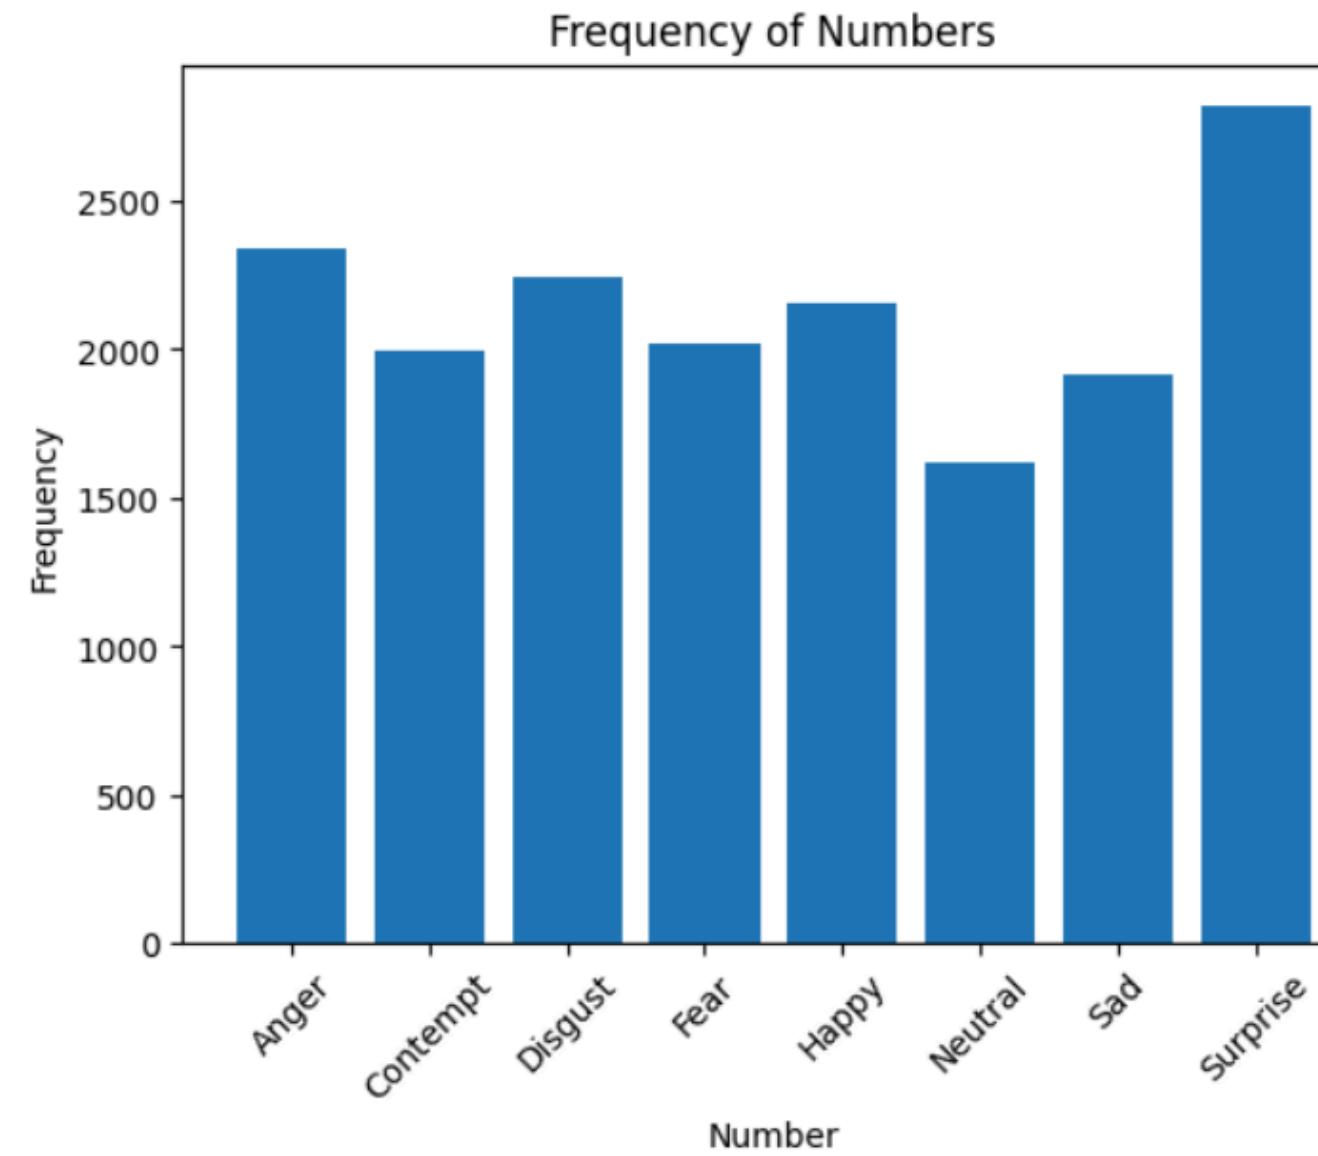
```
- affectnet-yolo-format
  - YOLO_format
    - test
      - images
      - labels
    - train
      - images
      - labels
    - valid
      - images
      - labels
  - data.yaml
```



# Data Understanding

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

```
- affectnet-yolo-format
  - YOLO_format
    - test
      - images
      - labels
    - train
      - images
      - labels
    - valid
      - images
      - labels
  - data.yaml
```



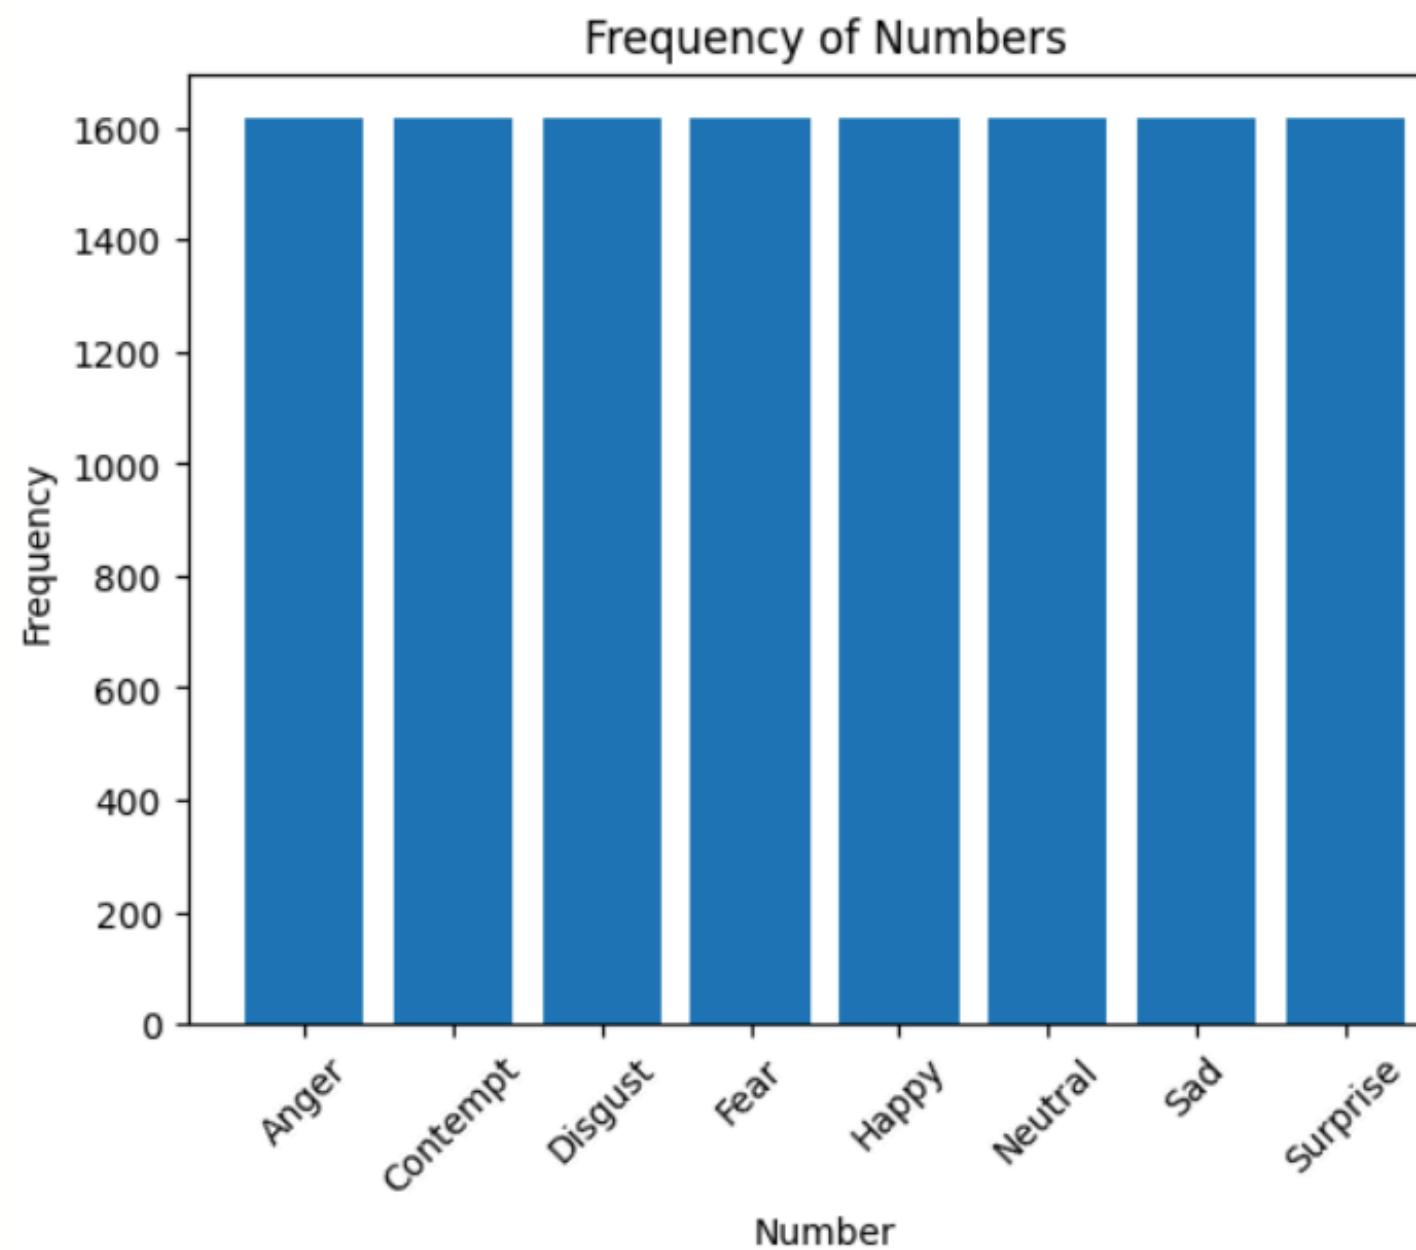
Valid Files: 5406

Test Files: 2755

Train Files: 17101

# Data Preparation

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



Balanced dataset creation:

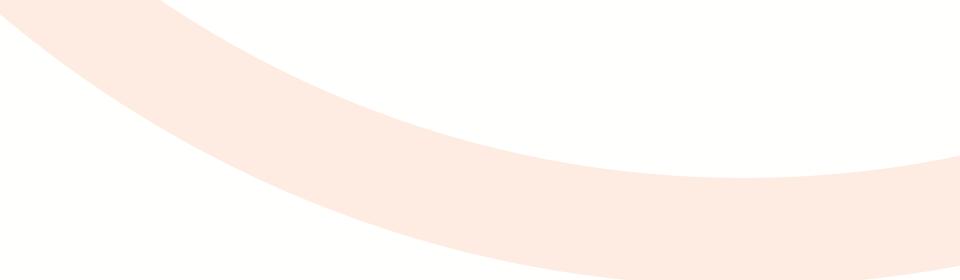
- Created a new yolo-data folder.
- Copied the valid and test folders.
- Balanced the train folder.
- Updated the data path in data.yaml.

# HyperParameter Fine-tuning

Using **Optuna** to test **20 combinations** of hyperparameters and maximize the model's performance the validation metric (**mAP**).

- Learning rate: between 0.0001 et 0.1 (log-uniform).
- Momentum: between 0.8 and 0.98.
- Weight decay: between 0.000001 et 0.001 (log-uniform).

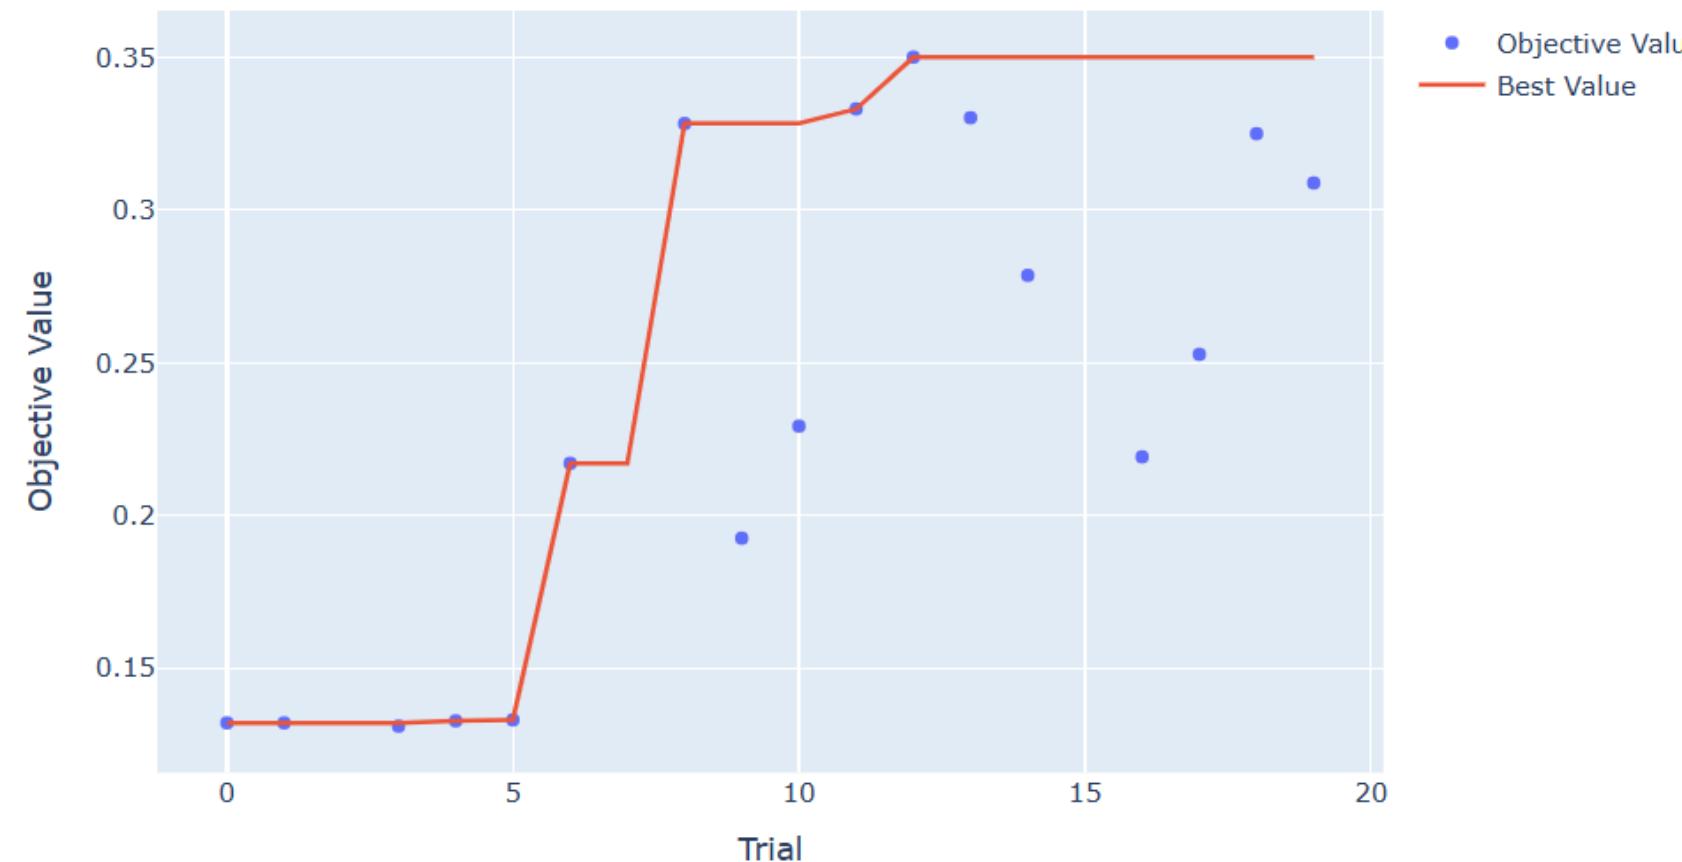
Training the model for **2 epochs** to evaluate its performance.

- 
- 
- 1
  - 2
  - 3
  - 4
  - 5
  - 6
  - 7
  - 8
  - 9

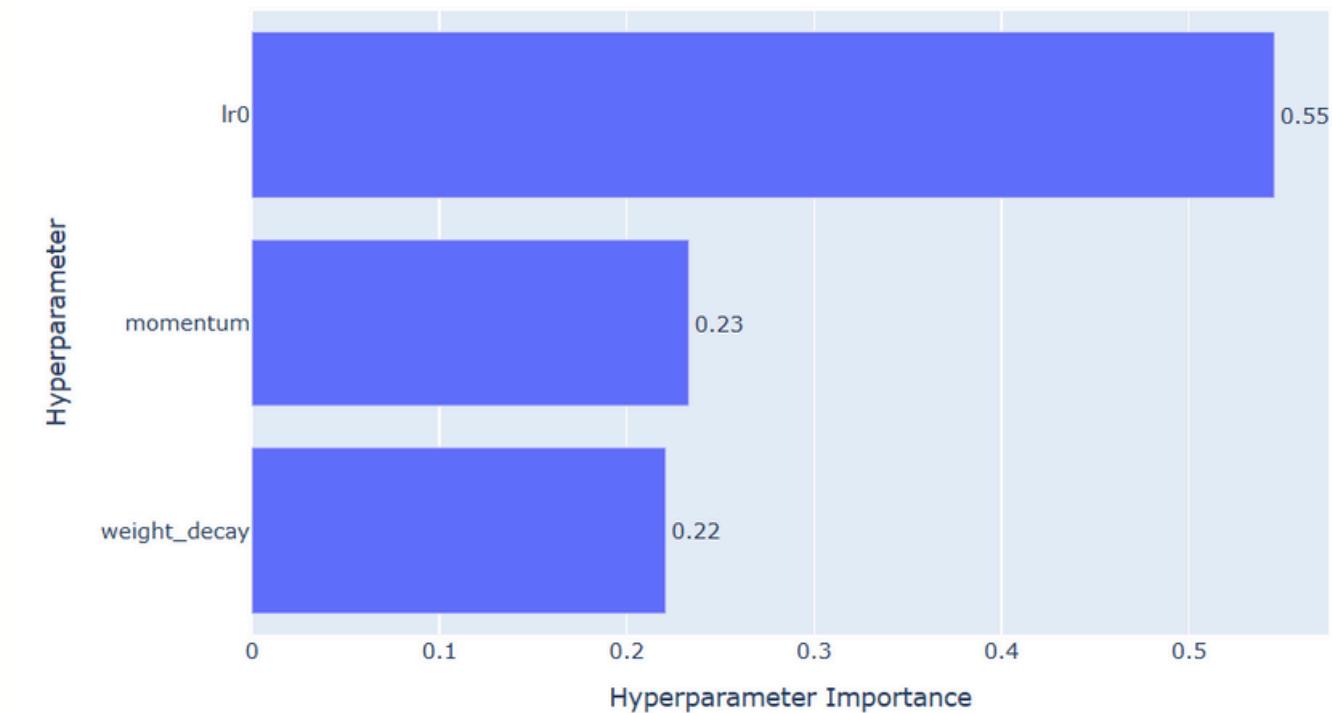
# HyperParameter Fine-tuning

- Lr0: 0.00010446991982747811
- momentum: 0.881082406956181
- weight\_decay: 7.273003893562258e-06

Optimization History Plot



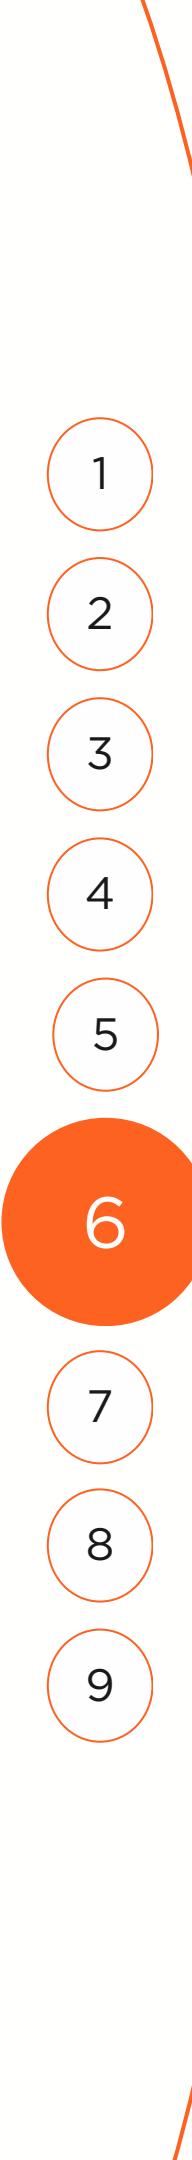
Hyperparameter Importances



1  
2  
3  
4  
5  
6  
7  
8  
9

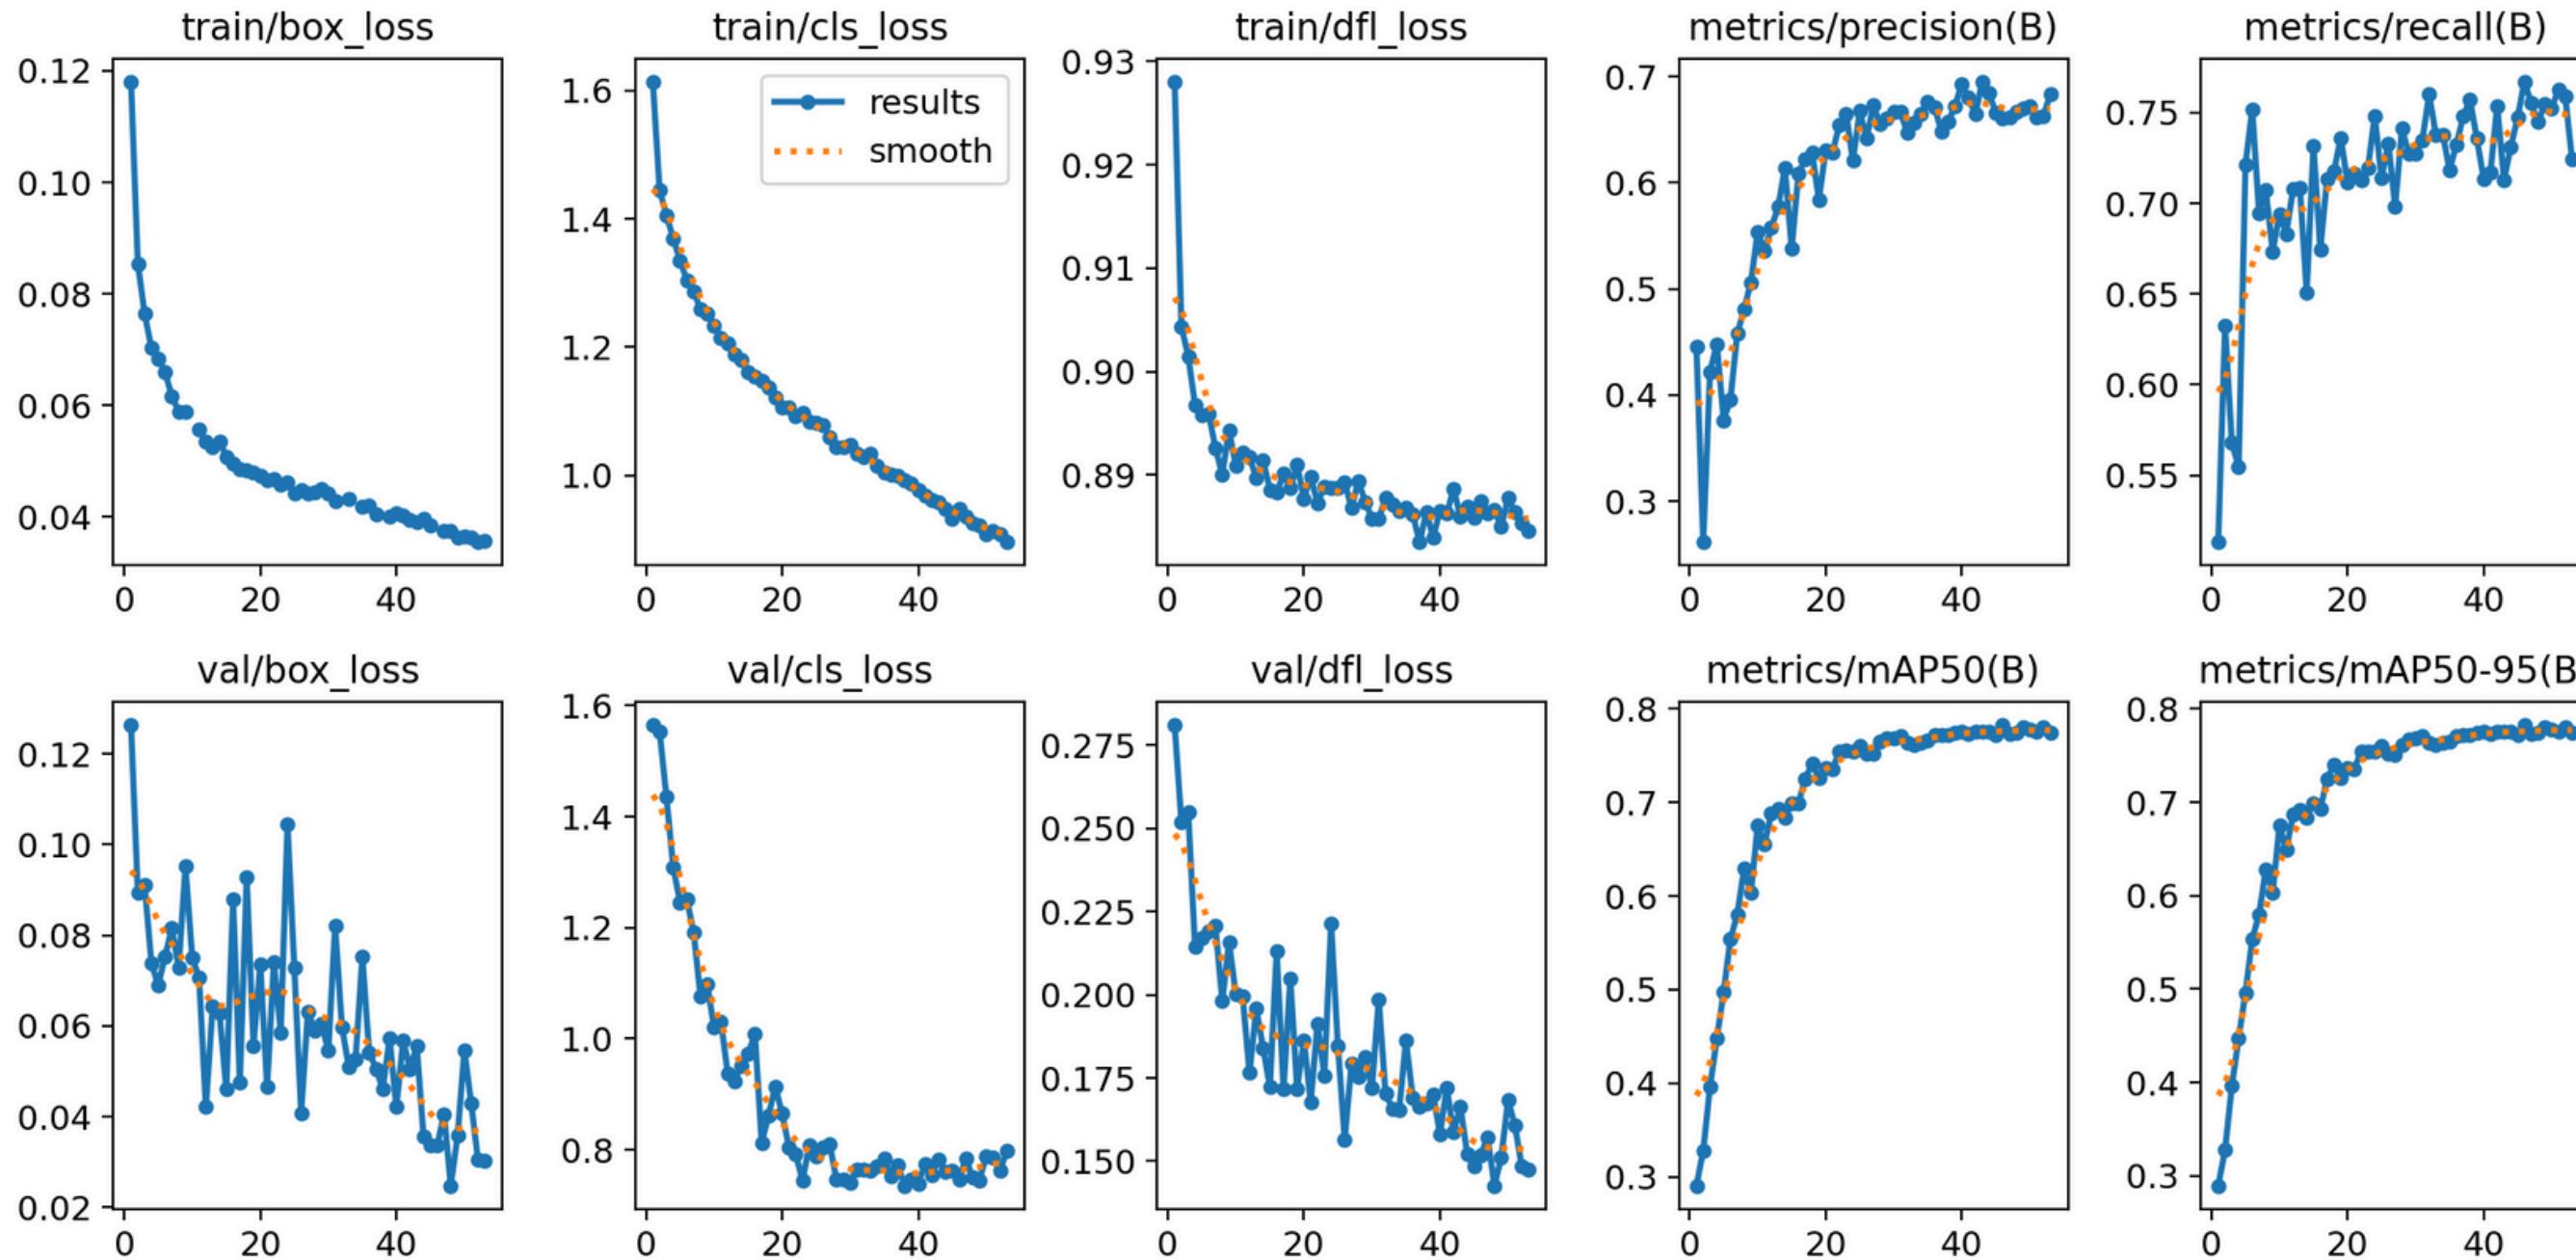
# Model Training

Optuna automatically identifies the best hyperparameter combination, which is then used to train the model for **100 epochs** with an early stopping **patience of 7 epochs**.

- 
- 1
  - 2
  - 3
  - 4
  - 5
  - 6
  - 7
  - 8
  - 9

# Model Evaluation

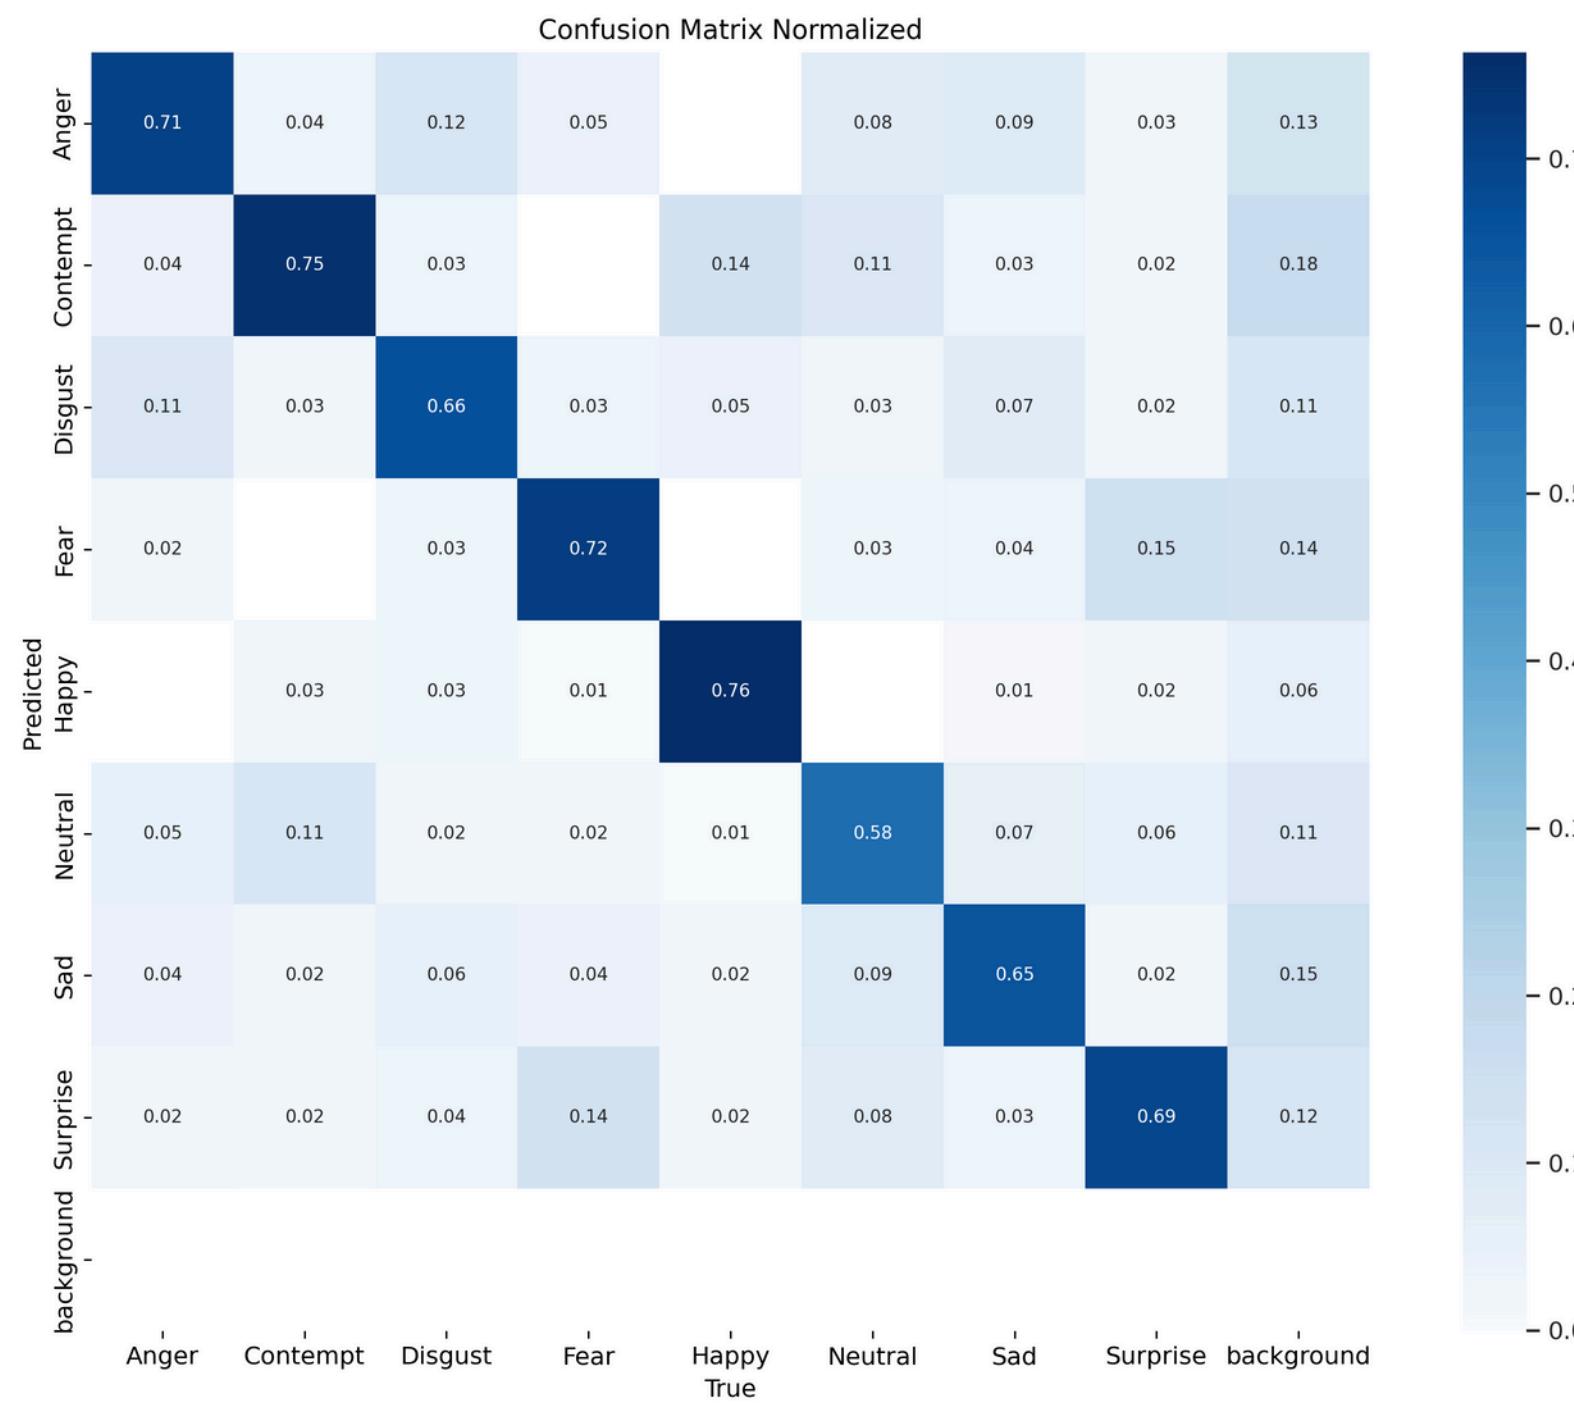
YOLO automatically logs metrics like loss and mAP during training.



- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

# Model Evaluation

We get the mAP: 0.78

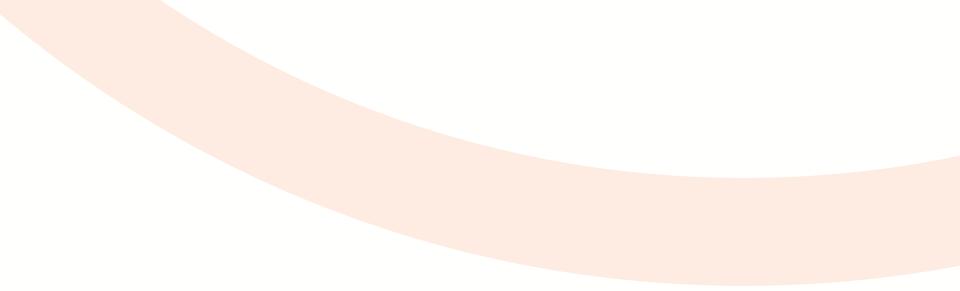


- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

# Testing





# A local demonstration application.

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



**THANK  
YOU**